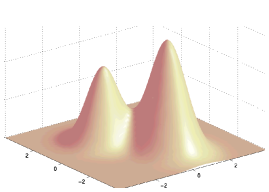
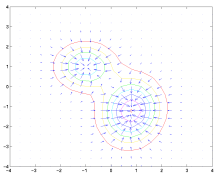
Mathematics of Machine Learning
Optimisation

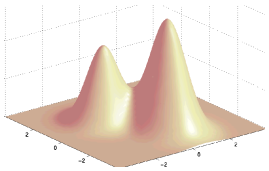



$$z = \mathrm{e}^{-(x+1)^2-(y-1)^2} + 0.6\mathrm{e}^{-(x-1)^2-0.5(y+1)^2+0.1(x-3)(y-3)}$$

Gradient descent, quadratic minima, differing length scales

Outline

1. **Motivation**
2. Gradient Descent
3. Why Gradient Descent is Difficult

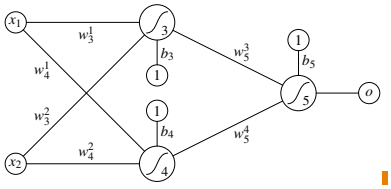


ML = Optimisation

- Many learning machines can be thought of as functions of the form
$$\hat{y} = f(\boldsymbol{x}|\boldsymbol{w})$$
(or more generally $\hat{\boldsymbol{y}} = \boldsymbol{f}(\boldsymbol{x}|\boldsymbol{w})$)
- Given an input pattern (set of features) $\boldsymbol{x}$ the learning machine makes a prediction $\hat{y}$
- We try to choose the parameters $\boldsymbol{w}$ so that the predictions are good
- In practice training a learning machine comes down to optimising some loss function

MLP

- We can depict a neural network such as an MLP by a diagram



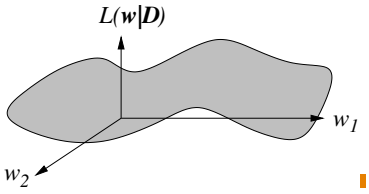
- Stands for the function ($\hat{y} = f(\boldsymbol{x}|\boldsymbol{w})$)
$$\hat{y} = g(w_5^3 g(w_3^1 x_1 + w_3^2 x_2 + b_3) + w_5^4 g(w_4^1 x_1 + w_4^2 x_2 + b_4) + b_5)$$
where, for example, $g(V) = \frac{1}{1+\mathrm{e}^{-V}}$

Training

- Given a (labelled) training dataset
$$\mathcal{D} = \{(\boldsymbol{x}_k, y_k) | k = 1, \ldots, m\}$$
- We define an error or loss function that we want to minimise
$$L(\boldsymbol{w}|\mathcal{D}) = \frac{1}{m}\sum_{k=1}^{m}\left(f(\boldsymbol{x}_k|\boldsymbol{w}) - y_k\right)^2$$
- We then use the machine with the weights $\boldsymbol{w}^*$ which minimise $L(\boldsymbol{w}|\mathcal{D})$

Computing Gradients

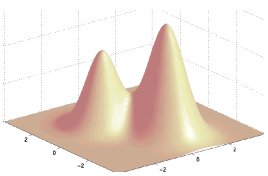
- $L(\boldsymbol{w}|\mathcal{D})$ is a complex function of the weights $\boldsymbol{w}$



- To minimise $L(\boldsymbol{w}|\mathcal{D})$ we compute the gradient $\nabla L(\boldsymbol{w}|\mathcal{D})$
- For neural networks an efficient algorithm for computing the gradient is known as back-propagation of errors (*back-prop*)

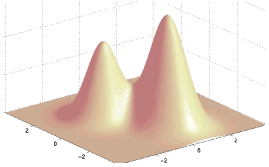
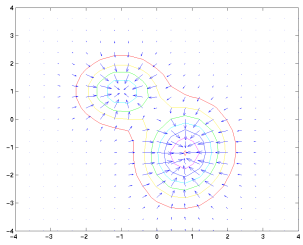
Outline

1. Motivation
2. **Gradient Descent**
3. Why Gradient Descent is Difficult



Gradient Optimisation

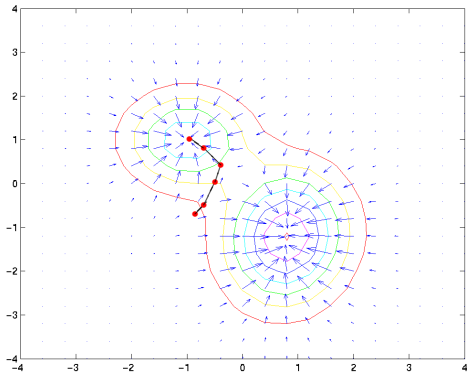
- At a maximum or minimum (or a saddle point) $\nabla L(\boldsymbol{w}|\mathcal{D}) = \boldsymbol{0}$
- E.g.
$$L = \mathrm{e}^{-(x+1)^2-(y-1)^2} + 0.6\mathrm{e}^{-(x-1)^2-0.5(y+1)^2+0.1(x-3)(y-3)}$$

Gradient Descent

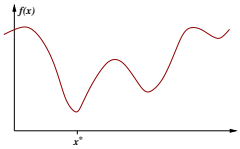
- For a simple function $L(\boldsymbol{w}|\mathcal{D})$ we can solve $\nabla L(\boldsymbol{w}|\mathcal{D}) = \mathbf{0}$ explicitly. E.g. the linear perceptron.
- For non-linear functions we usually can't solve this set of simultaneous equations.
- We can find a maximum or minimum **iteratively**.
- If we know the gradient then we can follow the gradient.
  - Maximisation: $\boldsymbol{w} \rightarrow \boldsymbol{w}' = \boldsymbol{w} + r\nabla L(\boldsymbol{w}|\mathcal{D})$
  - Minimisation: $\boldsymbol{w} \rightarrow \boldsymbol{w}' = \boldsymbol{w} - r\nabla L(\boldsymbol{w}|\mathcal{D})$

Hill-Climbing



What Goes Right

- Almost all minima are quadratic (Morse's theorem).



- Taylor expanding around a minimum x^*

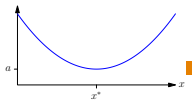
$$\begin{aligned} f(x) &= f(x^*) + (x-x^*)f'(x^*) + \frac{1}{2}(x-x^*)^2 f''(x^*) + \cdots \\ &= f(x^*) + \frac{1}{2}(x-x^*)^2 f''(x^*) + \frac{1}{3!}(x-x^*)^3 f'''(x^*) + \cdots \end{aligned}$$

- If $x - x^*$ is sufficiently small the higher order terms are negligible.

Newton's Method

- If we were in a quadratic minimum

$$f(x) = a + \frac{b}{2}(x-x^*)^2$$



- then

$$f'(x) = b(x-x^*), \qquad f''(x) = b$$

- so

$$x - x^* = \frac{f'(x)}{b} = \frac{f'(x)}{f''(x)}$$

- or

$$x^* = x - \frac{f'(x)}{f''(x)}$$

Newton's Method

- This is Newton's method.
- For non-quadratic functions Newton's method converges **quadratically** provided we are sufficiently close to a minimum.
- If we are at a distance $x - x^* = \epsilon$ from the minimum then after one iteration we will be a distance of order ϵ^2, after two iterations of order ϵ^4, etc.
- If we are too far from a minimum we might go anywhere!
- We should follow the gradient until we are near the minimum.

Taylor's Expansion in High Dimensions

- We can generalise these results to many dimensions.
- The Taylor expansion of a function $f(\boldsymbol{x})$ about $\boldsymbol{x}_0$

$$f(\boldsymbol{x}) = f(\boldsymbol{x}_0) + (\boldsymbol{x}-\boldsymbol{x}_0)^\mathsf{T}\nabla f(\boldsymbol{x}_0) + \frac{1}{2}(\boldsymbol{x}-\boldsymbol{x}_0)^\mathsf{T}\mathbf{H}(\boldsymbol{x}-\boldsymbol{x}_0) + \cdots$$

  where $\mathbf{H}$ is the **Hessian** matrix with elements

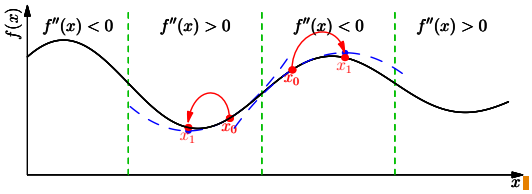
$$H_{ij} = \frac{\partial^2 f(\boldsymbol{x}_0)}{\partial x_i \partial x_j}$$

- Newton's method in high dimension is

$$\boldsymbol{x}^* = \boldsymbol{x} - \mathbf{H}^{-1}\nabla f(\boldsymbol{x})$$

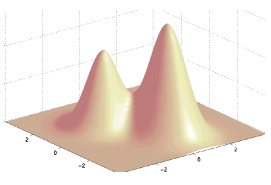
Using the Second Derivative

- If we are optimising n parameters the Hessian is an $n \times n$ matrix.
- It is time-consuming to compute (and prone to errors when coding)—for deep learning it is impossible even to store the Hessian.
- Away from minima they can be misleading



Outline

1. Motivation
2. Gradient Descent
3. **Why Gradient Descent is Difficult**



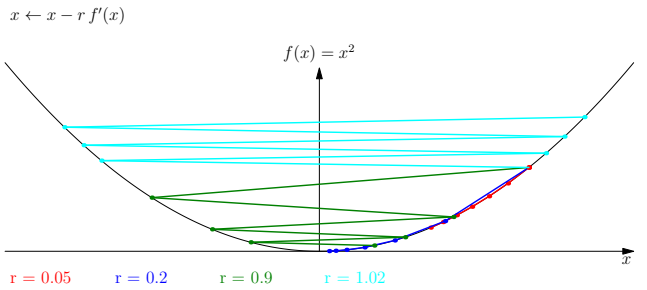
Step Size

- Gradient descent

$$\boldsymbol{x}' = \boldsymbol{x} - r\nabla f(\boldsymbol{x})$$

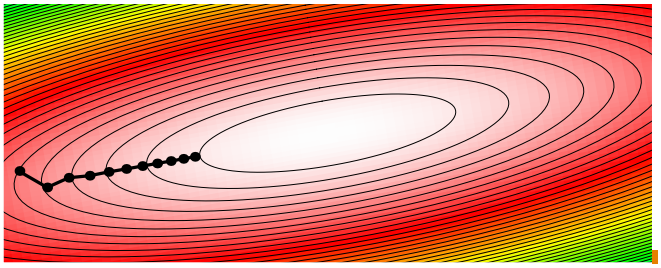
- Need to choose the learning rate or step size, r
- Too small a step takes lots of time
- Too large a step takes you away from a minimum

Step Size



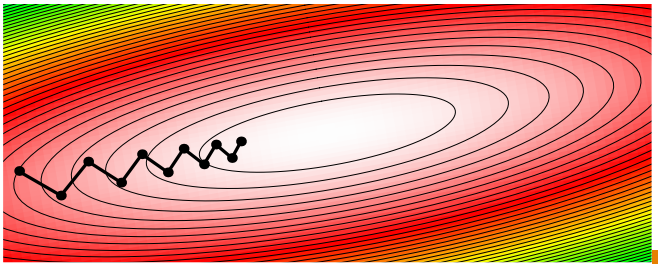
Higher Dimensions

- In higher dimensions the problem is that there are some directions you need to move a long way



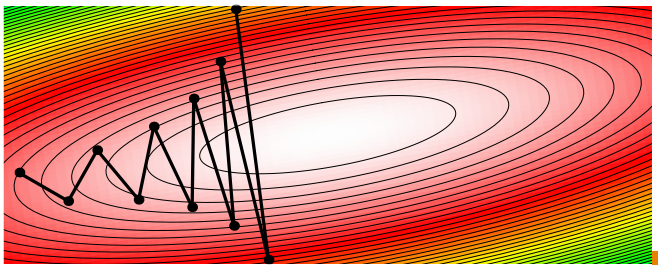
Getting There Quicker

- Increasing the step size speeds up convergence, but the direction of steepest descent doesn't point to the minimum



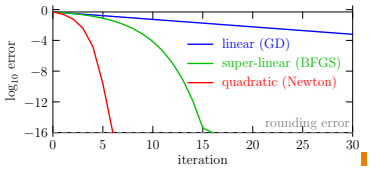
More Haste Less Speed

- Increasing the step size, just a little further, increases the rate of convergence in one direction, but . . .



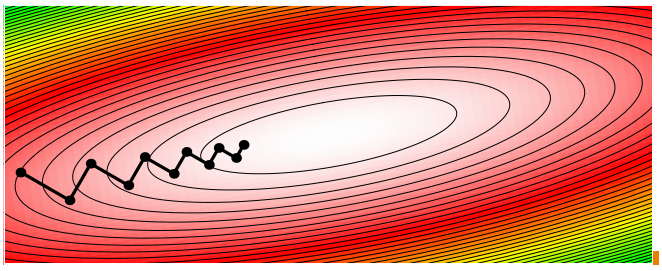
Linear and Quadratic Convergence

- Near a minimum, along each eigenvector of the Hessian, a gradient step multiplies the error by $(1 - r\lambda_i)$
- Stable only if $r < 2/\lambda_{\max}$, so the flattest direction shrinks by at best $1 - 2\lambda_{\min}/\lambda_{\max} = 1 - 2/\kappa$ per step (κ is the condition number)
- A constant factor per step is **linear** convergence; Newton squares the error: **quadratic** convergence



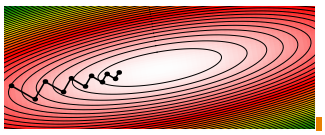
Line Minimisation

- We can systematically seek the minimum along a line of the gradient



Zig-Zag

- Note that in narrow valleys line minimisation zigzags (each step is at right angles to the last)



- If we computed the Hessian and used Newton's method we would jump straight to the minimum if we were in a quadratic potential
- However computing the Hessian is time consuming and misleading if we are not in a quadratic potential (i.e. far from the optimum)

Better Optimisation Algorithms

- Good optimisation algorithms use second-order (curvature) information without computing the full Hessian
- E.g. Conjugate gradient
  - ⋆ Performs Line Minimisation
  - ⋆ Uses gradient, but does not go along it
  - ⋆ For a quadratic minimum in d dimensions it reaches the minimum in d steps
- E.g. Levenberg-Marquardt
  - ⋆ Used on least squares problem only
  - ⋆ Uses linear approximation of function to approximate Hessian
  - ⋆ Adapts from hill-climbing to Newton method
  - ⋆ Avoids line-minimisation

BFGS and L-BFGS

- BFGS (Broyden–Fletcher–Goldfarb–Shanno) is the most widely used quasi-Newton method
- It builds an approximation $\mathbf{B} \approx \mathbf{H}^{-1}$ from how the gradient changes from one step to the next

$$\mathbf{B}\big(\boldsymbol{\nabla} f(\boldsymbol{x}_{t+1}) - \boldsymbol{\nabla} f(\boldsymbol{x}_t)\big) = \boldsymbol{x}_{t+1} - \boldsymbol{x}_t$$

- Steps $\boldsymbol{x}_{t+1} = \boldsymbol{x}_t - \alpha \mathbf{B} \boldsymbol{\nabla} f(\boldsymbol{x}_t)$, with a line search for α
- Converges **super-linearly** without ever computing the Hessian
- $\mathbf{B}$ is $N \times N$; **L-BFGS** keeps only the last ten or so steps instead, so its memory is $O(N)$

Levenberg-Marquardt

- Want to minimise $\|\boldsymbol{\epsilon}(\boldsymbol{w})\|^2$ where $\epsilon_i(\boldsymbol{w}) = f(\boldsymbol{x}_i|\boldsymbol{w}) - y_i$
- Use linear approximation

$$\epsilon_i(\boldsymbol{w}) \approx \epsilon_i(\boldsymbol{w}^{(k)}) + (\boldsymbol{w} - \boldsymbol{w}^{(k)})^\mathsf{T} \boldsymbol{\nabla} \epsilon_i(\boldsymbol{w}^{(k)})$$

  with $\boldsymbol{\nabla} \epsilon_i(\boldsymbol{w}^{(k)}) = \boldsymbol{\nabla} f(\boldsymbol{x}_i|\boldsymbol{w}^{(k)})$
- Solve quadratic minimisation of approximate error $\operatorname{argmin}_{\boldsymbol{w}} L_{approx}(\boldsymbol{w})$ with $\mathbf{J} = \boldsymbol{\nabla}\boldsymbol{\epsilon}(\boldsymbol{w}^{(k)})$

$$\begin{aligned} L_{approx}(\boldsymbol{w}) &= \|\boldsymbol{\epsilon}(\boldsymbol{w}^{(k)}) + \mathbf{J}(\boldsymbol{w} - \boldsymbol{w}^{(k)})\|^2 \\ &= \boldsymbol{\epsilon}(\boldsymbol{w}^{(k)})^\mathsf{T}\boldsymbol{\epsilon}(\boldsymbol{w}^{(k)}) + 2(\boldsymbol{w} - \boldsymbol{w}^{(k)})^\mathsf{T}\mathbf{J}^\mathsf{T}\boldsymbol{\epsilon}(\boldsymbol{w}^{(k)}) \\ &\quad + (\boldsymbol{w} - \boldsymbol{w}^{(k)})^\mathsf{T}\mathbf{J}^\mathsf{T}\mathbf{J}(\boldsymbol{w} - \boldsymbol{w}^{(k)}) \end{aligned}$$

Trust Region

- Solution given by $\boldsymbol{\nabla}_{\boldsymbol{w}} L_{approx}(\boldsymbol{w}) = 0$ gives

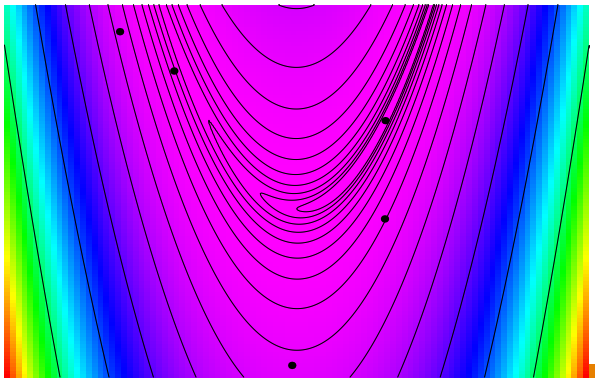
$$\boldsymbol{w}^{(k+1)} = \boldsymbol{w}^{(k)} - \big(\mathbf{J}^\mathsf{T}\mathbf{J}\big)^{-1}\mathbf{J}^\mathsf{T}\boldsymbol{\epsilon}(\boldsymbol{w}^{(k)})$$

- Can lead us in the wrong direction
- Instead use $\boldsymbol{w}^{(k+1)} = \operatorname{argmin}_{\boldsymbol{w}} L_{approx}(\boldsymbol{w}) + \nu\|\boldsymbol{w} - \boldsymbol{w}^{(k)}\|^2$

$$\boldsymbol{w}^{(k+1)} = \boldsymbol{w}^{(k)} - \big(\mathbf{J}^\mathsf{T}\mathbf{J} + \nu\mathbf{I}\big)^{-1}\mathbf{J}^\mathsf{T}\boldsymbol{\epsilon}(\boldsymbol{w}^{(k)})$$

- ν limits the step size
- If predicted reduction in error is accurate then reduce ν, else if predicted reduction in error is very poor increase ν

Rosenbrock: $\epsilon_1 = 10(x_2 - x_1^2)$, $\epsilon_2 = 1 - x_1$



Where Are These Used?

- For up to thousands of parameters (L-BFGS: millions), with gradients computed from all the data
- **Logistic regression**: scikit-learn's default solver is L-BFGS; Newton's method for generalised linear models
- **Gaussian processes**: fitting the kernel hyper-parameters with L-BFGS
- **Curve fitting** (non-linear least squares): Levenberg–Marquardt, e.g. SciPy's `curve_fit`
- **Support vector machines**: a quadratic programme (later lectures)
- With millions to billions of parameters and mini-batches of data we need different methods: next lecture

Summary

- There are some **non-gradient methods** (Nelder Mead, evolutionary strategies, Powell's method), but in very high dimensions these are not very competitive
- There are **gradient methods** (first order methods) that suffer from the problem of having to choose a single step size with conflicting requirements in different directions; they converge **linearly**
- **Newton's method** (a second order method) requires computing the Hessian matrix, converges **quadratically**, but can take you in the wrong direction if you are not sufficiently close to a minimum
- There exist a number of **quasi-Newton methods** (BFGS, conjugate gradient, Levenberg-Marquardt, etc.) that approximate Newton's method often without explicitly computing the Hessian