


Principal Component Analysis (PCA)

1.6 -1.1 -1.6 2.1 -0.52 2.8 0.72 0.7 -0.68 -0.41 -1.4 -1.5 -0.54 -0.62 1.3 -1.4 -0.27 0.74 0.77 -1




Covariance matrices, dimensionality reduction, PCA, Duality

Adam Prügel-Bennett AICE3001/6002 Advanced Machine Learning <https://ecs-vlc.github.io/aice3001/> 1

Spread of Data

- Often data varies significantly in only some directions




- Reduce dimensions by projecting onto low dimensional subspace with maximum variation

Adam Prügel-Bennett AICE3001/6002 Advanced Machine Learning <https://ecs-vlc.github.io/aice3001/> 3

Dimensionality Reduction

- Often helpful to consider only directions where data varies significantly
- Want to find directions along which data has its greatest variation



Adam Prügel-Bennett AICE3001/6002 Advanced Machine Learning <https://ecs-vlc.github.io/aice3001/> 5

Direction of Maximum Variation

- Expanding the Lagrangian


$$\begin{aligned}\mathcal{L} &= \frac{1}{m-1} \sum_{k=1}^m (v^T(x_k - \mu))^2 - \lambda(\|v\|^2 - 1) \\ &= \frac{1}{m-1} \sum_{k=1}^m (v^T(x_k - \mu)(x_k - \mu)^T v) - \lambda(\|v\|^2 - 1) \\ &= v^T \left(\frac{1}{m-1} \sum_{k=1}^m (x_k - \mu)(x_k - \mu)^T \right) v - \lambda(\|v\|^2 - 1) \\ &= v^T C v - \lambda(v^T v - 1)\end{aligned}$$

- Extrema of the Lagrangian

$$\nabla \mathcal{L} = 2(Cv - \lambda v) = 0 \Rightarrow Cv = \lambda v$$

Adam Prügel-Bennett AICE3001/6002 Advanced Machine Learning <https://ecs-vlc.github.io/aice3001/> 7


- Covariance Matrices
- Principal Component Analysis
- Duality



Adam Prügel-Bennett AICE3001/6002 Advanced Machine Learning <https://ecs-vlc.github.io/aice3001/> 2

Looking is not Enough

Can't spot low dimensional data by looking at numbers




Adam Prügel-Bennett AICE3001/6002 Advanced Machine Learning <https://ecs-vlc.github.io/aice3001/> 4

Direction of Maximum Variation

- Look for the vector v with $\|v\|^2 = 1$ to maximise

$$\sigma^2 = \frac{1}{m-1} \sum_{i=1}^m (v^T(x_i - \mu))^2$$



- This is a constrained optimisation problem
- Solve by maximising Lagrangian


$$\mathcal{L} = \frac{1}{m-1} \sum_{k=1}^m (v^T(x_k - \mu))^2 - \lambda(\|v\|^2 - 1)$$

- λ is a Lagrange multiplier

Adam Prügel-Bennett AICE3001/6002 Advanced Machine Learning <https://ecs-vlc.github.io/aice3001/> 6

Direction of Maximum Variation

- The eigenvectors are directions that are extrema of the variance



- The variance in direction v is equal to

$$\begin{aligned}\sigma^2 &= \frac{1}{m-1} \sum_{i=1}^m (v^T(x_i - \mu))^2 \\ &= v^T C v = \lambda v^T v = \lambda\end{aligned}$$

- The variance is maximised by the eigenvector with the maximum eigenvalue

Adam Prügel-Bennett AICE3001/6002 Advanced Machine Learning <https://ecs-vlc.github.io/aice3001/> 8

Covariance Matrix

- The **covariance matrix** is defined as

$$\mathbf{C} = \frac{1}{m-1} \sum_{k=1}^m (\mathbf{x}_k - \boldsymbol{\mu})(\mathbf{x}_k - \boldsymbol{\mu})^\top$$

- The components C_{ij} measure how the i^{th} and j^{th} components co-vary

$$C_{ij} = \frac{1}{m-1} \sum_{k=1}^m (x_{ik} - \mu_i)(x_{jk} - \mu_j)$$

- C.f. covariance of random variables

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$$

Matrix Form

- The covariance matrix is

$$\mathbf{C} = \frac{1}{m-1} \sum_{k=1}^m (\mathbf{x}_k - \boldsymbol{\mu})(\mathbf{x}_k - \boldsymbol{\mu})^\top$$

- Define the matrix

$$\mathbf{X} = \frac{1}{\sqrt{m-1}} (\mathbf{x}_1 - \boldsymbol{\mu}, \mathbf{x}_2 - \boldsymbol{\mu}, \dots, \mathbf{x}_m - \boldsymbol{\mu})$$

- We can write the covariance matrix as

$$\mathbf{C} = \mathbf{X}\mathbf{X}^\top$$

Eigenvalue Decomposition

- The eigenvectors of $\mathbf{C}$ with the largest eigenvalues are known as the **principal components**
- The eigenvalues are all greater than or equal to zero
- Recall an eigenvector $\mathbf{v}$ satisfies the equation

$$\mathbf{C}\mathbf{v} = \lambda\mathbf{v}$$


- Multiplying both sides by $\mathbf{v}^\top$

$$\mathbf{v}^\top \mathbf{C} \mathbf{v} = \lambda \mathbf{v}^\top \mathbf{v} = \lambda \|\mathbf{v}\|^2$$

but $\mathbf{v}^\top \mathbf{C} \mathbf{v} \geq 0$ and $\|\mathbf{v}\|^2 > 0$ so

$$\lambda = \frac{\mathbf{v}^\top \mathbf{C} \mathbf{v}}{\|\mathbf{v}\|^2} \geq 0$$

Ellipsoid and Eigen Space



Outer Product

- Remember that the outer-product of two vectors is defined as

$$\mathbf{x}\mathbf{y}^\top = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} \begin{pmatrix} y_1 & y_2 & \dots & y_n \end{pmatrix} = \begin{pmatrix} x_1 y_1 & x_1 y_2 & \dots & x_1 y_n \\ x_2 y_1 & x_2 y_2 & \dots & x_2 y_n \\ \vdots & \vdots & \ddots & \vdots \\ x_n y_1 & x_n y_2 & \dots & x_n y_n \end{pmatrix}$$

- C.f. Inner product

$$\mathbf{x}^\top \mathbf{y} = \begin{pmatrix} x_1 & x_2 & \dots & x_n \end{pmatrix} \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = x_1 y_1 + x_2 y_2 + \dots + x_n y_n$$

Properties of Covariance Matrix

- The **quadratic form** of a vector and matrix is defined as

$$\mathbf{v}^\top \mathbf{M} \mathbf{v}$$

- The quadratic form of a covariance matrix is non-negative for any vector

$$\mathbf{v}^\top \mathbf{C} \mathbf{v} = \mathbf{v}^\top \mathbf{X} \mathbf{X}^\top \mathbf{v} = \mathbf{u}^\top \mathbf{u} = \|\mathbf{u}\|^2 \geq 0$$

where $\mathbf{u} = \mathbf{X}^\top \mathbf{v}$

- Matrices with non-negative quadratic forms are known as **positive semi-definite**

Surface Defined by Matrix

- The set of vectors $\mathbf{x}$ such that


$$\mathbf{x}^\top \mathbf{C}^{-1} \mathbf{x} = 1$$

defines a surface

- The surface is an ellipsoid, $\mathcal{E}$
- The eigenvectors point in the direction of the principal axes of the ellipsoid
- The radii of the principal axes are equal to the square root of the eigenvalues

Spanning Input Space

- A covariance matrix will have a zero eigenvalue only if there is no variation in the direction of the corresponding eigenvector
- A covariance matrix will have zero eigenvalues if the number of patterns are less than or equal to the number of dimensions
- A covariance matrix formed from $p+1$ patterns that are linearly independent (i.e. you cannot form any one out of p of the other patterns) will have no zero eigenvalues




- Matrices with no zero eigenvalues are called **full rank** matrices (as opposed to rank deficient)
- Full rank matrices are invertible, rank deficient matrices are singular and non-invertible
- Full rank covariance matrices that have positive eigenvalues only and are said to be **positive definite**
- We would expect that when $m > p$ the covariance matrix will be positive definite unless there are some symmetries that linearly constrain the patterns


Principal Component Analysis

- PCA occurs as follows
 - Construct the covariance matrix
 - Find the eigenvalues and eigenvectors
 - Keep the eigenvectors with the largest eigenvalues (principal components)
 - Project the inputs into the space spanned by the principal components
- We then use the projected inputs as inputs to our learning machine


Subspace Projection



Eigenvectors



- Covariance Matrices
- Principal Component Analysis**
- Duality



Projection Matrix

- To project the inputs construct the projection matrix


$$P = \begin{pmatrix} v_1^T \\ v_2^T \\ \vdots \\ v_k^T \end{pmatrix}$$

- $k < p$ is the number of principal components we keep
- Given a p -dimensional input pattern x we can construct a k -dimensional representation z

$$z = P(x - \mu)$$

- Use z as our new inputs

Hand Written Digits




Reconstruction

- Projecting into a subspace of eigenvectors can be seen as approximating the inputs by


$$\hat{x}_i = \mu + \sum_{j=1}^k z_j^i v_j, \quad z_j^i = v_j^T (x_i - \mu), \quad \|v_j\| = 1$$

- Principle component analysis projects the data into a subspace of size m with the minimal approximation error $\mathbb{E}[\| \hat{x}_i - x_i \|^2]$
- The loss of “energy” (or squared error) is equal to the sum of the eigenvalues in the directions that are ignored

Eigenvalues for Digits




Reconstruction from Eigenvectors



Outline

1. Covariance Matrices
2. Principal Component Analysis
3. Duality



PCA for Images

- An image often contains around $p = 256 \times 256 = 64k$ pixels
- In standard PCA we would create an $p \times p$ matrix with over 4×10^9 elements
- This is intractable
- m images span at most a $m - 1$ dimensional subspace
- Usually this subspace will be much smaller than the space of all images $m \ll p$

Dual Matrix

- The covariance $C = XX^T$ is a $p \times p$ matrix
- Consider the $m \times m$ matrix $D = X^T X$
- Suppose v is an eigenvector of D


$$\begin{aligned} Dv &= \lambda v \\ X^T X v &= \lambda v \\ XX^T X v &= \lambda X v \\ CX v &= \lambda X v \Rightarrow Cu = \lambda u \end{aligned}$$

- $u = Xv$ (and $v \propto X^T u$)

What Does a Subspace Look Like?

- Consider $y^1 = \begin{pmatrix} 2 \\ 4 \end{pmatrix}$, $y^2 = \begin{pmatrix} 8 \\ 6 \end{pmatrix}$ with mean $\mu = \begin{pmatrix} 5 \\ 3 \end{pmatrix}$
- Subtracting the mean $x^i = y^i - \mu$ we can construct matrix

$$X = \begin{pmatrix} x_1^1 & x_1^2 \\ x_2^1 & x_2^2 \\ x_3^1 & x_3^2 \end{pmatrix} = \begin{pmatrix} -3 & 3 \\ -1 & 1 \\ 2 & -2 \end{pmatrix}$$



Dual Matrix

- Matrices $C = XX^T$ and $D = X^T X$ have the same eigenvalues
- Can use the dual $m \times m$ matrix D to find eigenvalues and eigenvectors of C
- Note that $D = X^T X$ has components $D_{kl} \propto (x_k - \mu)^T (x_l - \mu)$
- Takes $O(p \times m \times m)$ time to construct D
- We work in a “dual space” which is the space spanned by the examples

Summary

- PCA allows us to reduce the dimensionality of the inputs
- We project the inputs into a sub-space where the data varies the most
- We can work in either the original space (XX^T) or the dual space ($X^T X$)
- When we have many more features than examples (i.e. $p \gg m$) then it is more efficient working in the dual space
- We will see examples of dual spaces again when we look at SVMs