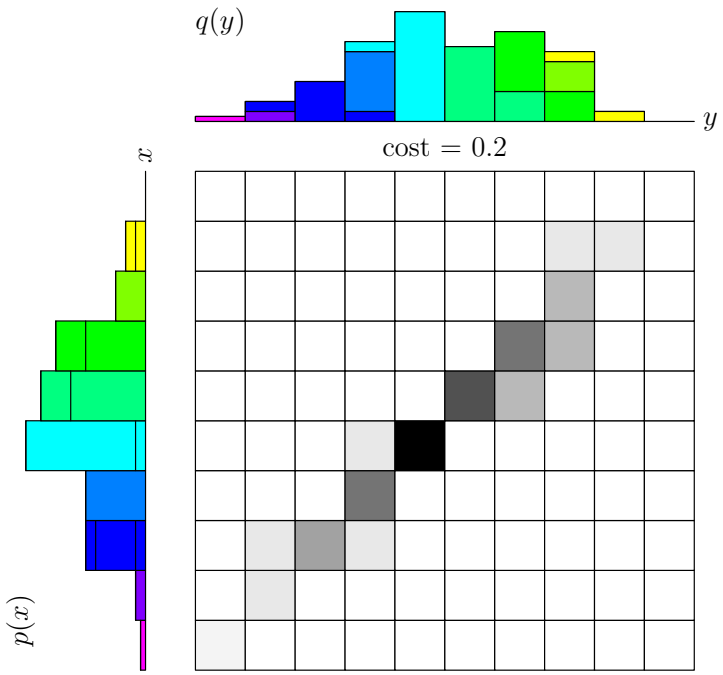
Mathematics of Machine Learning

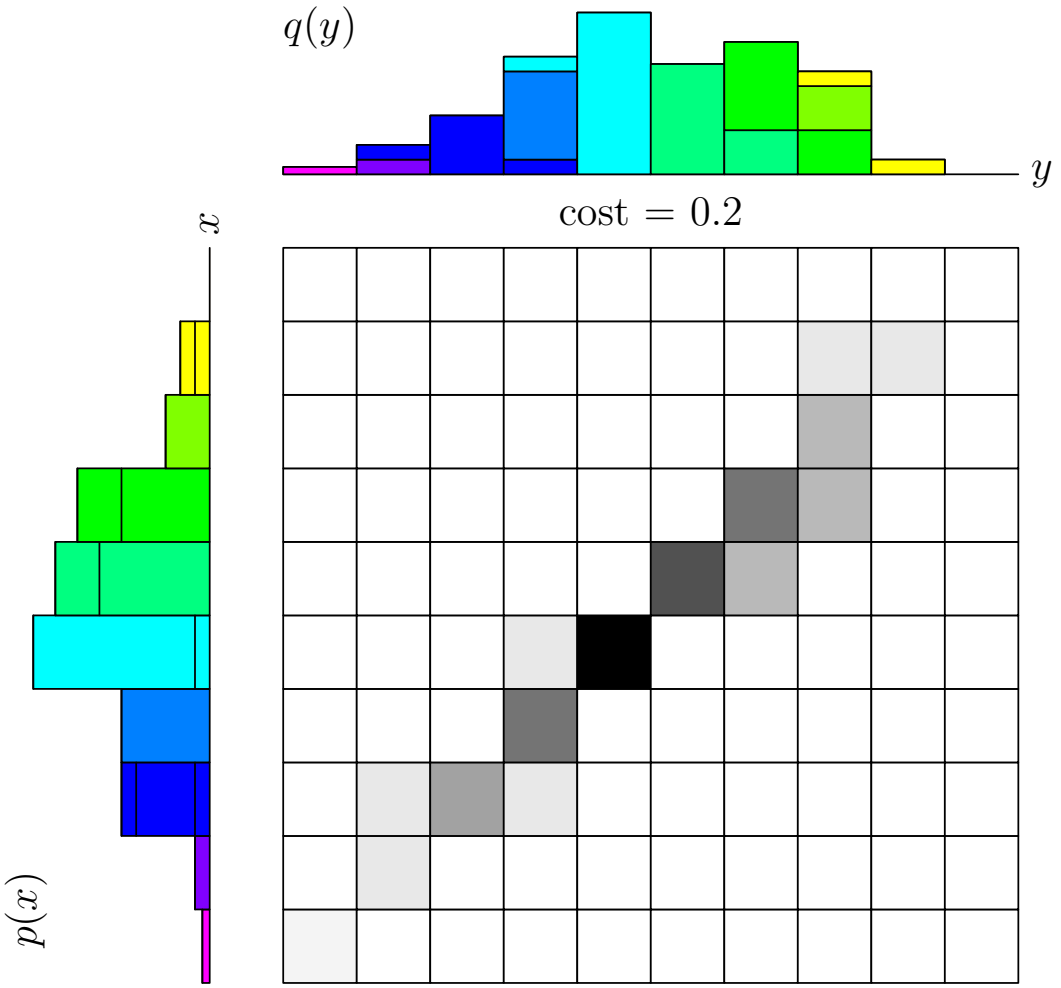
Optimal Transport and the Wasserstein Distance



Wasserstein distance, optimal transport, duality, Lipschitz functions, WGANs, Sinkhorn, FID

Outline

1. **Comparing Distributions**
2. Optimal Transport
3. Duality
4. Applications



Measuring Distances Between Distributions

- In many machine learning tasks we want to minimise the distance between two probability distributions (e.g. make a generative model's samples look like the data)
- This requires that we can measure distances between probability distributions
- One prominent measure is the Kullback-Leibler or KL divergence

$$\mathrm{KL}(p\|q) = \int p(\boldsymbol{x}) \log\left(\frac{p(\boldsymbol{x})}{q(\boldsymbol{x})}\right) \mathrm{d}\boldsymbol{x}$$

- This is very commonly used in ML (e.g. VAEs, Variational Approximation)

Measuring Distances Between Distributions

- In many machine learning tasks we want to minimise the distance between two probability distributions (e.g. make a generative model's samples look like the data)
- This requires that we can measure distances between probability distributions
- One prominent measure is the Kullback-Leibler or KL divergence

$$\mathrm{KL}(p\|q) = \int p(\boldsymbol{x}) \log\left(\frac{p(\boldsymbol{x})}{q(\boldsymbol{x})}\right) \mathrm{d}\boldsymbol{x}$$

- This is very commonly used in ML (e.g. VAEs, Variational Approximation)

Measuring Distances Between Distributions

- In many machine learning tasks we want to minimise the distance between two probability distributions (e.g. make a generative model's samples look like the data)
- This requires that we can measure distances between probability distributions
- One prominent measure is the Kullback-Leibler or KL divergence

$$\mathrm{KL}(p\|q) = \int p(\boldsymbol{x}) \log\left(\frac{p(\boldsymbol{x})}{q(\boldsymbol{x})}\right) \mathrm{d}\boldsymbol{x}$$

- This is very commonly used in ML (e.g. VAEs, Variational Approximation)

Measuring Distances Between Distributions

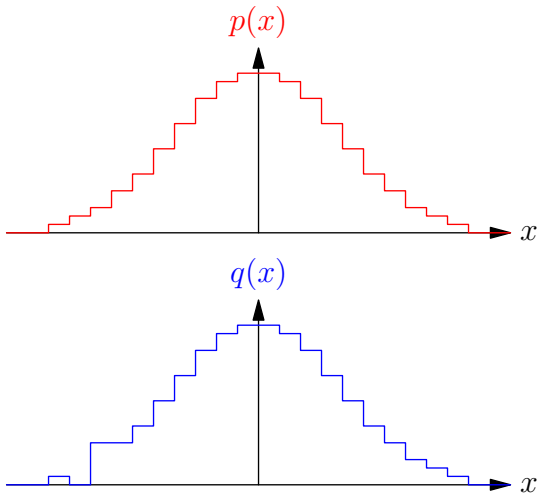
- In many machine learning tasks we want to minimise the distance between two probability distributions (e.g. make a generative model's samples look like the data)
- This requires that we can measure distances between probability distributions
- One prominent measure is the Kullback-Leibler or KL divergence

$$\mathrm{KL}(p\|q) = \int p(\boldsymbol{x}) \log\left(\frac{p(\boldsymbol{x})}{q(\boldsymbol{x})}\right) \mathrm{d}\boldsymbol{x}$$

- This is very commonly used in ML (e.g. VAEs, Variational Approximation)

Trouble with KL

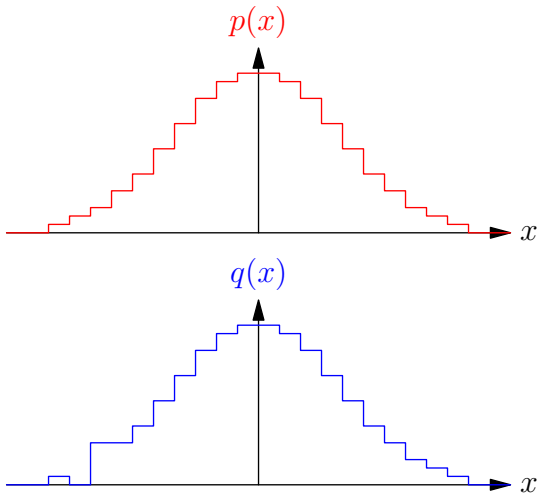
- KL-divergences are non-negative quantities that are minimised when the two probability distributions are the same
- They are not distances (they aren't symmetric and they don't satisfy the triangular inequality)
- We don't really care about this, but what we do care about is that if $q(\boldsymbol{x}) = 0$ when $p(\boldsymbol{x}) \neq 0$ then $\log\left(\frac{p(\boldsymbol{x})}{q(\boldsymbol{x})}\right)$ diverges



- We can therefore have distributions that seem very similar but their KL-divergence is huge (or infinite)

Trouble with KL

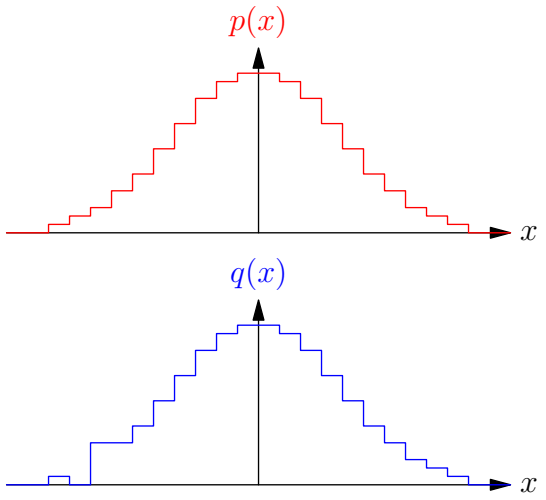
- KL-divergences are non-negative quantities that are minimised when the two probability distributions are the same
- They are not distances (they aren't symmetric and they don't satisfy the triangular inequality)
- We don't really care about this, but what we do care about is that if $q(\boldsymbol{x}) = 0$ when $p(\boldsymbol{x}) \neq 0$ then $\log\left(\frac{p(\boldsymbol{x})}{q(\boldsymbol{x})}\right)$ diverges



- We can therefore have distributions that seem very similar but their KL-divergence is huge (or infinite)

Trouble with KL

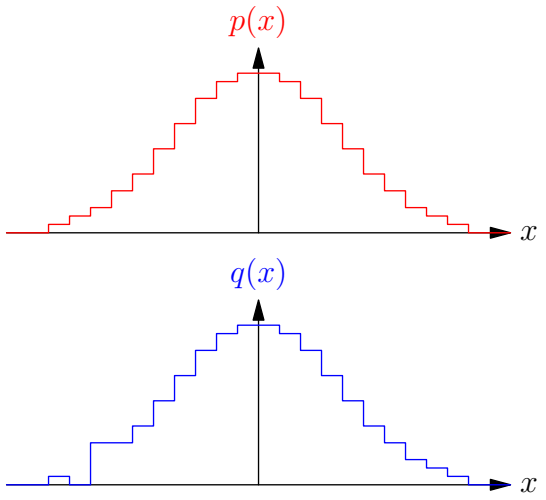
- KL-divergences are non-negative quantities that are minimised when the two probability distributions are the same
- They are not distances (they aren't symmetric and they don't satisfy the triangular inequality)
- We don't really care about this, but what we do care about is that if $q(\boldsymbol{x}) = 0$ when $p(\boldsymbol{x}) \neq 0$ then $\log\left(\frac{p(\boldsymbol{x})}{q(\boldsymbol{x})}\right)$ diverges



- We can therefore have distributions that seem very similar but their KL-divergence is huge (or infinite)

Trouble with KL

- KL-divergences are non-negative quantities that are minimised when the two probability distributions are the same
- They are not distances (they aren't symmetric and they don't satisfy the triangular inequality)
- We don't really care about this, but what we do care about is that if $q(\boldsymbol{x}) = 0$ when $p(\boldsymbol{x}) \neq 0$ then $\log\left(\frac{p(\boldsymbol{x})}{q(\boldsymbol{x})}\right)$ diverges

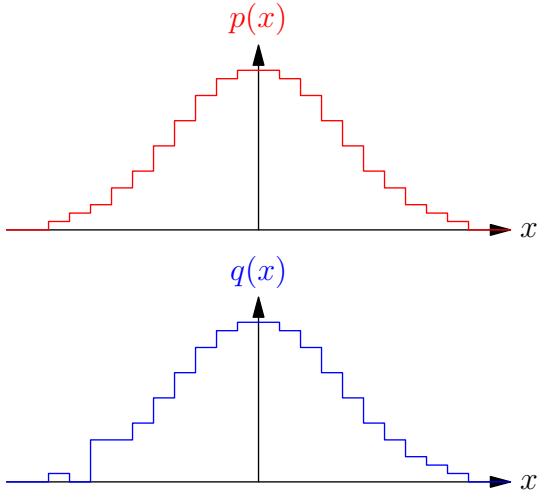


- We can therefore have distributions that seem very similar but their KL-divergence is huge (or infinite)

Wasserstein Distance

- A more benign measure of the differences between two probability distributions is the **Wasserstein** or **Earth Mover's** distance

- This is a true distance, but more importantly, it measures distance in a very natural way so that distributions that are close have a small Wasserstein distance

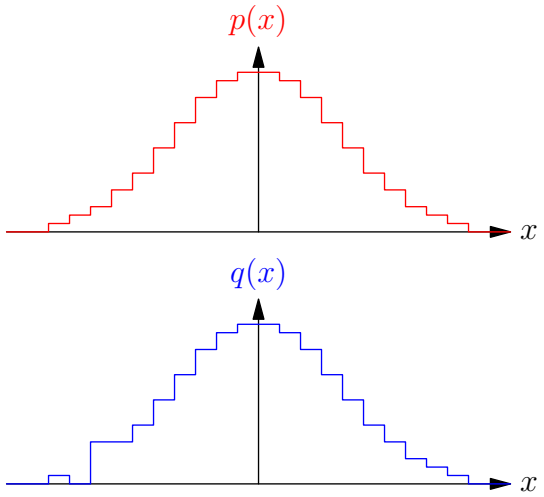


- This may seem contrived, but images (e.g. $128 \times 128 \times 3$ arrays of pixel values) lie close to a thin manifold, so two distributions of images can be close without overlapping: then the KL-divergence is useless but the Wasserstein distance is not

Wasserstein Distance

- A more benign measure of the differences between two probability distributions is the **Wasserstein** or **Earth Mover's** distance

- This is a true distance, but more importantly, it measures distance in a very natural way so that distributions that are close have a small Wasserstein distance

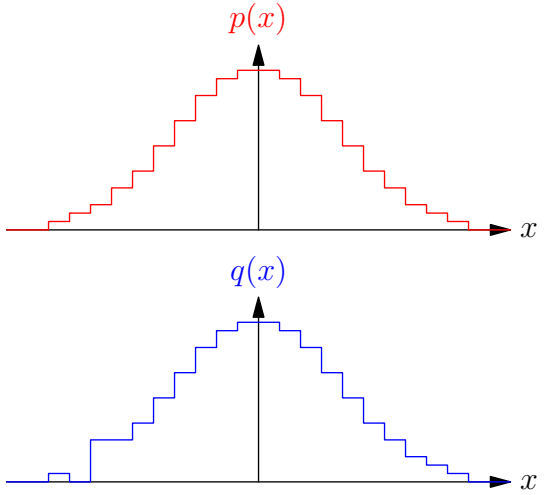


- This may seem contrived, but images (e.g. $128 \times 128 \times 3$ arrays of pixel values) lie close to a thin manifold, so two distributions of images can be close without overlapping: then the KL-divergence is useless but the Wasserstein distance is not

Wasserstein Distance

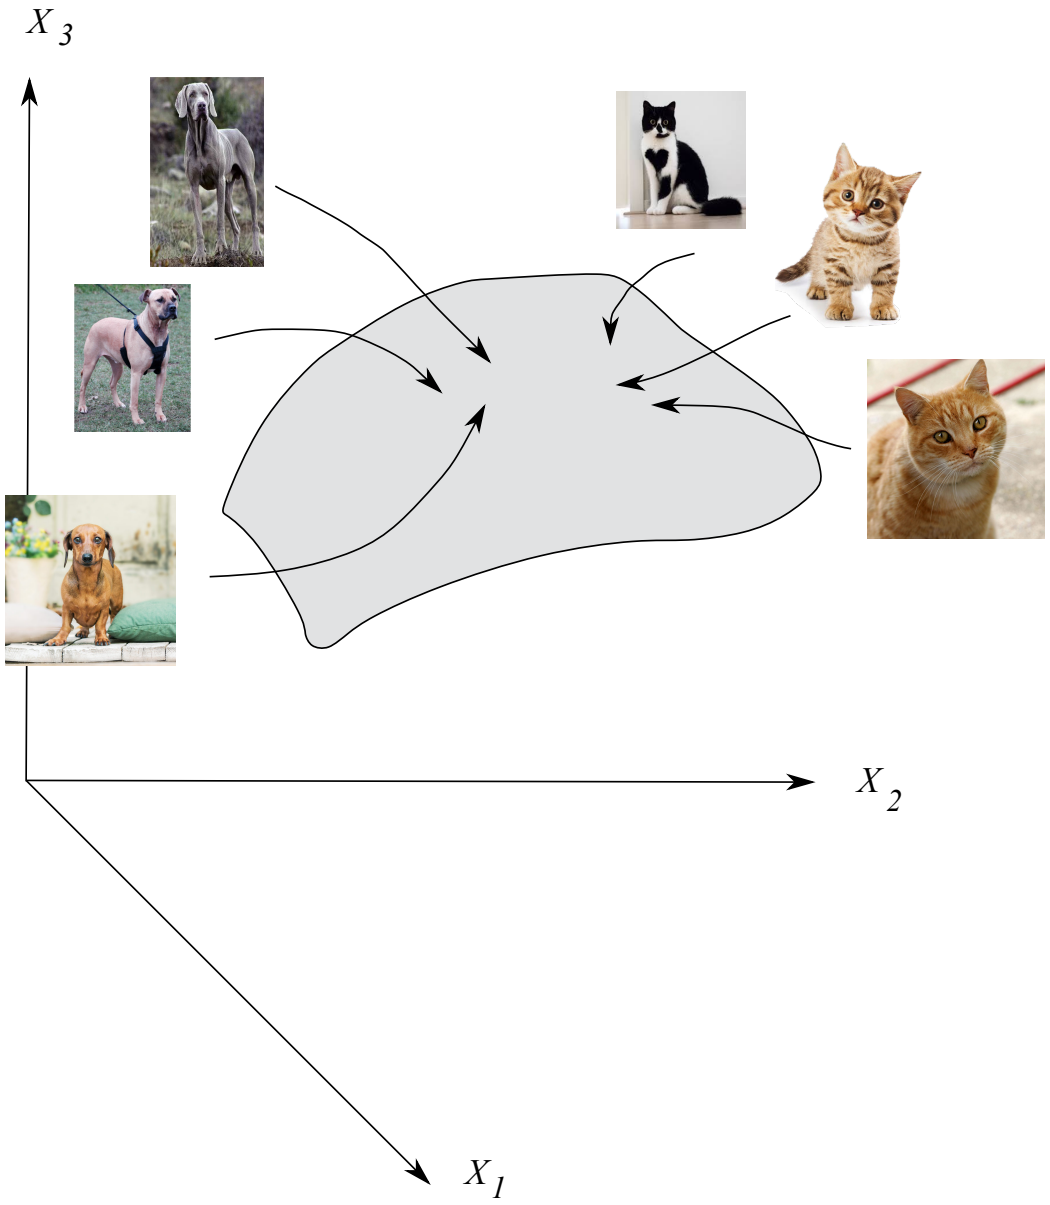
- A more benign measure of the differences between two probability distributions is the **Wasserstein** or **Earth Mover's** distance

- This is a true distance, but more importantly, it measures distance in a very natural way so that distributions that are close have a small Wasserstein distance



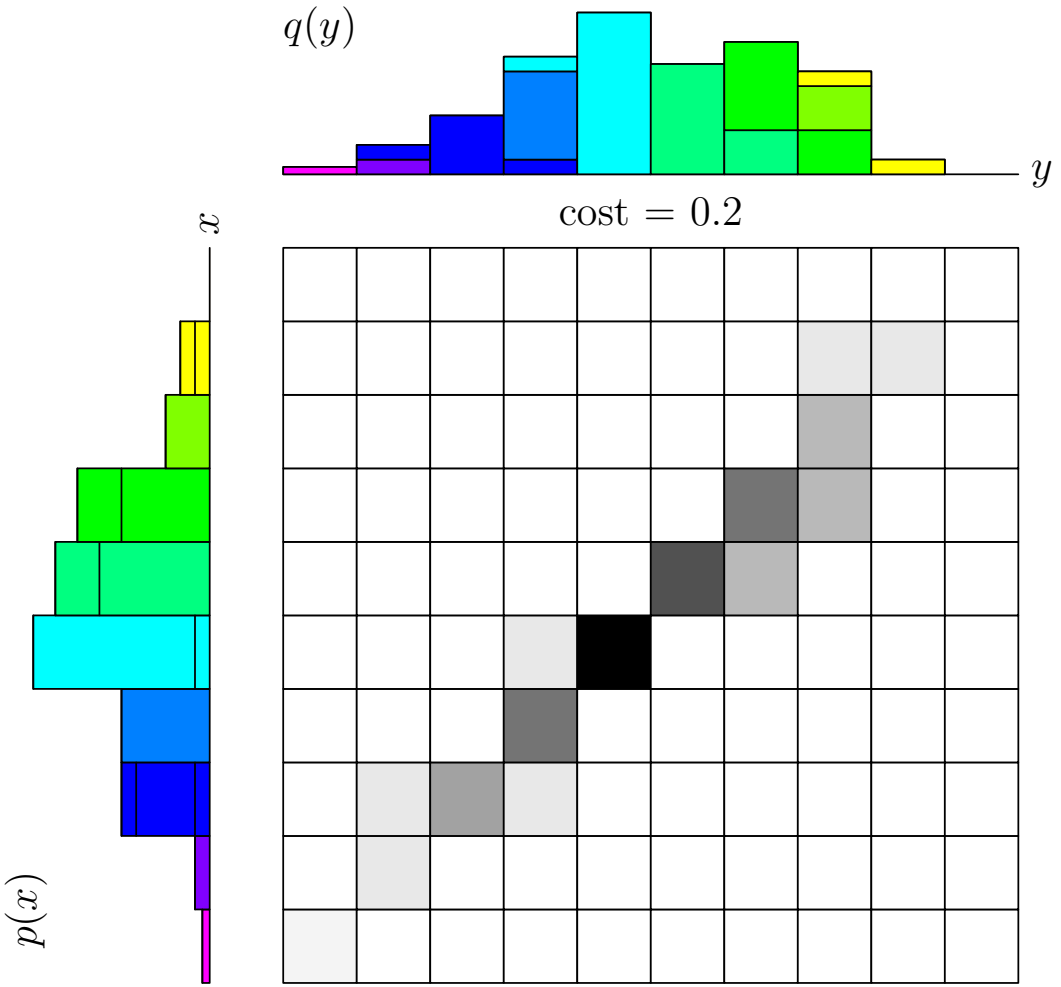
- This may seem contrived, but images (e.g. $128 \times 128 \times 3$ arrays of pixel values) lie close to a thin manifold, so two distributions of images can be close without overlapping: then the KL-divergence is useless but the Wasserstein distance is not

High Probability Manifold



Outline

1. Comparing Distributions
2. **Optimal Transport**
3. Duality
4. Applications



Transportation Policy

- But how do we formalise the Wasserstein distance?
- A good place to start is to define a transportation policy $\gamma(\boldsymbol{x},\boldsymbol{y})$ with

$$\int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{y} = p(\boldsymbol{x}) \qquad \int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x} = q(\boldsymbol{y})$$

- This looks like a joint probability distribution, but we interpret $\gamma(\boldsymbol{x},\boldsymbol{y})$ as the amount of probability mass/density that we transfer from $p(\boldsymbol{x})$ to $q(\boldsymbol{y})$

Transportation Policy

- But how do we formalise the Wasserstein distance?
- A good place to start is to define a transportation policy $\gamma(\boldsymbol{x},\boldsymbol{y})$ with

$$\int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{y} = p(\boldsymbol{x}) \qquad \int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x} = q(\boldsymbol{y})$$

- This looks like a joint probability distribution, but we interpret $\gamma(\boldsymbol{x},\boldsymbol{y})$ as the amount of probability mass/density that we transfer from $p(\boldsymbol{x})$ to $q(\boldsymbol{y})$

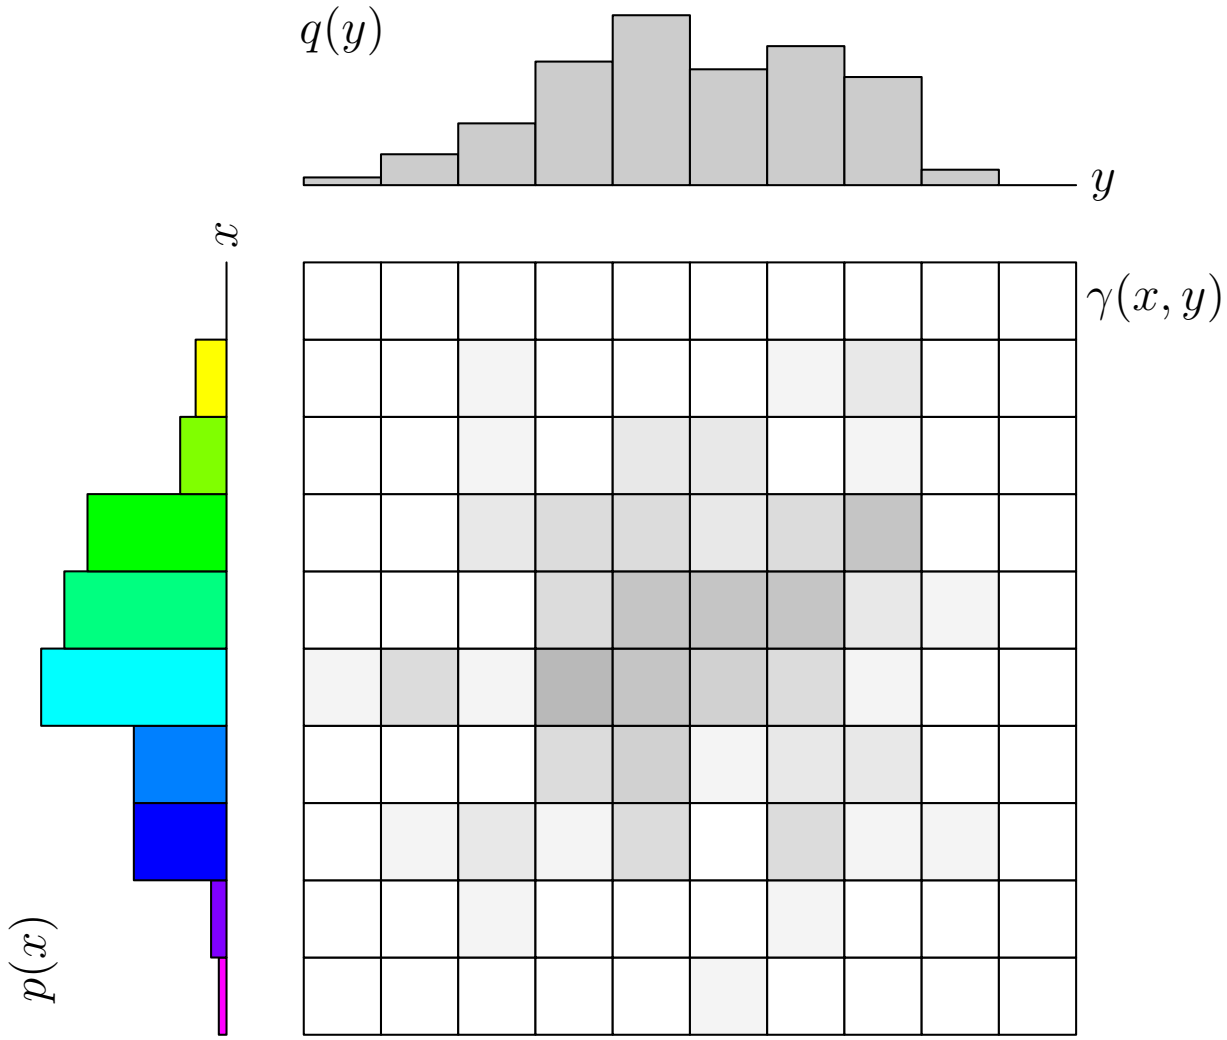
Transportation Policy

- But how do we formalise the Wasserstein distance?
- A good place to start is to define a transportation policy $\gamma(\boldsymbol{x},\boldsymbol{y})$ with

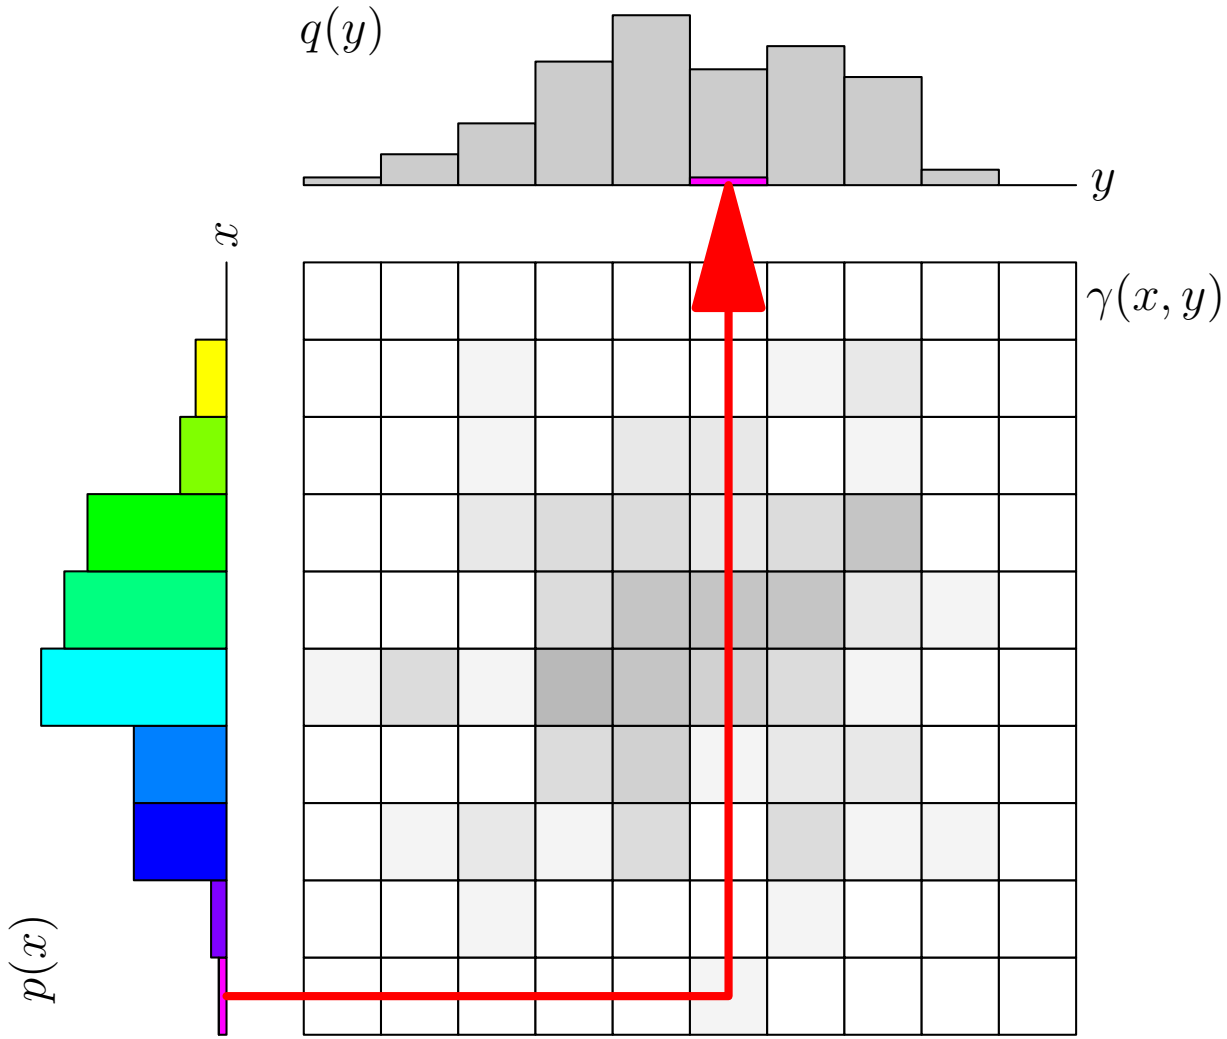
$$\int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{y} = p(\boldsymbol{x}) \qquad \int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x} = q(\boldsymbol{y})$$

- This looks like a joint probability distribution, but we interpret $\gamma(\boldsymbol{x},\boldsymbol{y})$ as the amount of probability mass/density that we transfer from $p(\boldsymbol{x})$ to $q(\boldsymbol{y})$

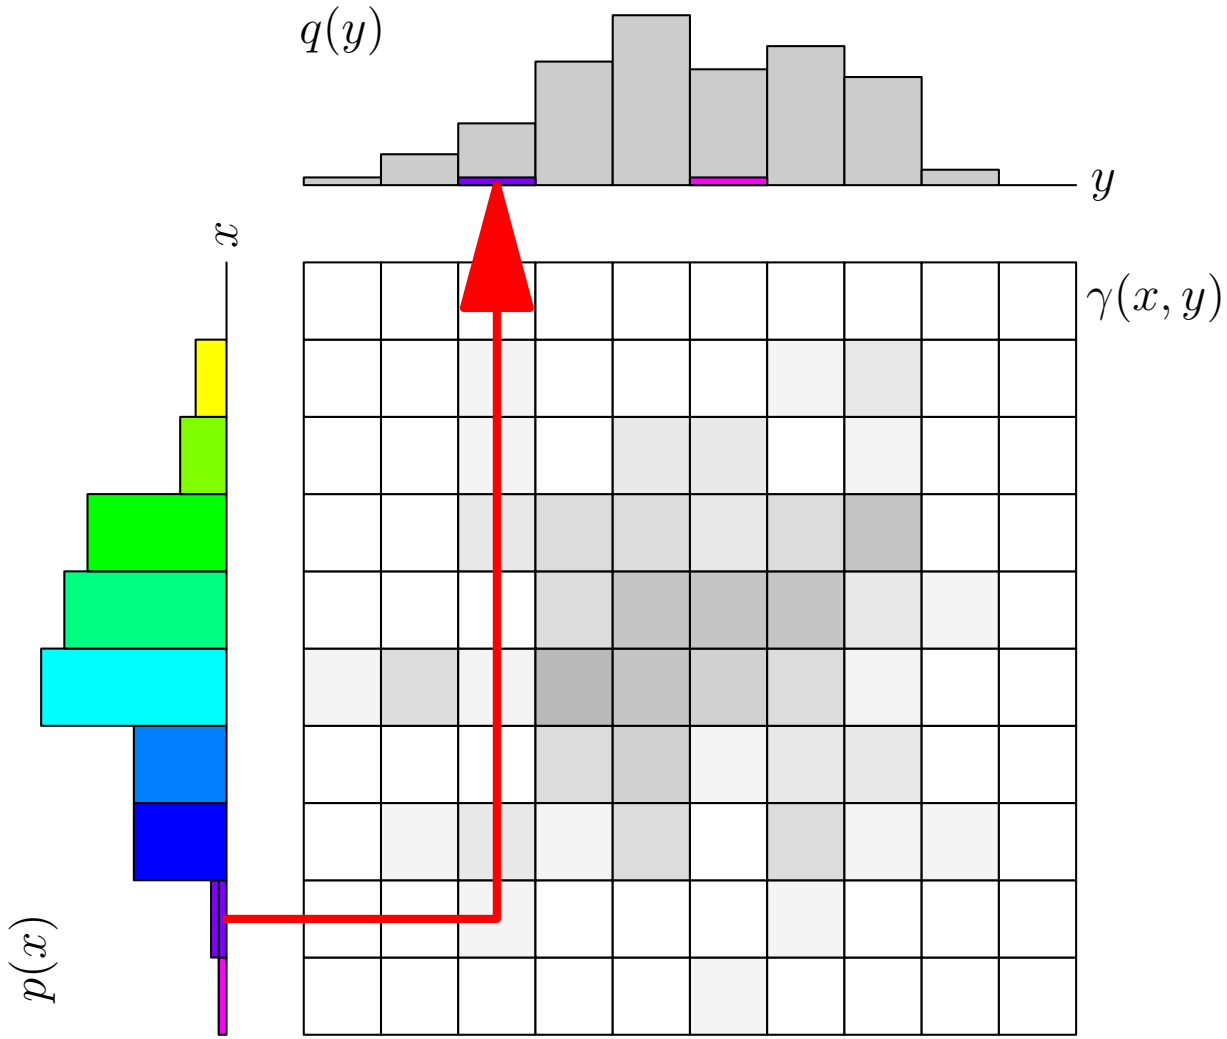
Transportation Policy



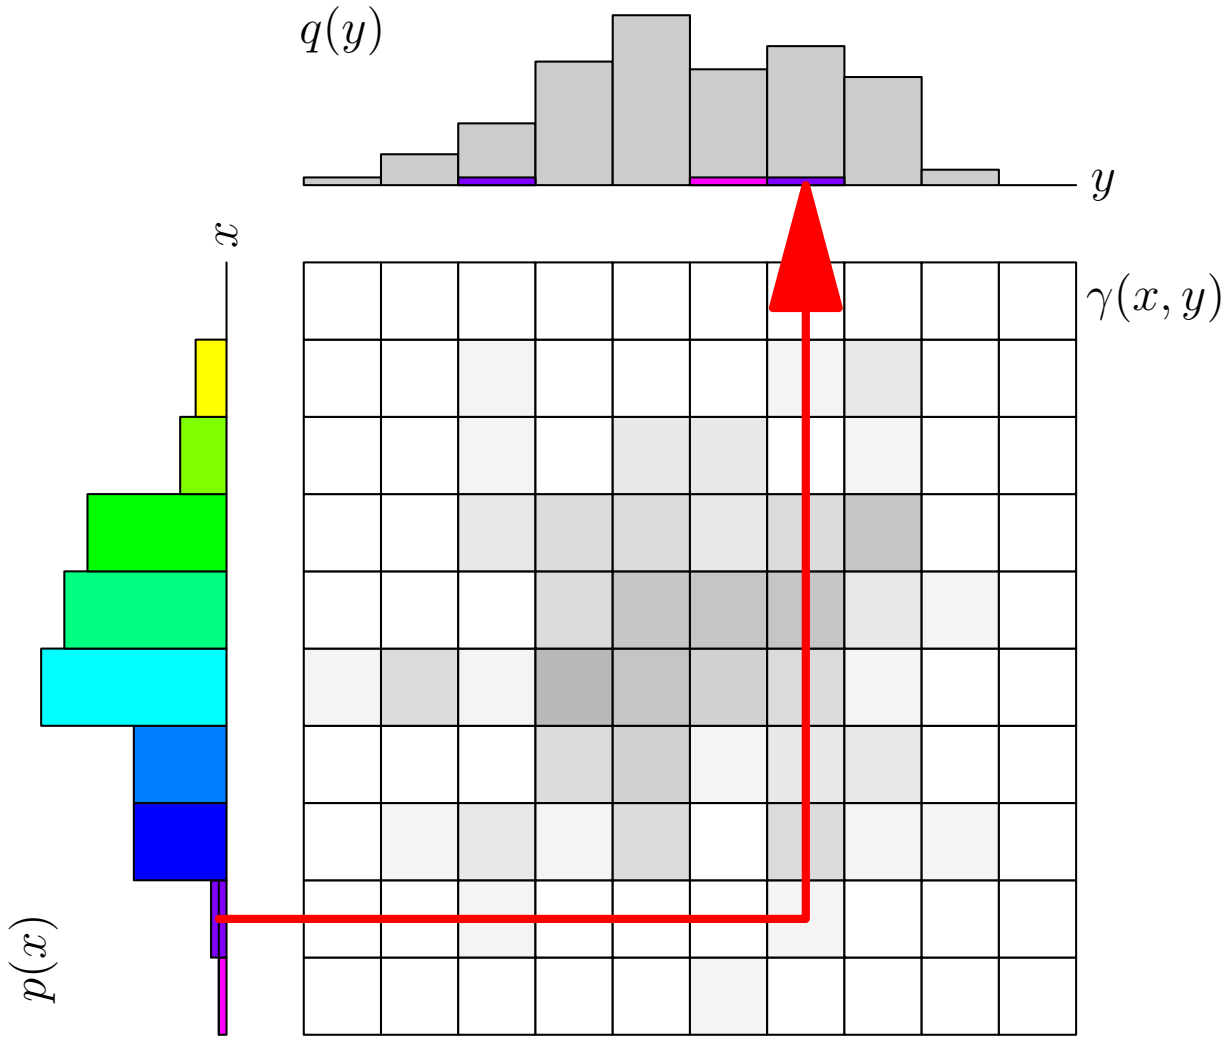
Transportation Policy



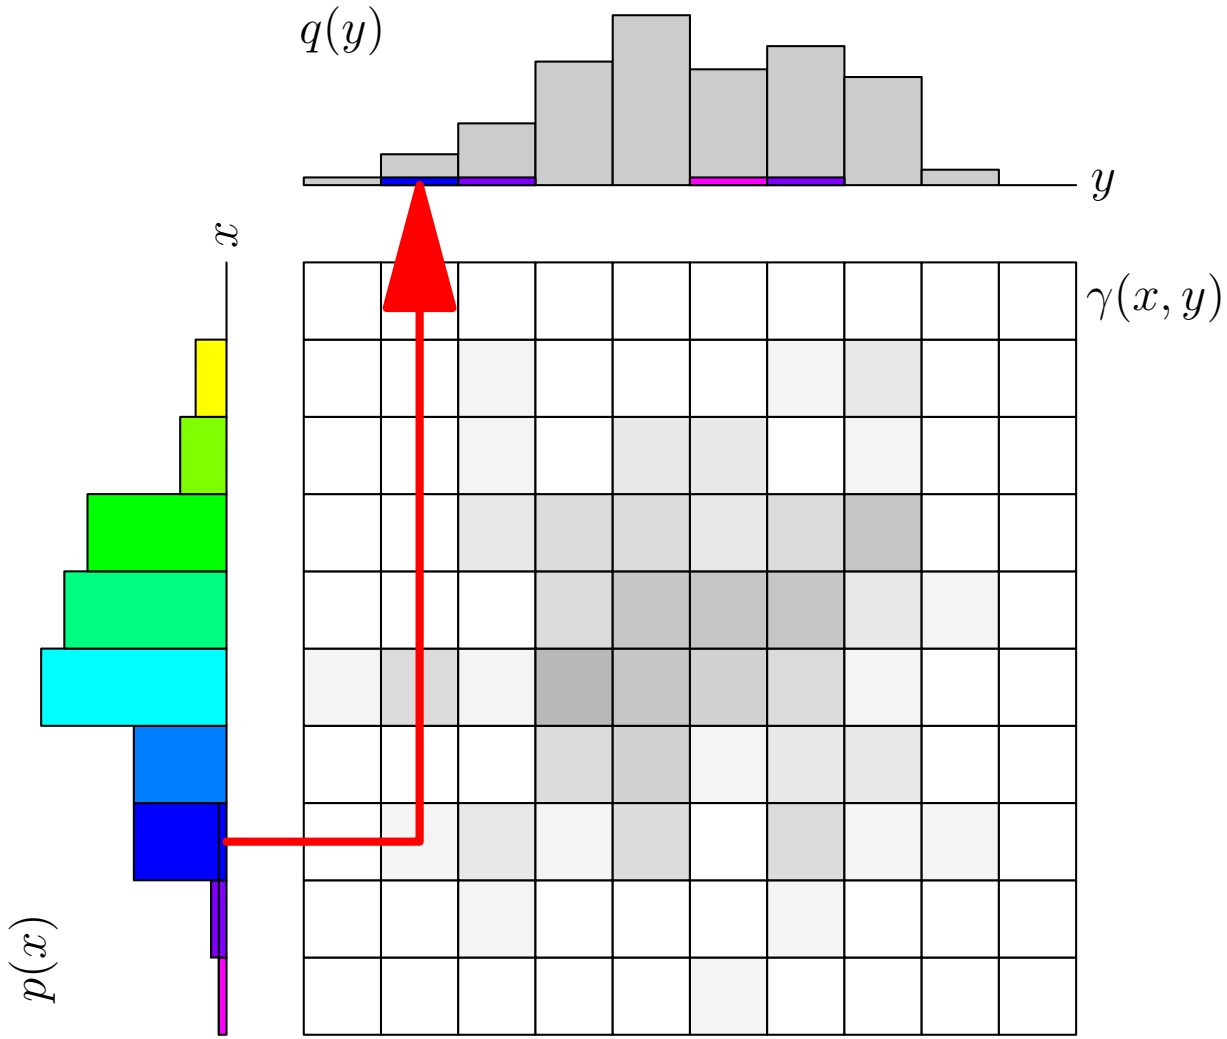
Transportation Policy



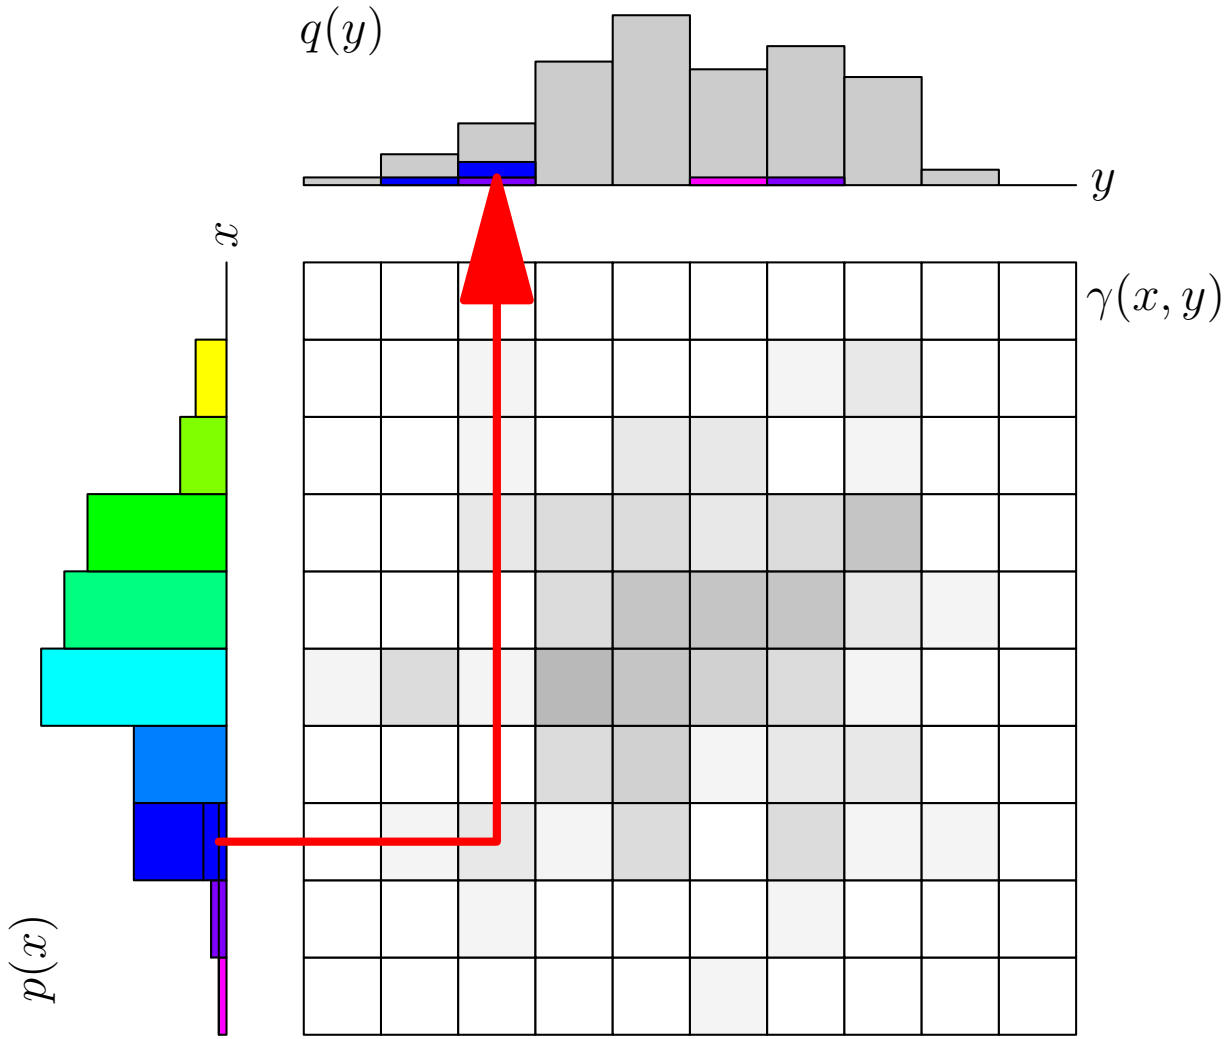
Transportation Policy



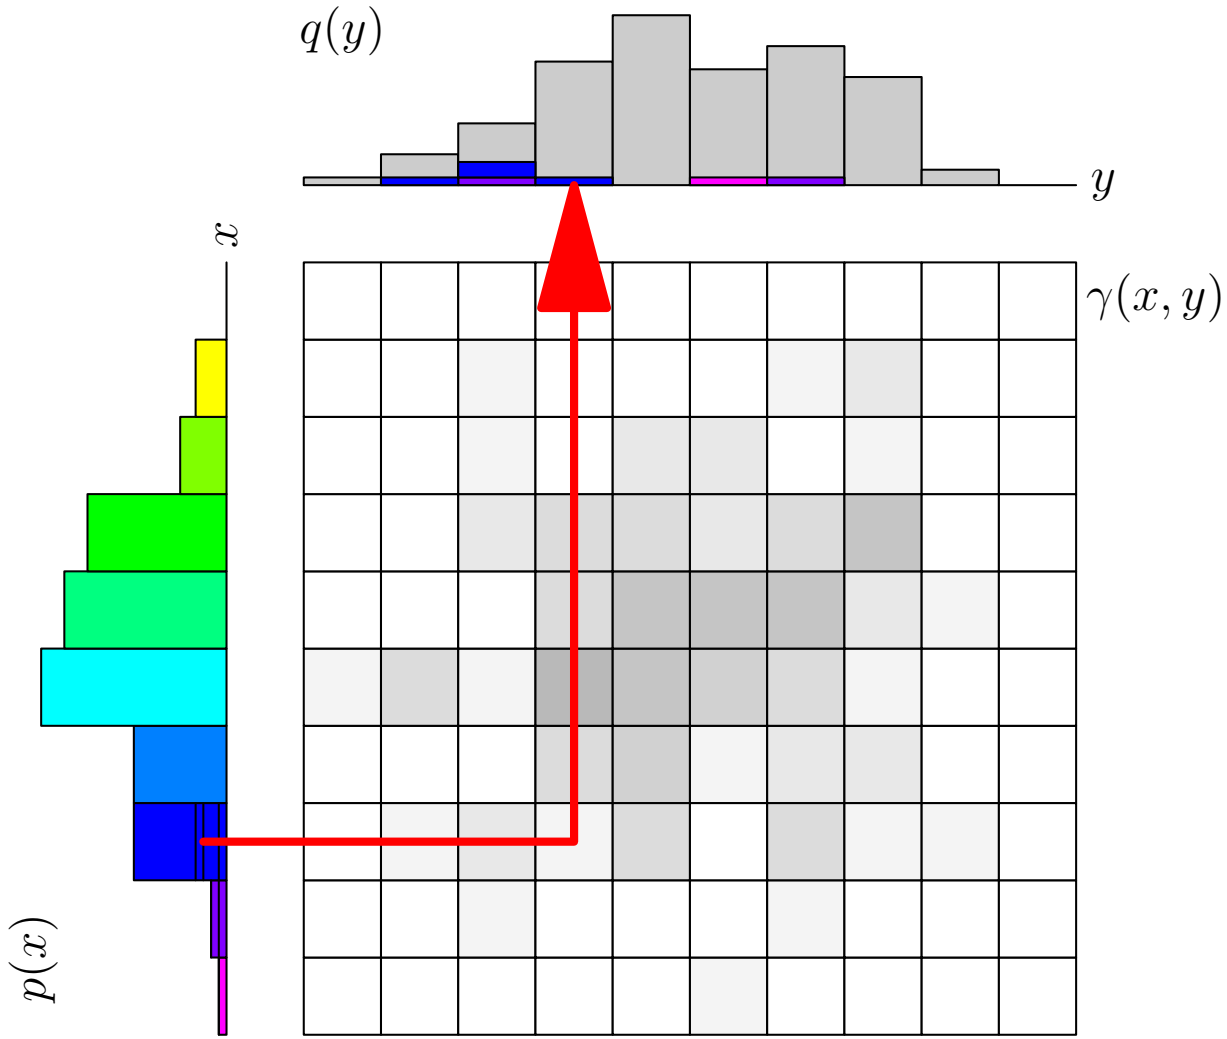
Transportation Policy



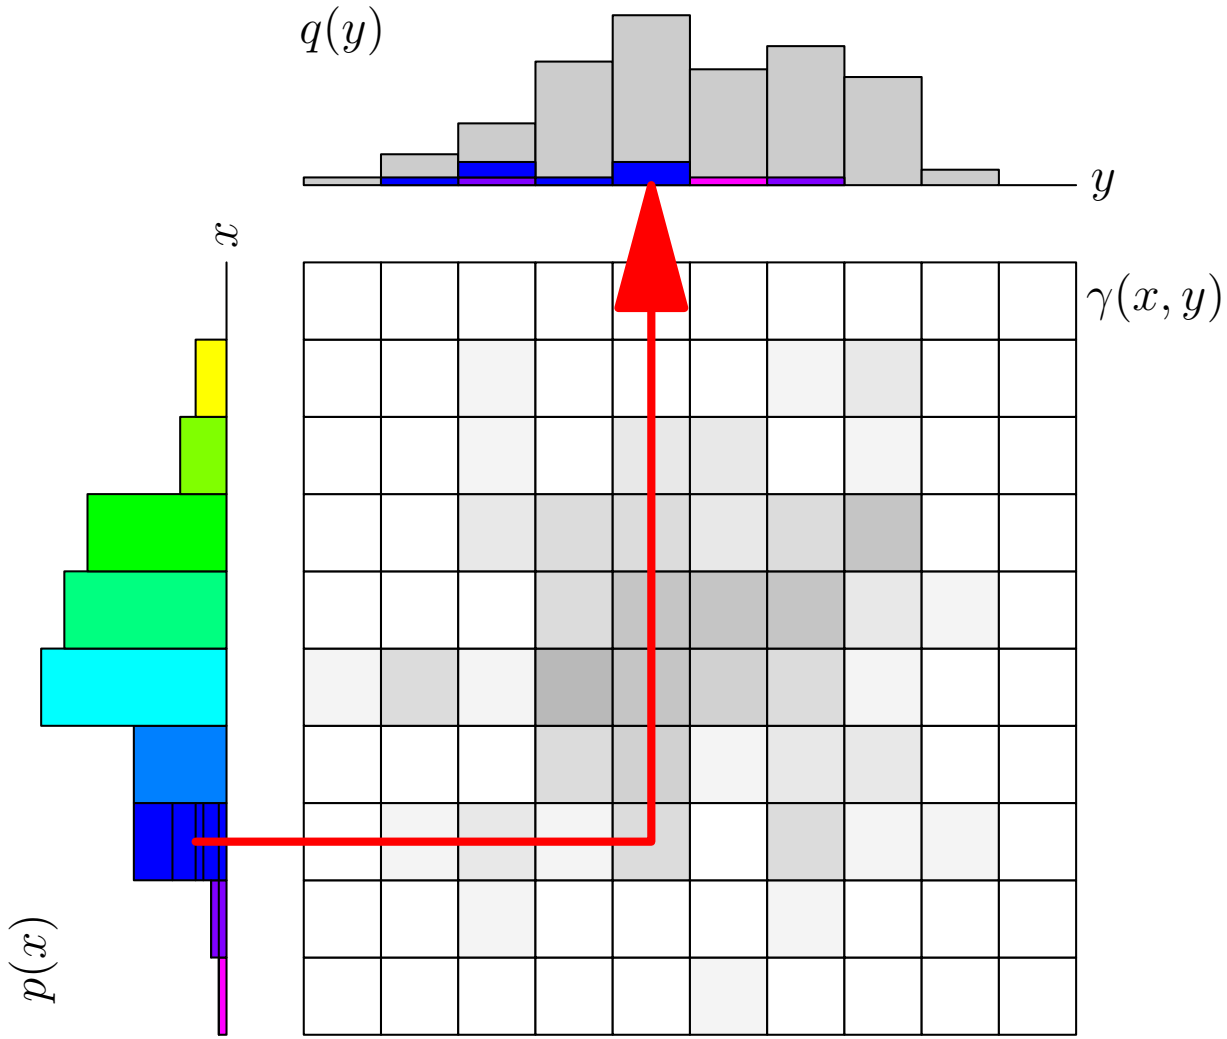
Transportation Policy



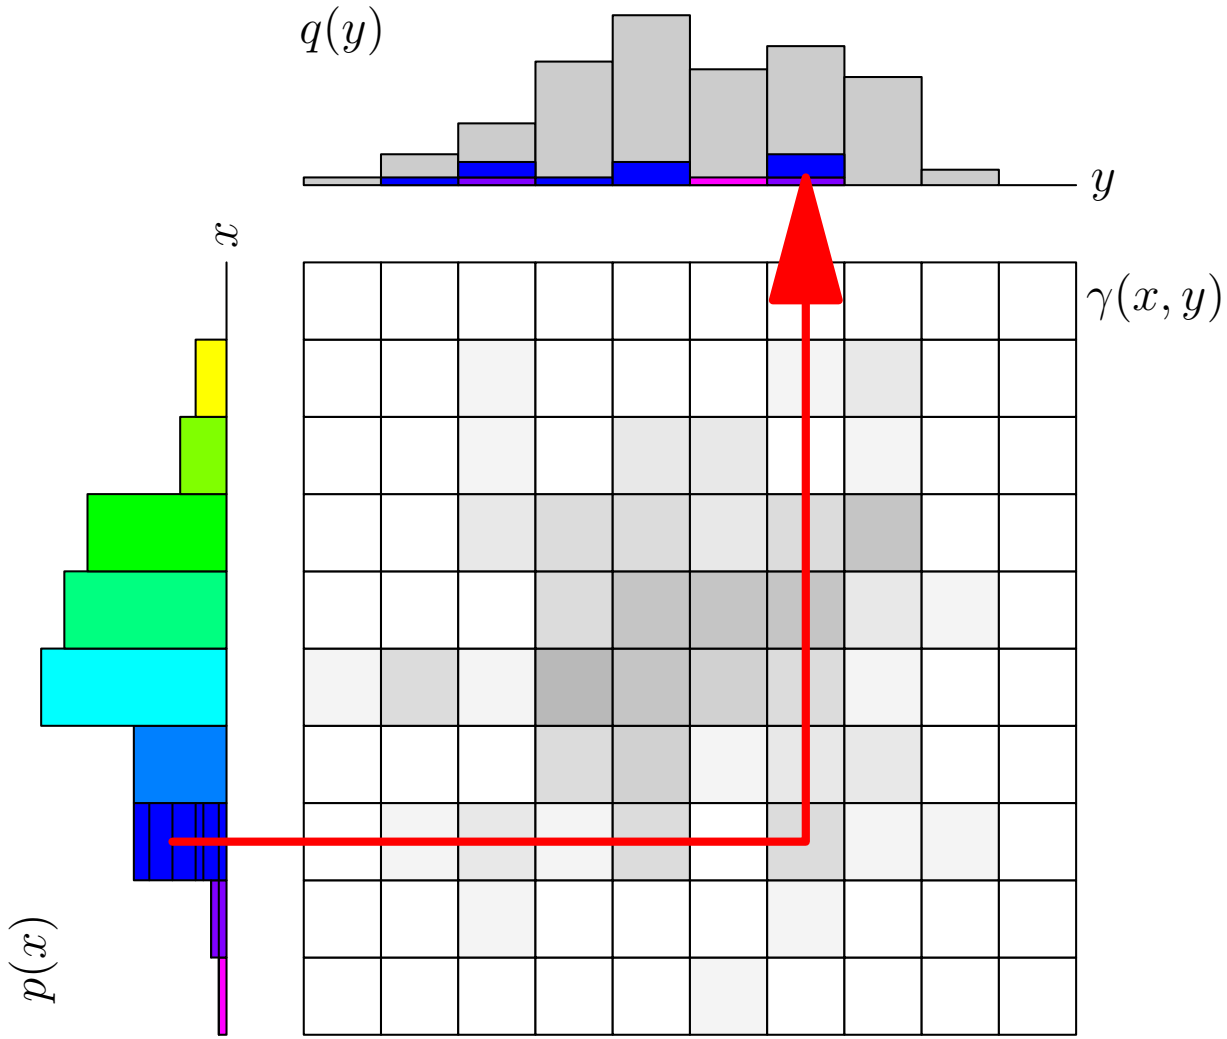
Transportation Policy



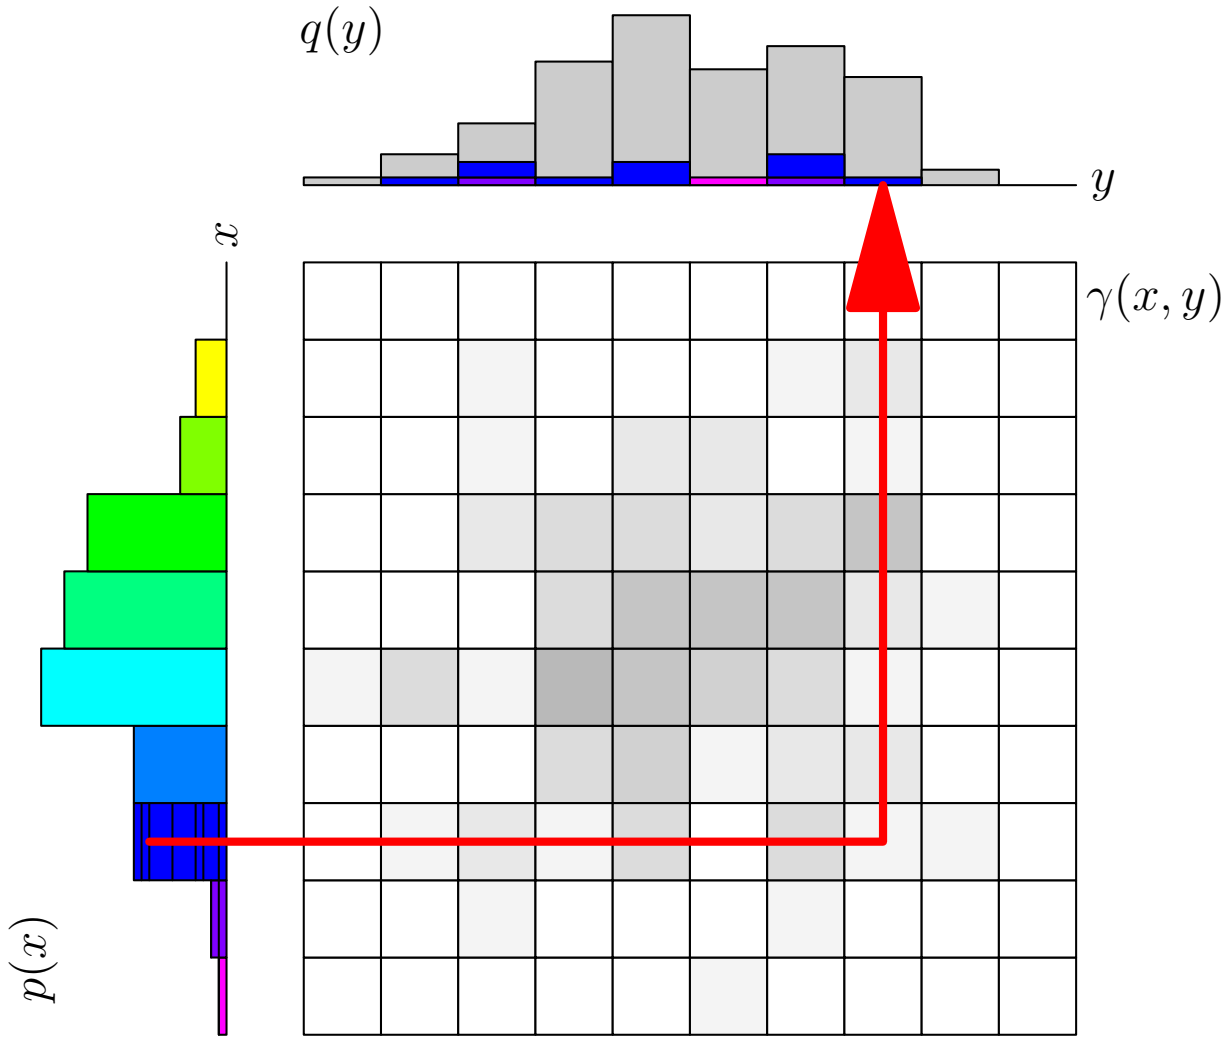
Transportation Policy



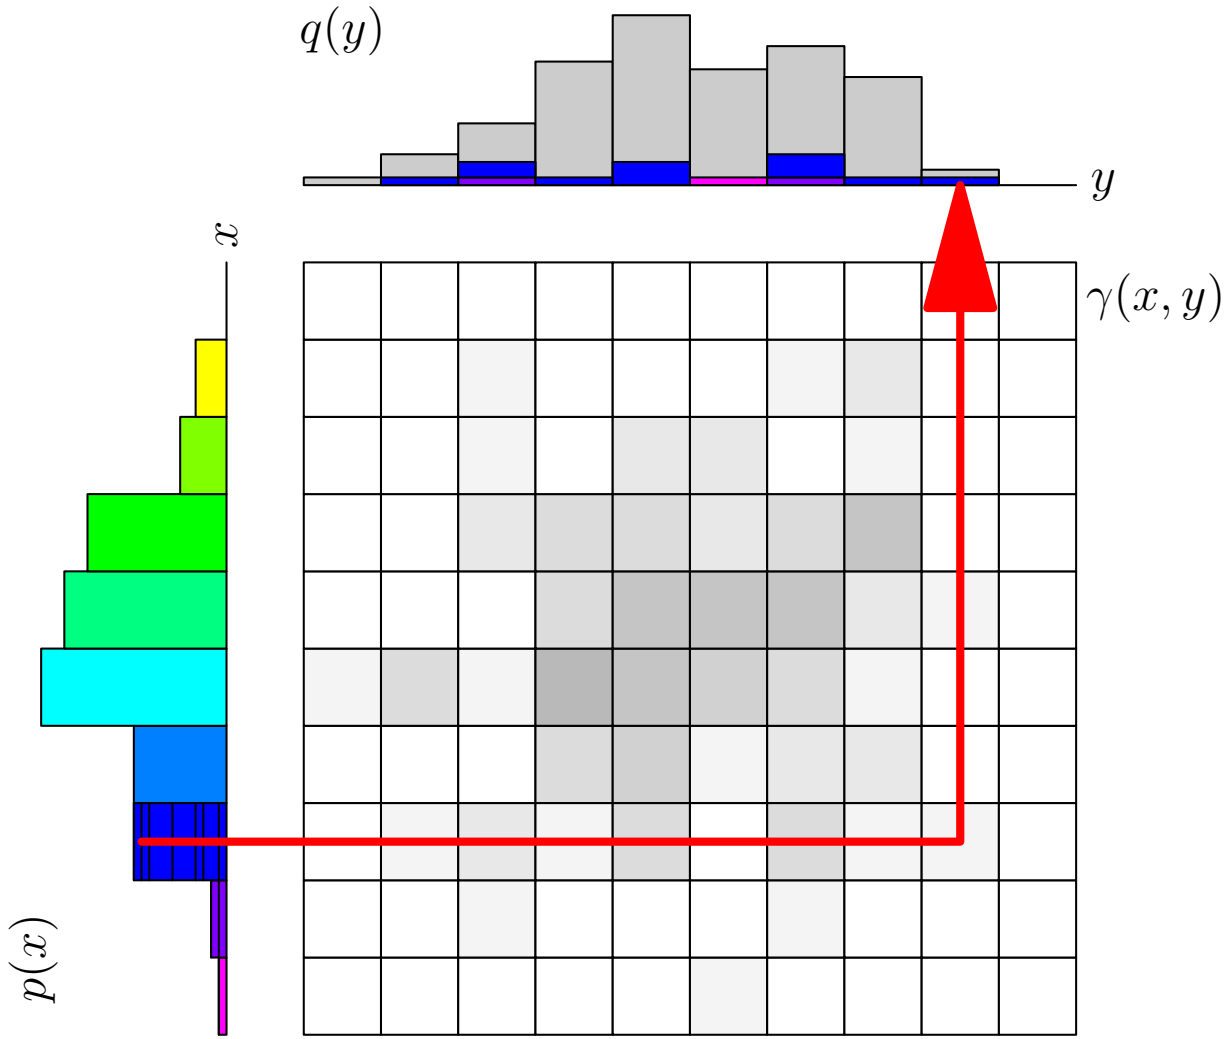
Transportation Policy



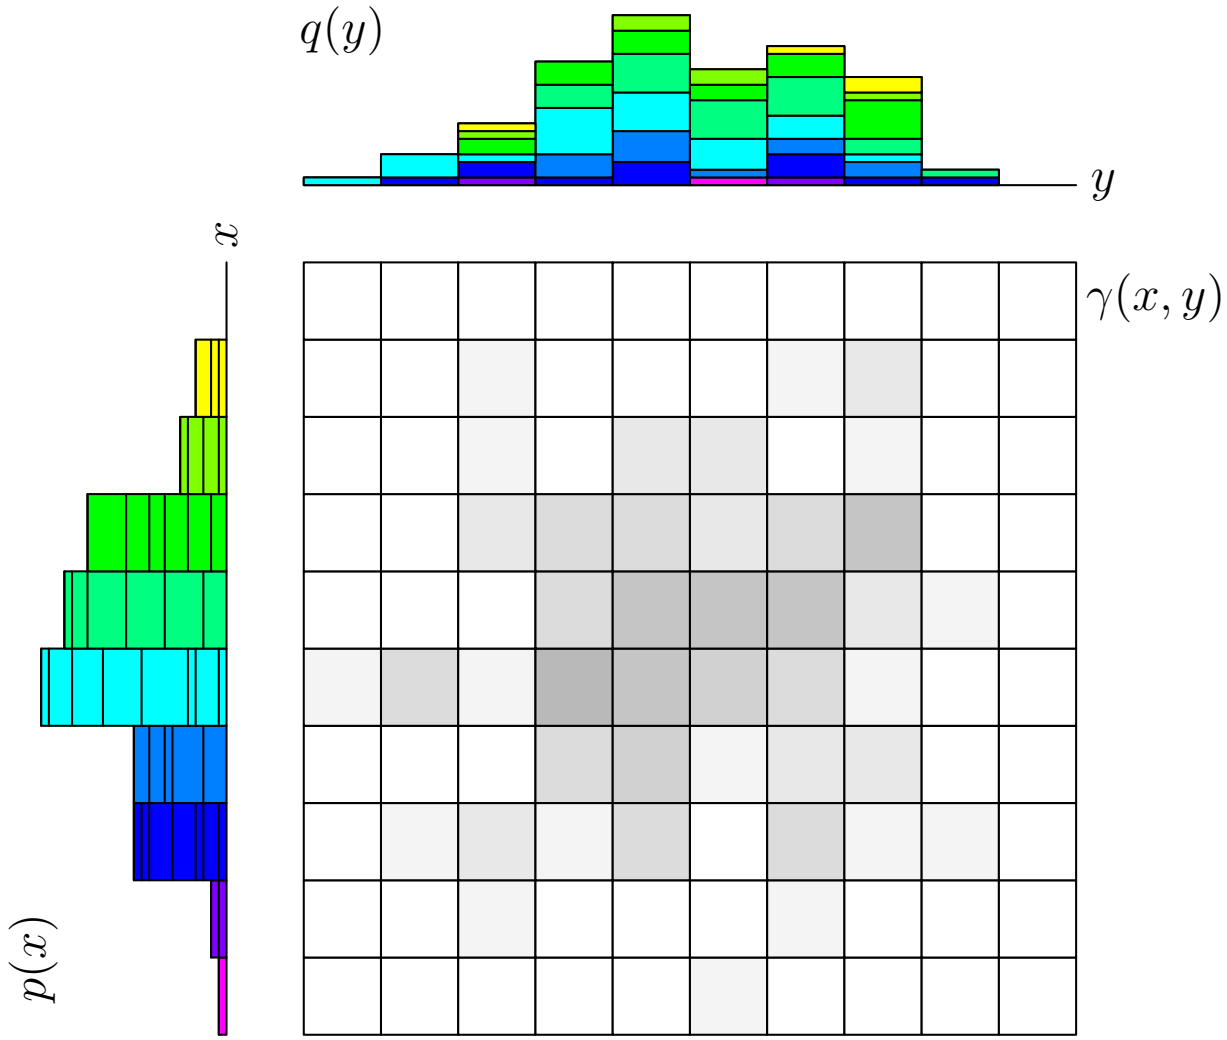
Transportation Policy



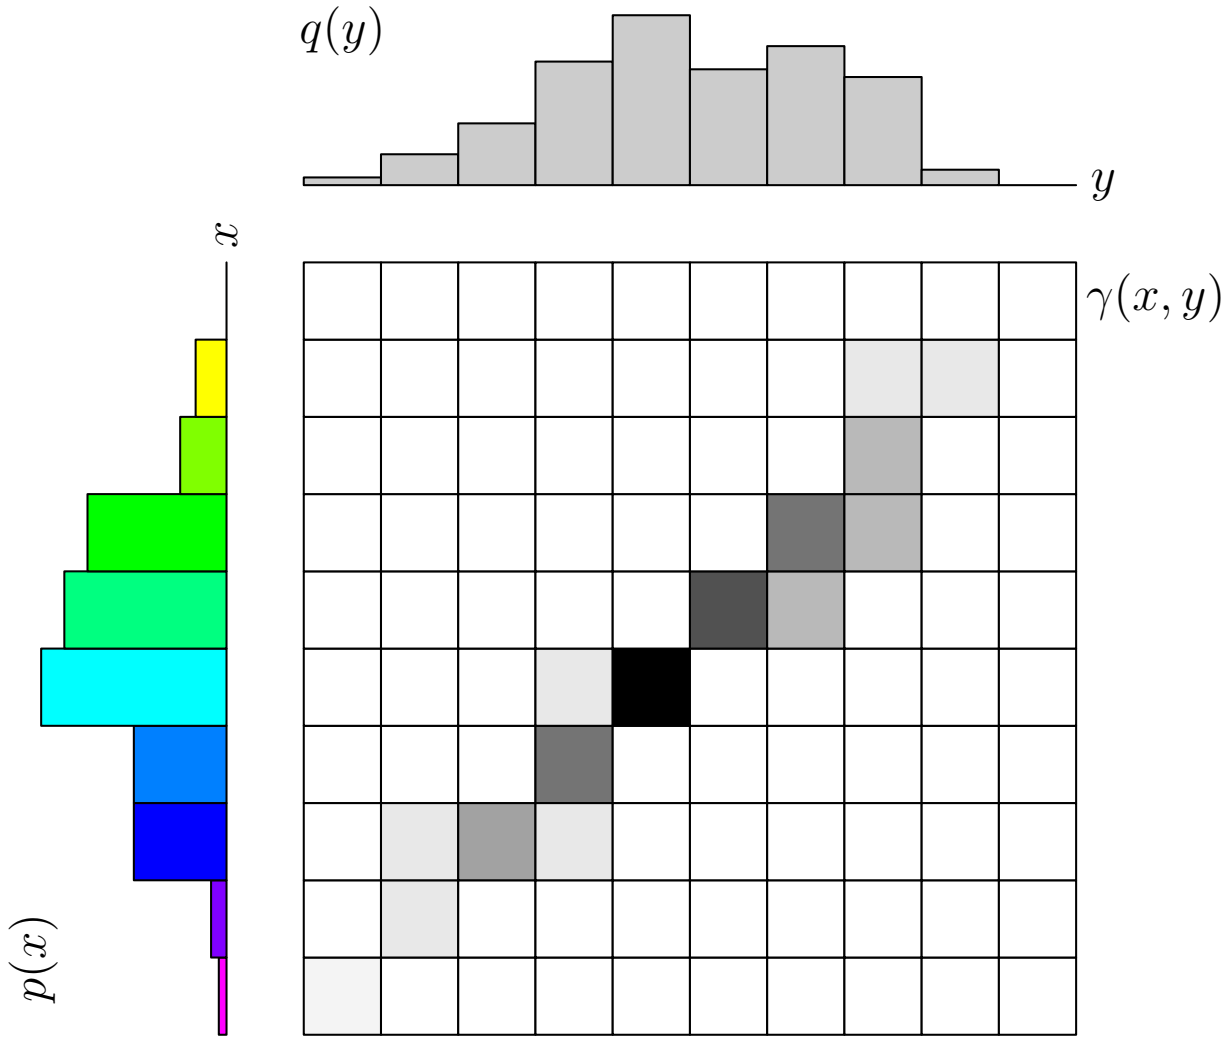
Transportation Policy



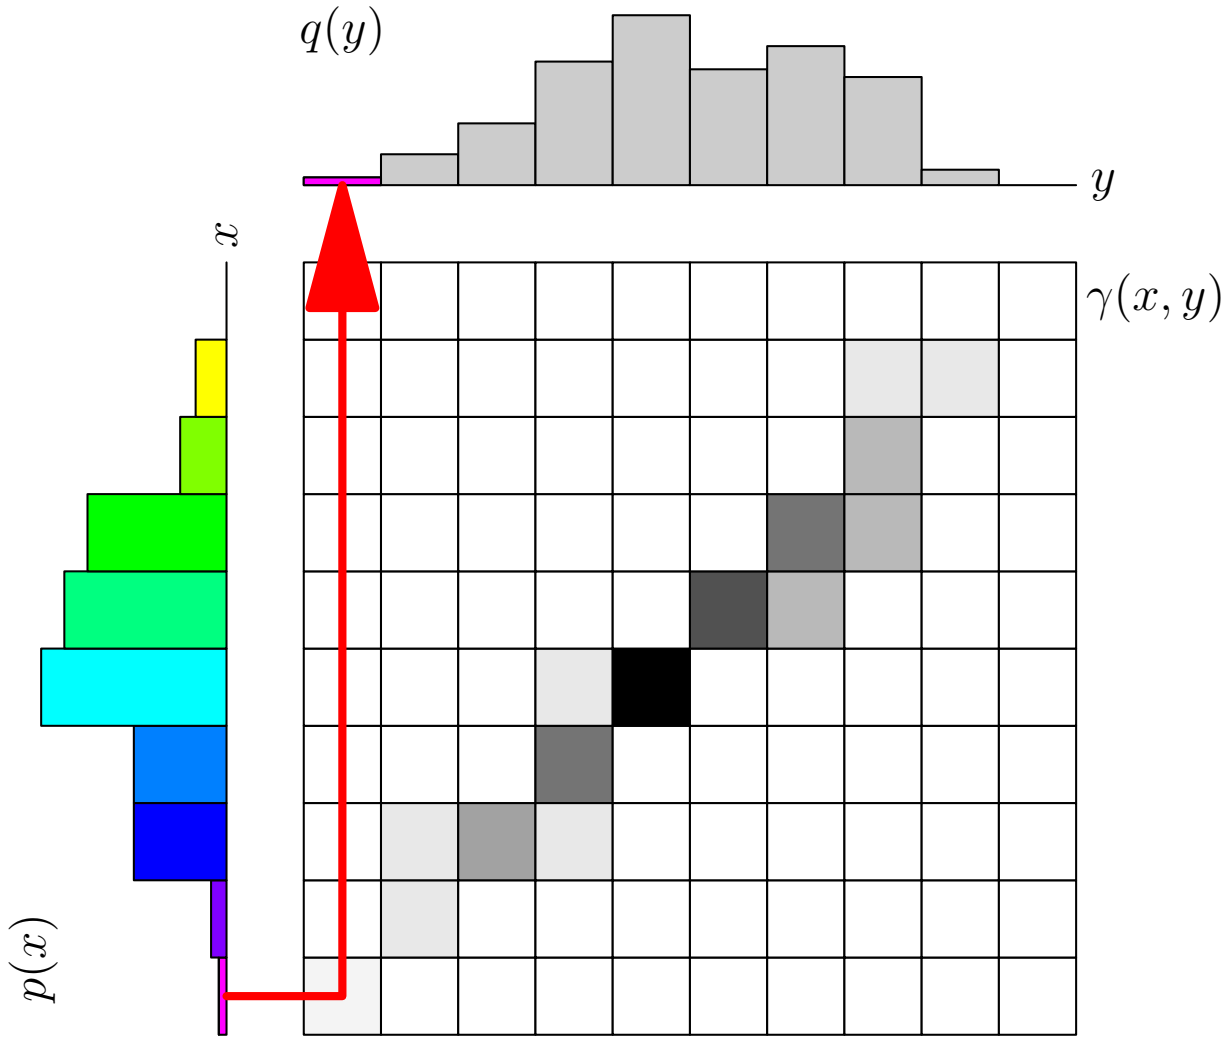
Transportation Policy



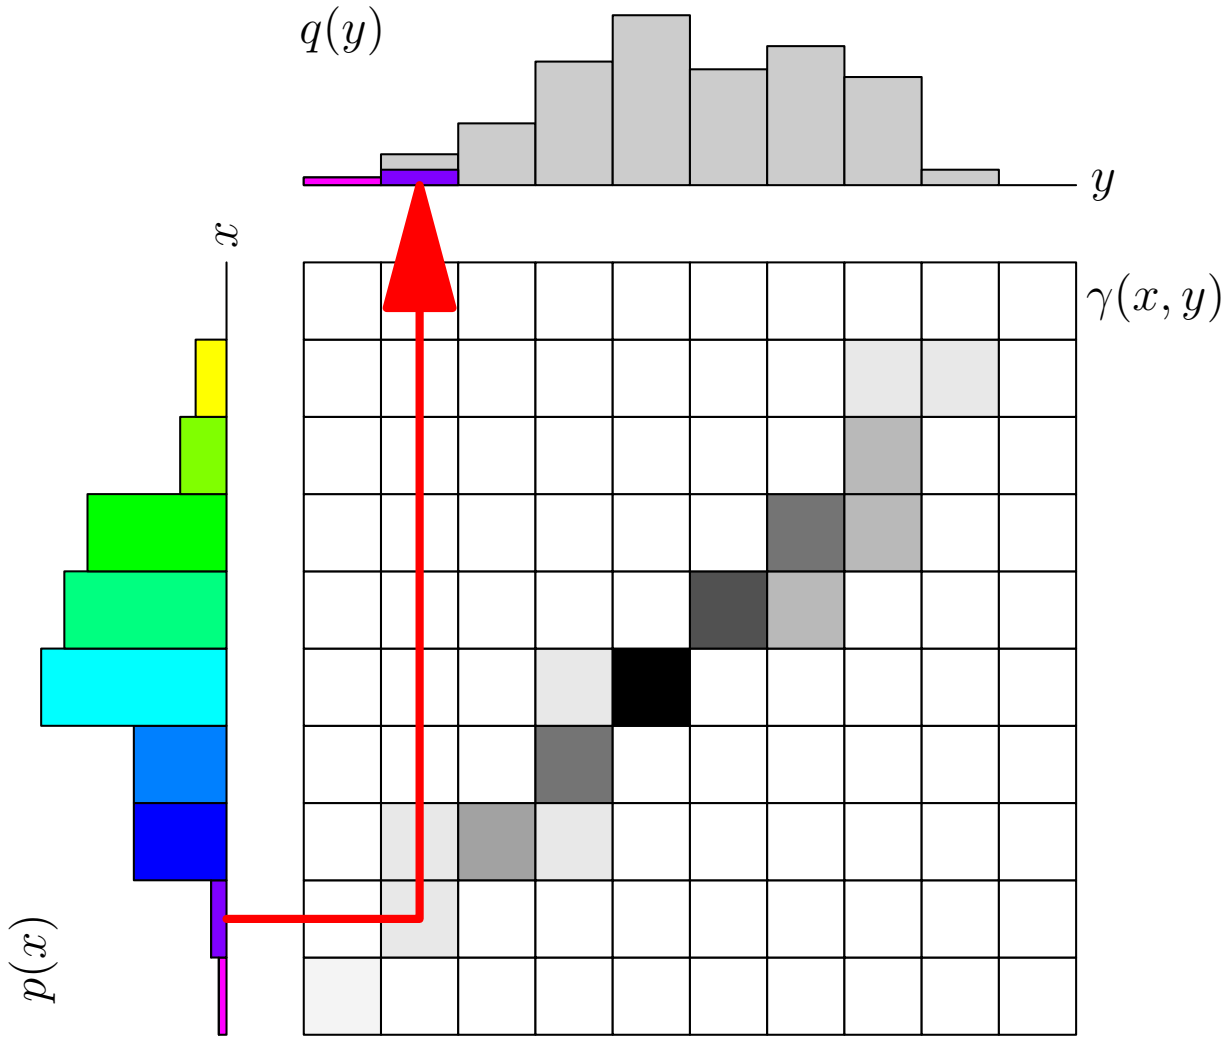
Transportation Policy



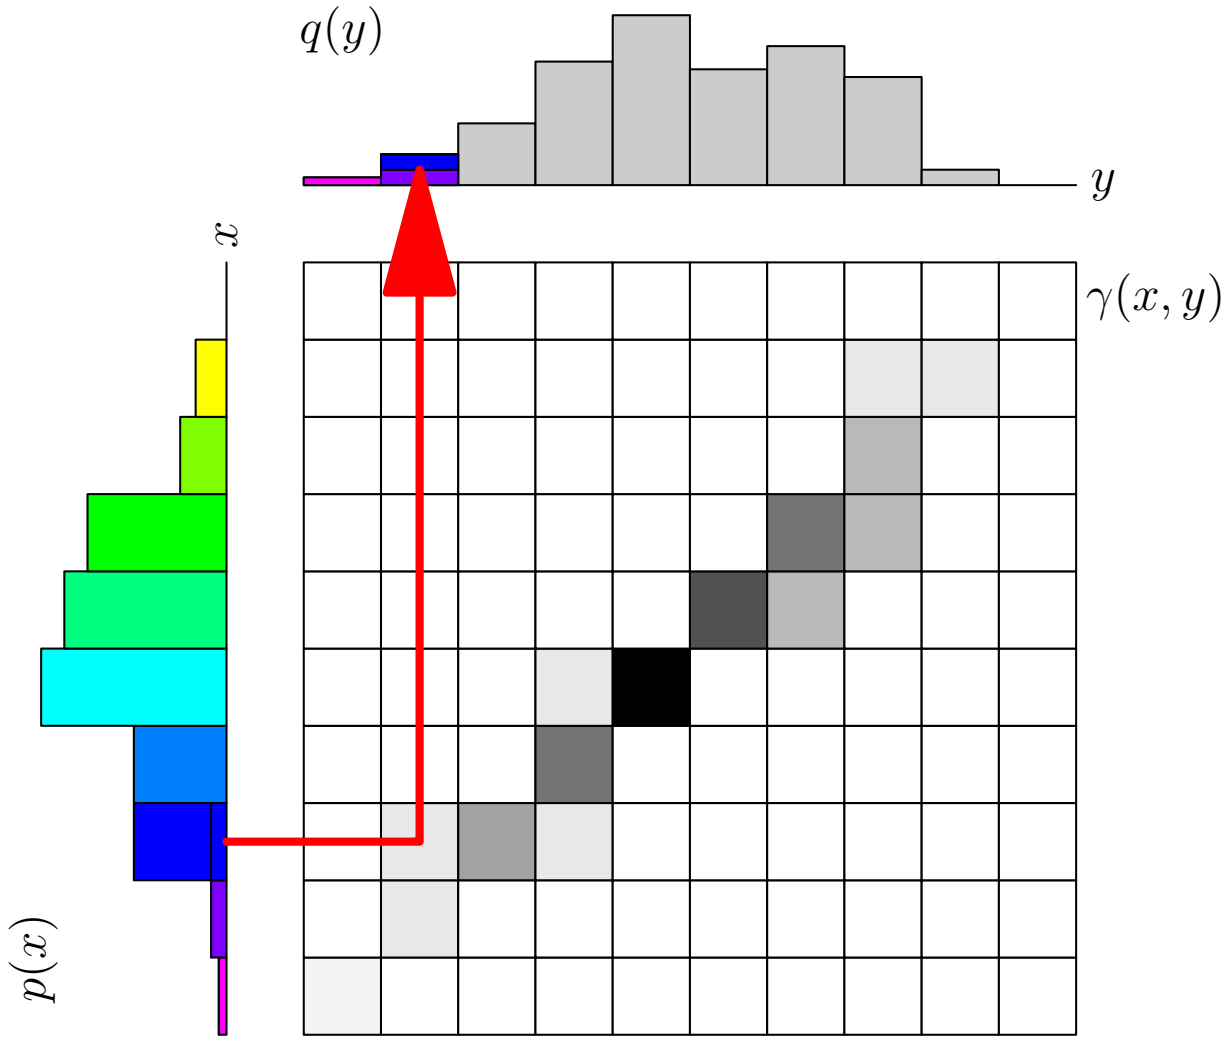
Transportation Policy



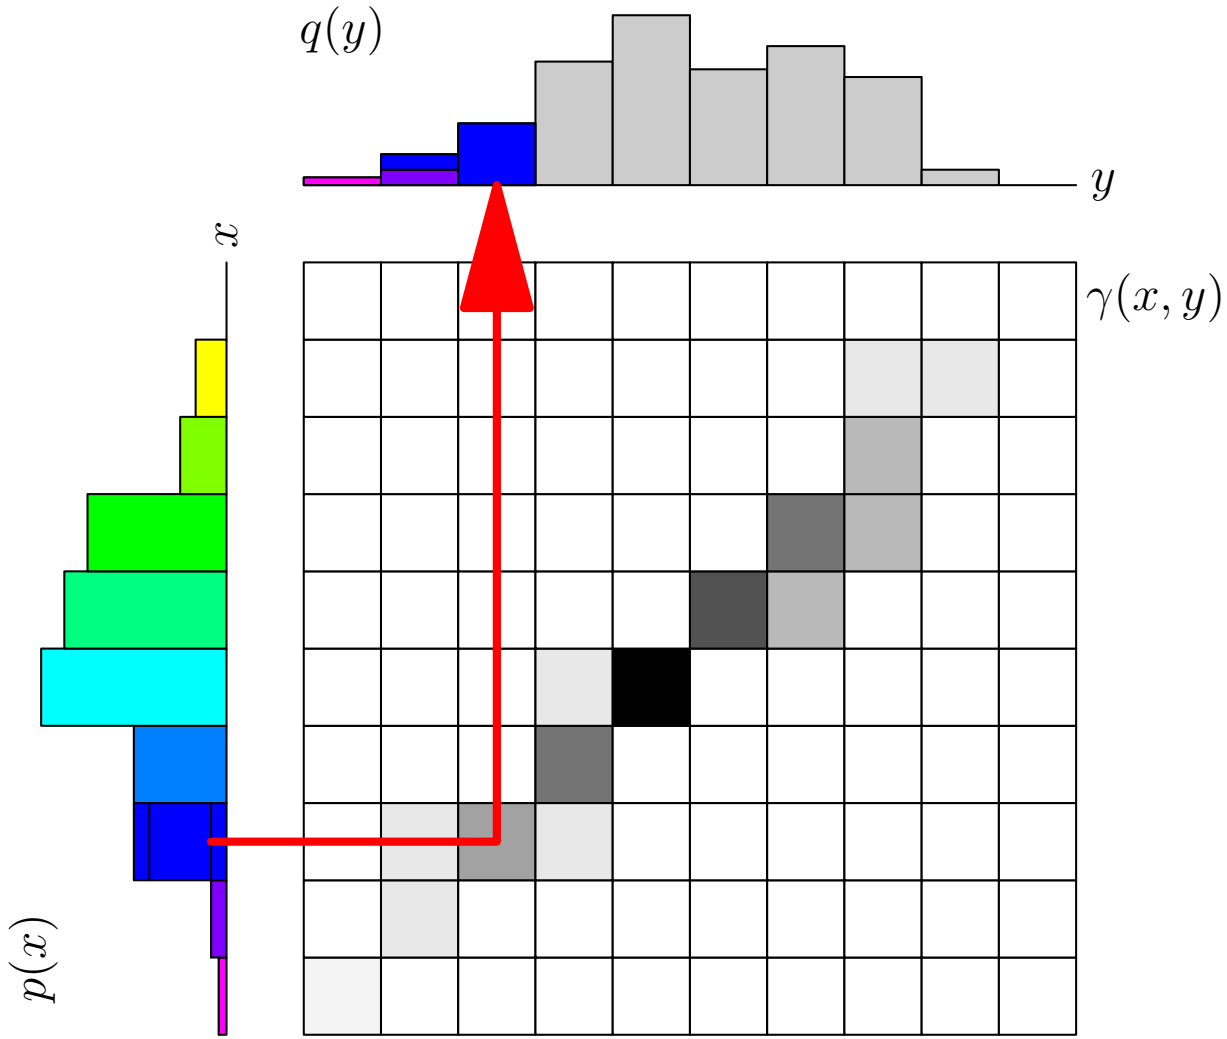
Transportation Policy



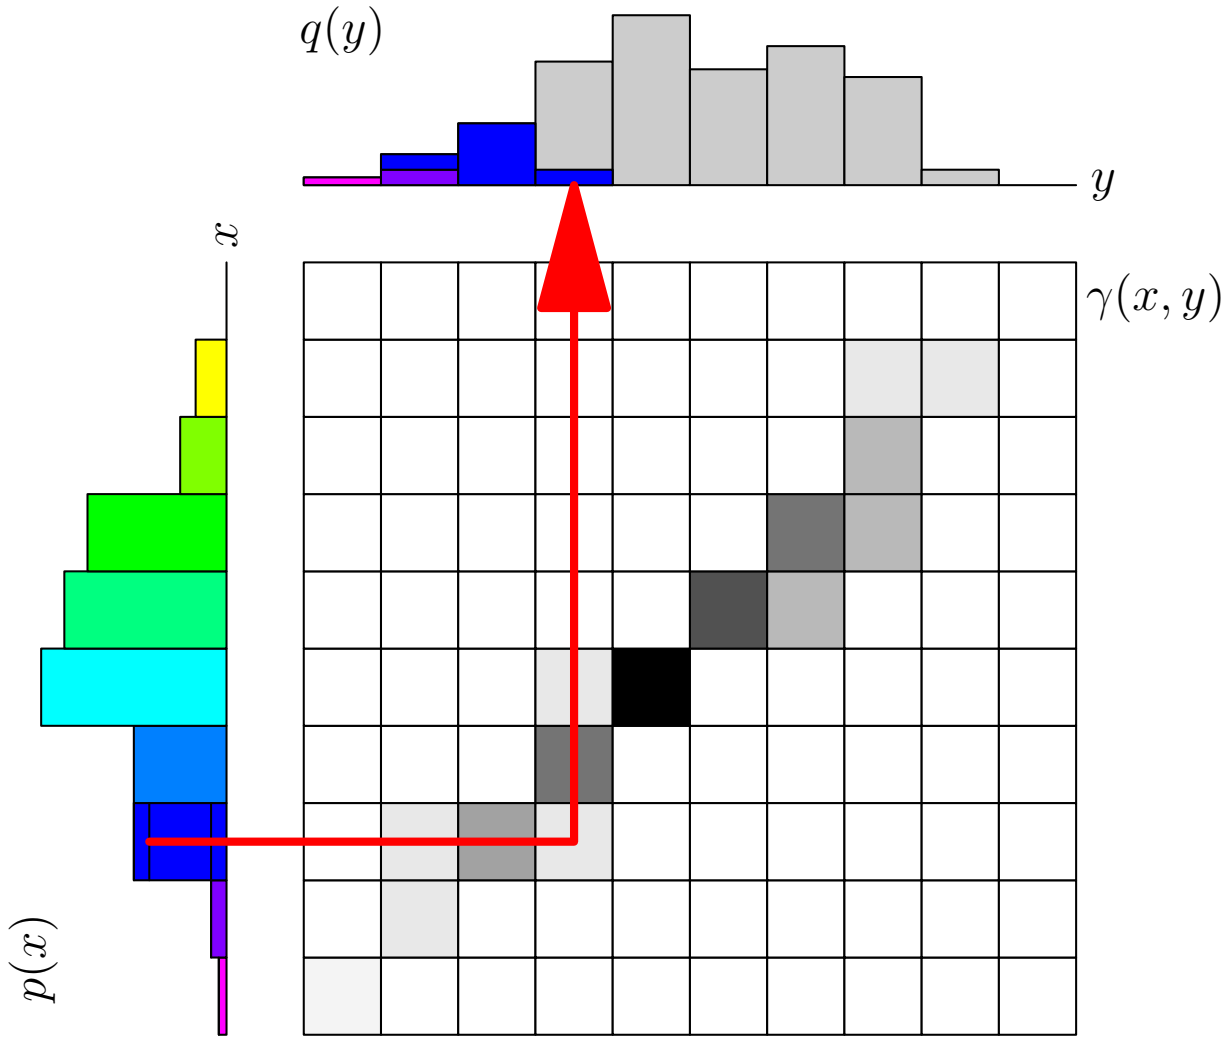
Transportation Policy



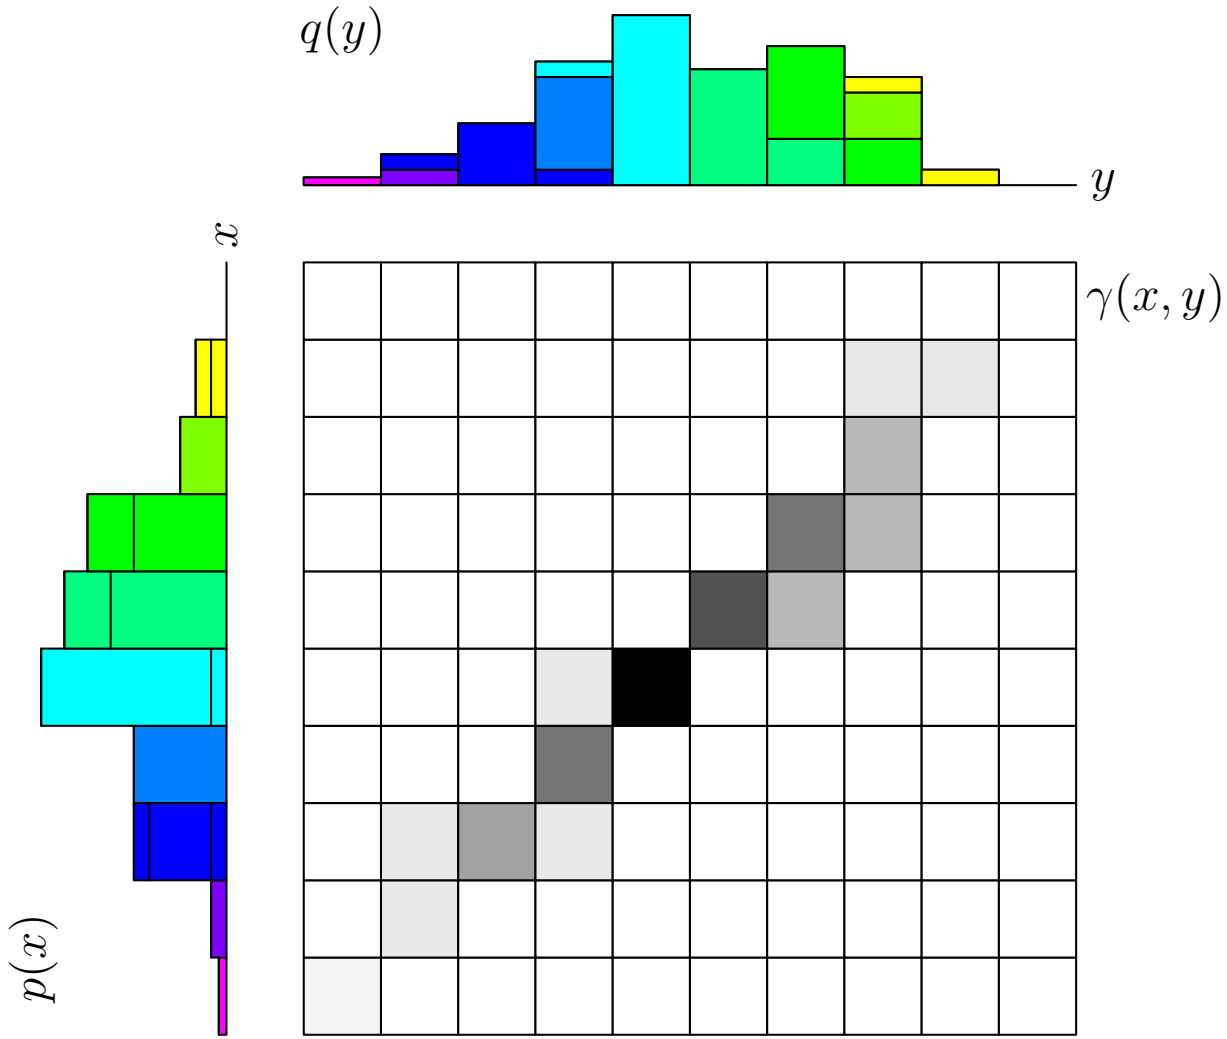
Transportation Policy



Transportation Policy



Transportation Policy



The Cost of Transport

- We want the transportation policy that minimises the total work: the probability mass moved times the distance it is moved
- Let $d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x}-\boldsymbol{y}\|$ be a distance measure then the cost of a transportation policy is

$$C(\gamma) = \int\int d(\boldsymbol{x},\boldsymbol{y})\gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x}\mathrm{d}\boldsymbol{y}$$

- Usually we take $d(\boldsymbol{x},\boldsymbol{y})$ to be the Euclidean distance, but we can choose any distance

The Cost of Transport

- We want the transportation policy that minimises the total work: the probability mass moved times the distance it is moved
- Let $d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x}-\boldsymbol{y}\|$ be a distance measure then the cost of a transportation policy is

$$C(\gamma) = \int\int d(\boldsymbol{x},\boldsymbol{y})\gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x}\mathrm{d}\boldsymbol{y}$$

- Usually we take $d(\boldsymbol{x},\boldsymbol{y})$ to be the Euclidean distance, but we can choose any distance

The Cost of Transport

- We want the transportation policy that minimises the total work: the probability mass moved times the distance it is moved
- Let $d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x}-\boldsymbol{y}\|$ be a distance measure then the cost of a transportation policy is

$$C(\gamma) = \int\int d(\boldsymbol{x},\boldsymbol{y})\gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x}\mathrm{d}\boldsymbol{y} = \mathbb{E}_{\gamma}[d(\boldsymbol{x},\boldsymbol{y})]$$

  where we interpret $\gamma(\boldsymbol{x},\boldsymbol{y})$ as a probability distribution

- Usually we take $d(\boldsymbol{x},\boldsymbol{y})$ to be the Euclidean distance, but we can choose any distance

The Cost of Transport

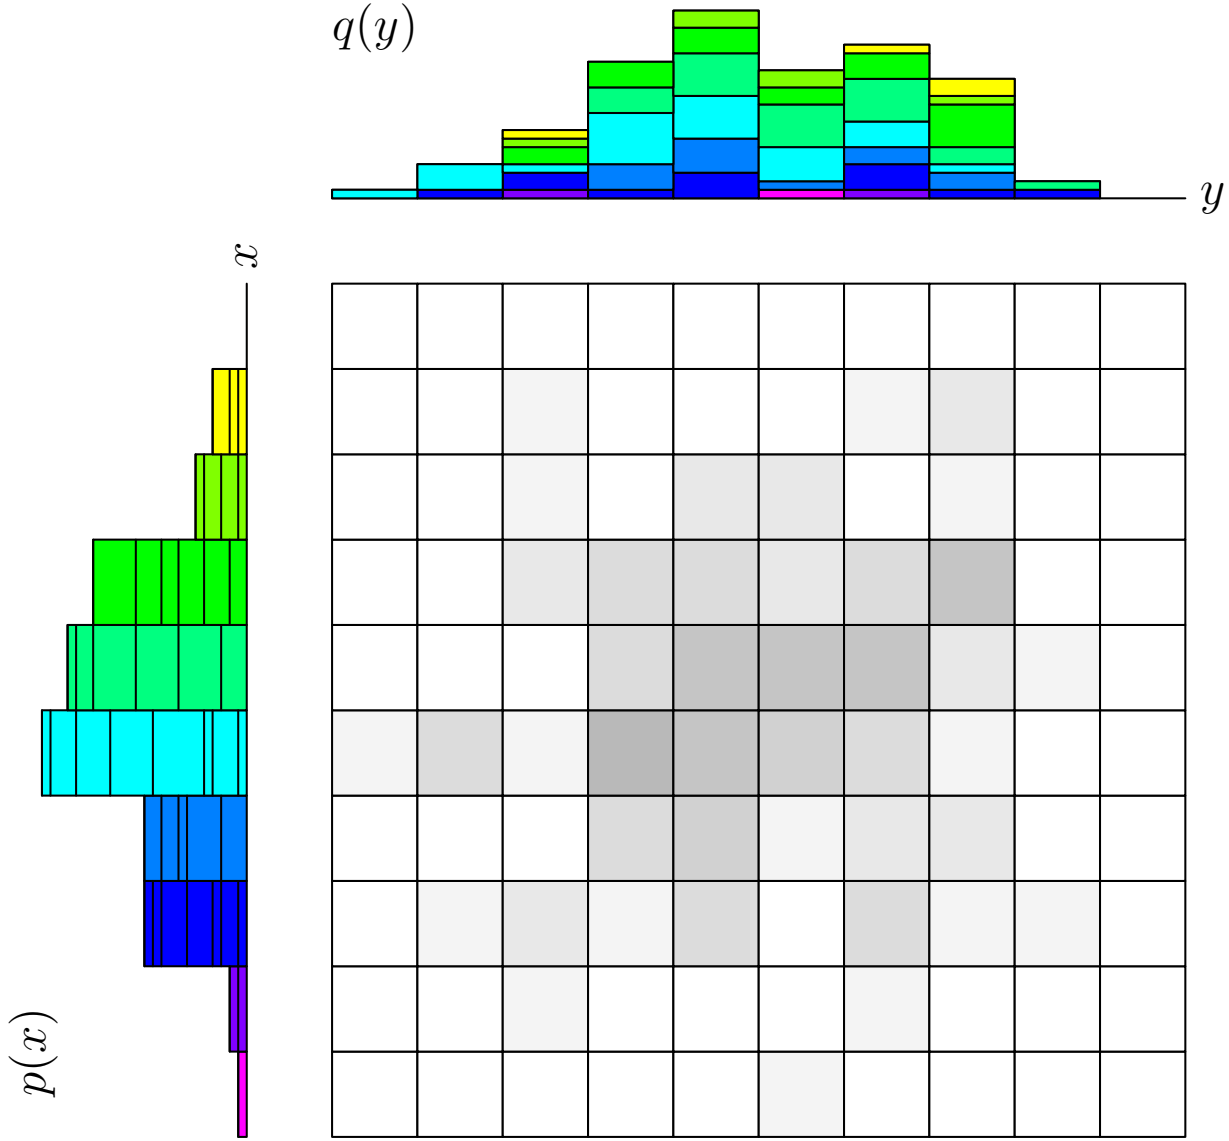
- We want the transportation policy that minimises the total work: the probability mass moved times the distance it is moved
- Let $d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x}-\boldsymbol{y}\|$ be a distance measure then the cost of a transportation policy is

$$C(\gamma) = \int\int d(\boldsymbol{x},\boldsymbol{y})\gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x}\mathrm{d}\boldsymbol{y} = \mathbb{E}_\gamma[d(\boldsymbol{x},\boldsymbol{y})]$$

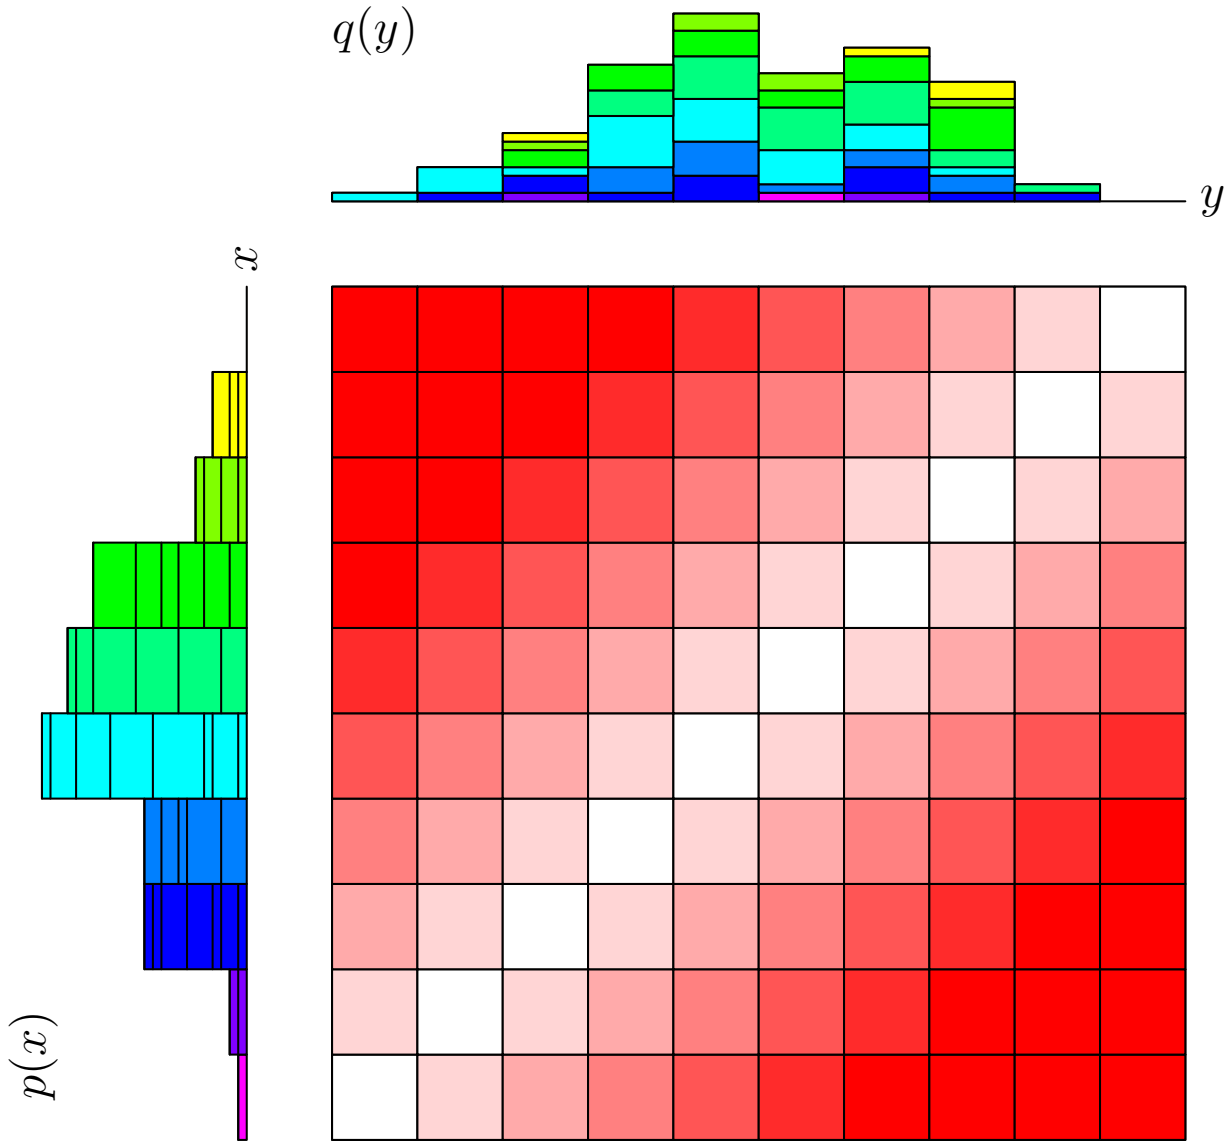
  where we interpret $\gamma(\boldsymbol{x},\boldsymbol{y})$ as a probability distribution

- Usually we take $d(\boldsymbol{x},\boldsymbol{y})$ to be the Euclidean distance, but we can choose any distance

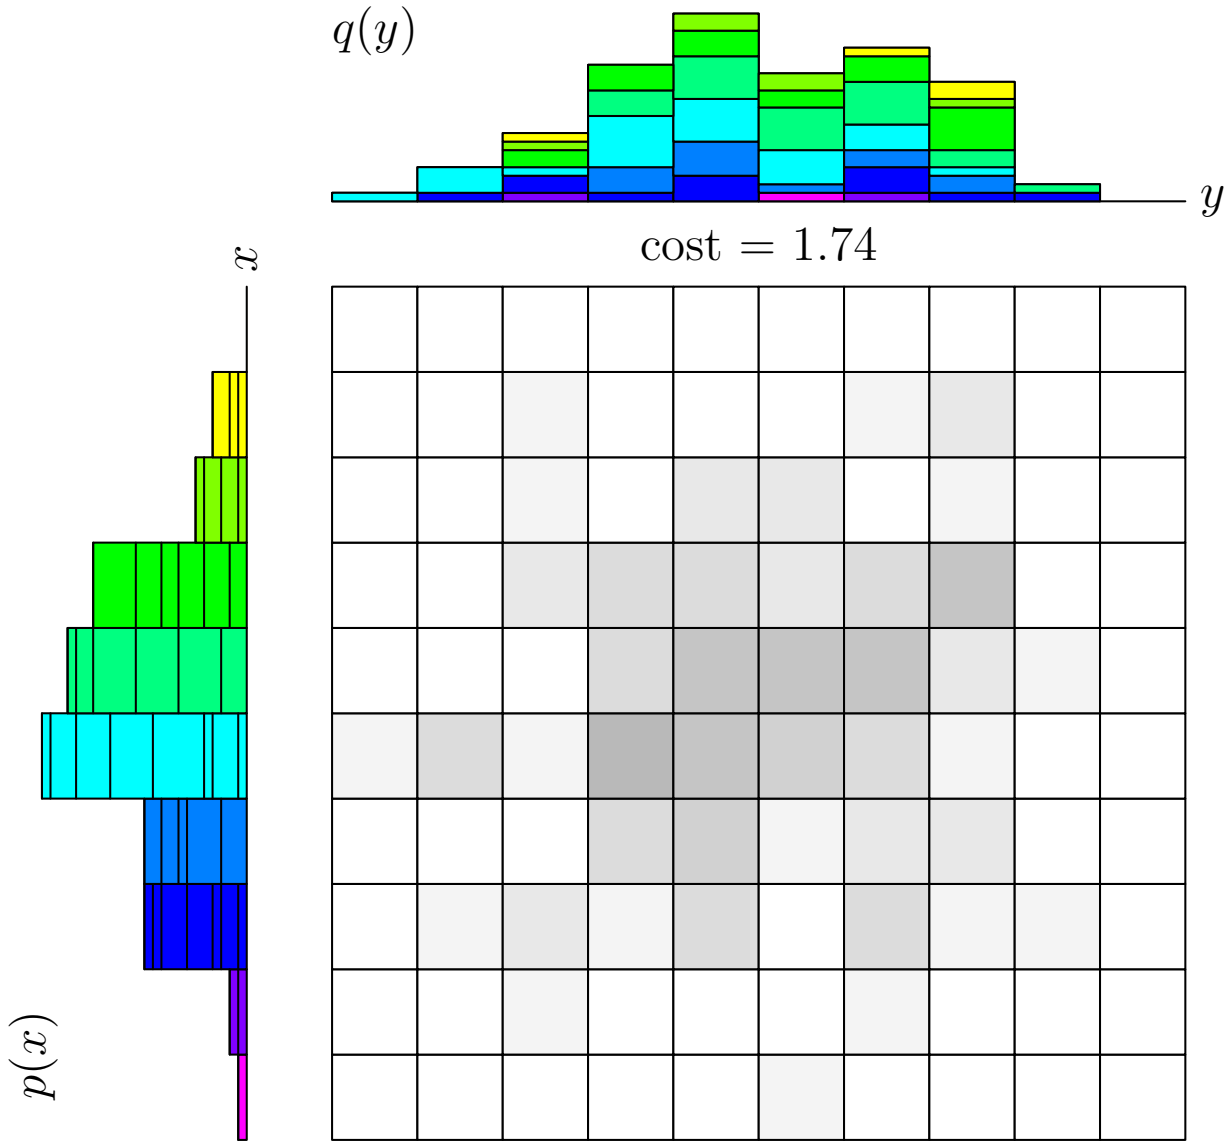
Transportation Cost



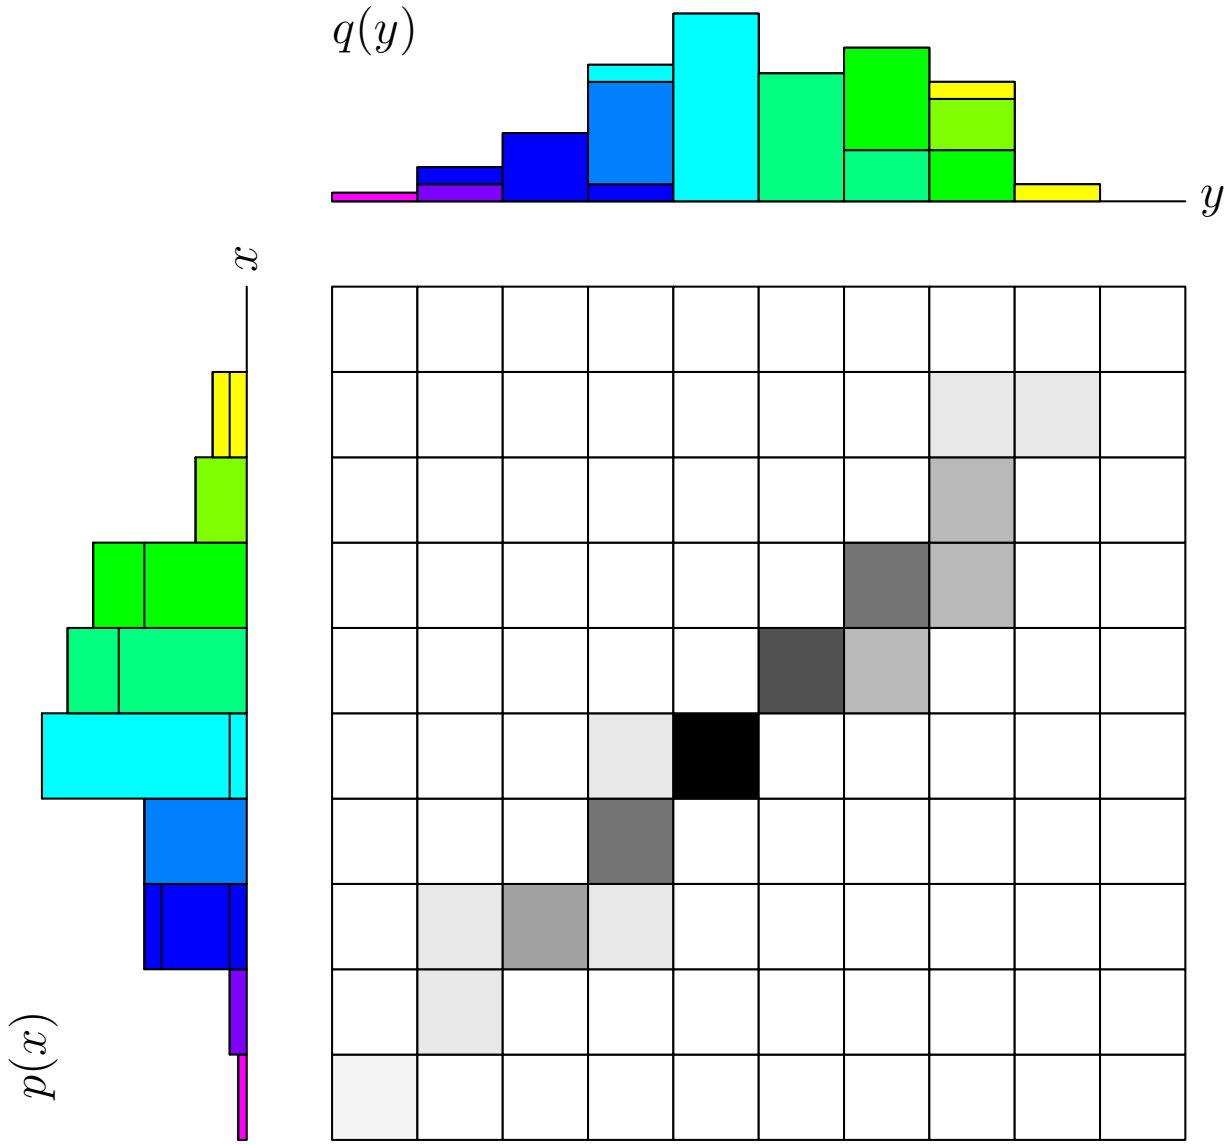
Transportation Cost



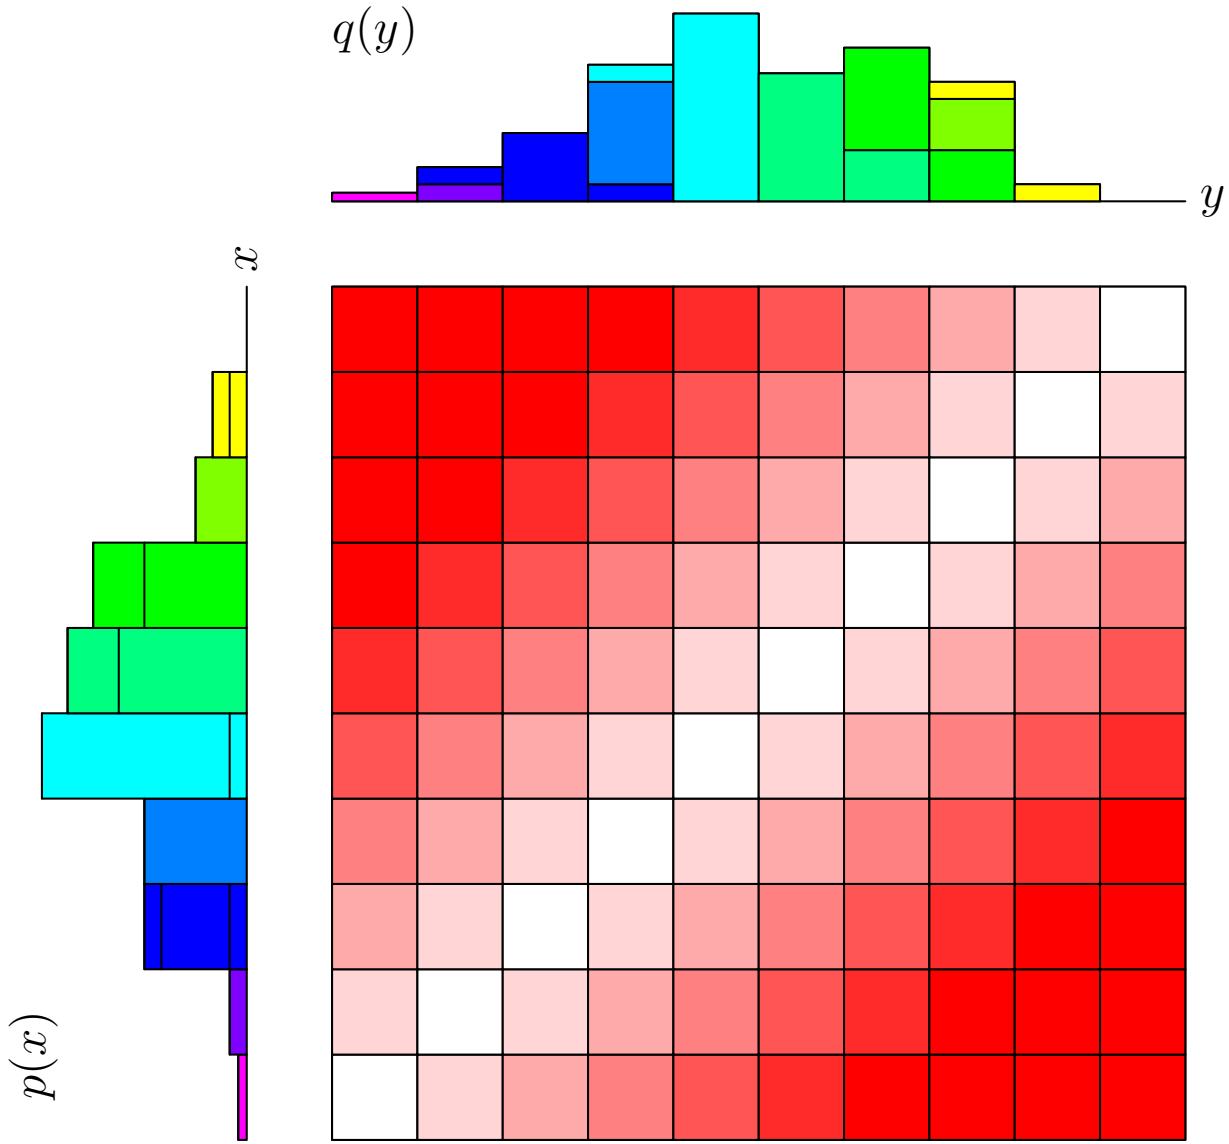
Transportation Cost



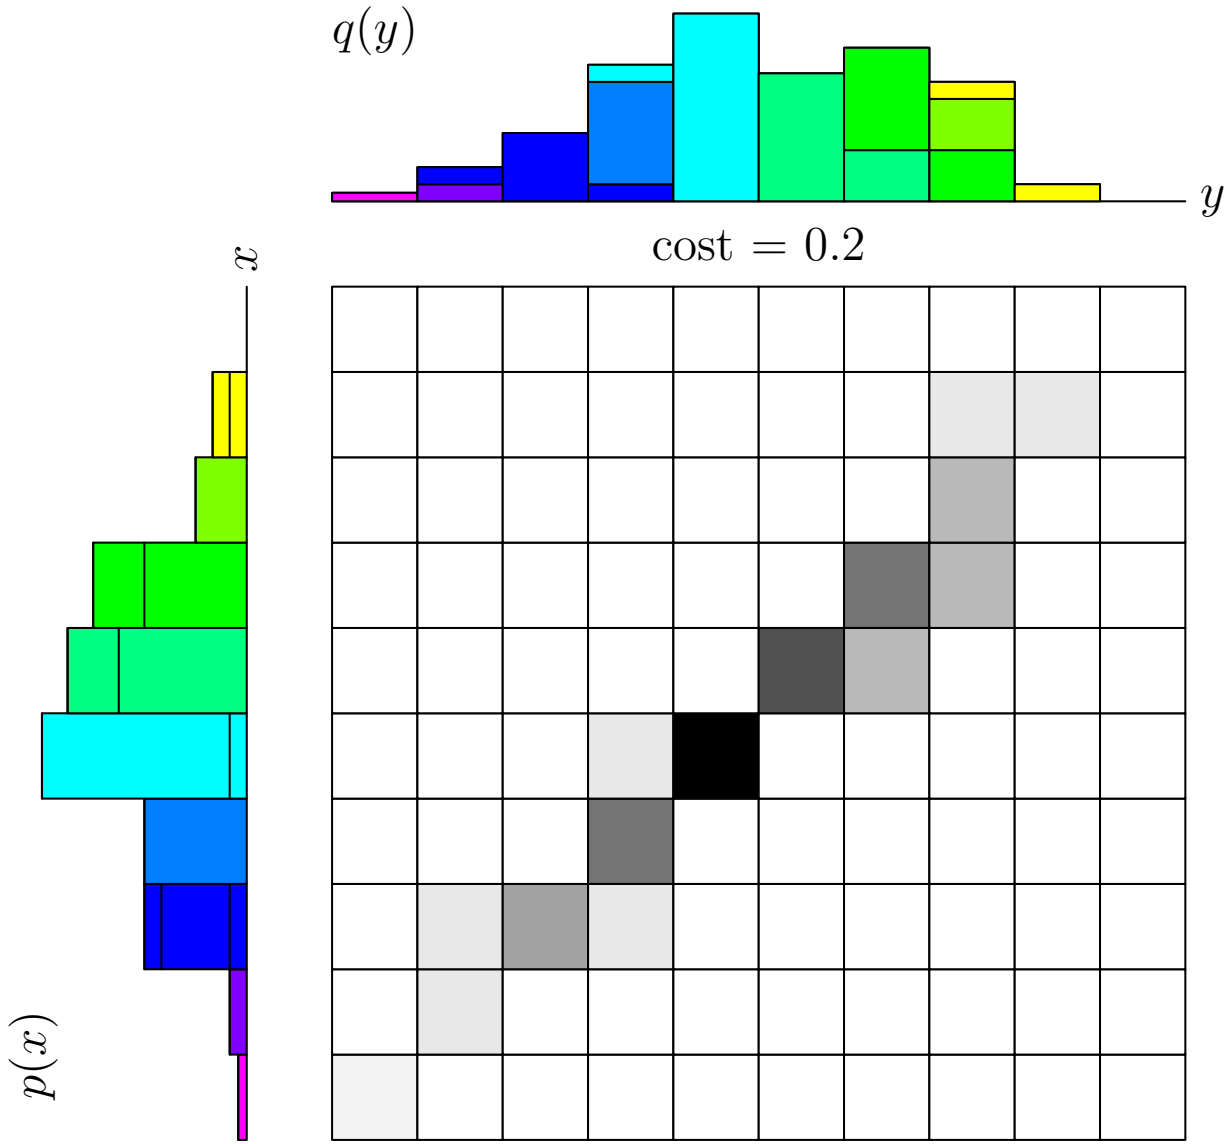
Transportation Cost



Transportation Cost



Transportation Cost



The Wasserstein Distance

- The Wasserstein distance $W(p,q)$ between two probability distributions is defined as

$$W(p,q) = \min_{\gamma \in \Lambda(p,q)} \mathbb{E}_\gamma[d(\boldsymbol{x},\boldsymbol{y})]$$

- Where $\Lambda(p,q)$ is the set of joint distributions $\gamma(\boldsymbol{x},\boldsymbol{y})$ such that

$$\int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{y} = p(\boldsymbol{x}) \qquad \int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x} = q(\boldsymbol{y})$$

- With cost d^k we get $W_k(p,q) = \left(\min_\gamma \mathbb{E}_\gamma\left[d(\boldsymbol{x},\boldsymbol{y})^k\right]\right)^{1/k}$; here we use $W = W_1$

The Wasserstein Distance

- The Wasserstein distance $W(p,q)$ between two probability distributions is defined as

$$W(p,q) = \min_{\gamma \in \Lambda(p,q)} \mathbb{E}_\gamma[d(\boldsymbol{x},\boldsymbol{y})]$$

- Where $\Lambda(p,q)$ is the set of joint distributions $\gamma(\boldsymbol{x},\boldsymbol{y})$ such that

$$\int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{y} = p(\boldsymbol{x}) \qquad \int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x} = q(\boldsymbol{y})$$

- With cost d^k we get $W_k(p,q) = \left(\min_\gamma \mathbb{E}_\gamma\left[d(\boldsymbol{x},\boldsymbol{y})^k\right]\right)^{1/k}$; here we use $W = W_1$

The Wasserstein Distance

- The Wasserstein distance $W(p,q)$ between two probability distributions is defined as

$$W(p,q) = \min_{\gamma \in \Lambda(p,q)} \mathbb{E}_\gamma[d(\boldsymbol{x},\boldsymbol{y})]$$

- Where $\Lambda(p,q)$ is the set of joint distributions $\gamma(\boldsymbol{x},\boldsymbol{y})$ such that

$$\int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{y} = p(\boldsymbol{x}) \qquad \int \gamma(\boldsymbol{x},\boldsymbol{y})\mathrm{d}\boldsymbol{x} = q(\boldsymbol{y})$$

- With cost d^k we get $W_k(p,q) = \left(\min_\gamma \mathbb{E}_\gamma\left[d(\boldsymbol{x},\boldsymbol{y})^k\right]\right)^{1/k}$; here we use $W = W_1$

One Dimension

- In one dimension the Wasserstein distance has a closed form in terms of the cumulative distribution functions F_p and F_q

$$W(p,q) = \int_{-\infty}^{\infty} \left| F_p(x) - F_q(x) \right| \mathrm{d}x$$

  the area between the two curves

- For two samples of the same size: sort both and match them in order

$$W = \frac{1}{n} \sum_{i=1}^{n} \left| x_{(i)} - y_{(i)} \right|$$

- In higher dimensions there is no such shortcut: we have to solve the optimisation

One Dimension

- In one dimension the Wasserstein distance has a closed form in terms of the cumulative distribution functions F_p and F_q

$$W(p,q) = \int_{-\infty}^{\infty} \left| F_p(x) - F_q(x) \right| \mathrm{d}x$$

the area between the two curves

- For two samples of the same size: sort both and match them in order

$$W = \frac{1}{n} \sum_{i=1}^{n} \left| x_{(i)} - y_{(i)} \right|$$

- In higher dimensions there is no such shortcut: we have to solve the optimisation

One Dimension

- In one dimension the Wasserstein distance has a closed form in terms of the cumulative distribution functions F_p and F_q

$$W(p,q) = \int_{-\infty}^{\infty} \left| F_p(x) - F_q(x) \right| \mathrm{d}x$$

  the area between the two curves

- For two samples of the same size: sort both and match them in order

$$W = \frac{1}{n} \sum_{i=1}^{n} \left| x_{(i)} - y_{(i)} \right|$$

- In higher dimensions there is no such shortcut: we have to solve the optimisation

One Dimension

- In one dimension the Wasserstein distance has a closed form in terms of the cumulative distribution functions F_p and F_q

$$W(p,q) = \int_{-\infty}^{\infty} \left| F_p(x) - F_q(x) \right| \mathrm{d}x$$

  the area between the two curves

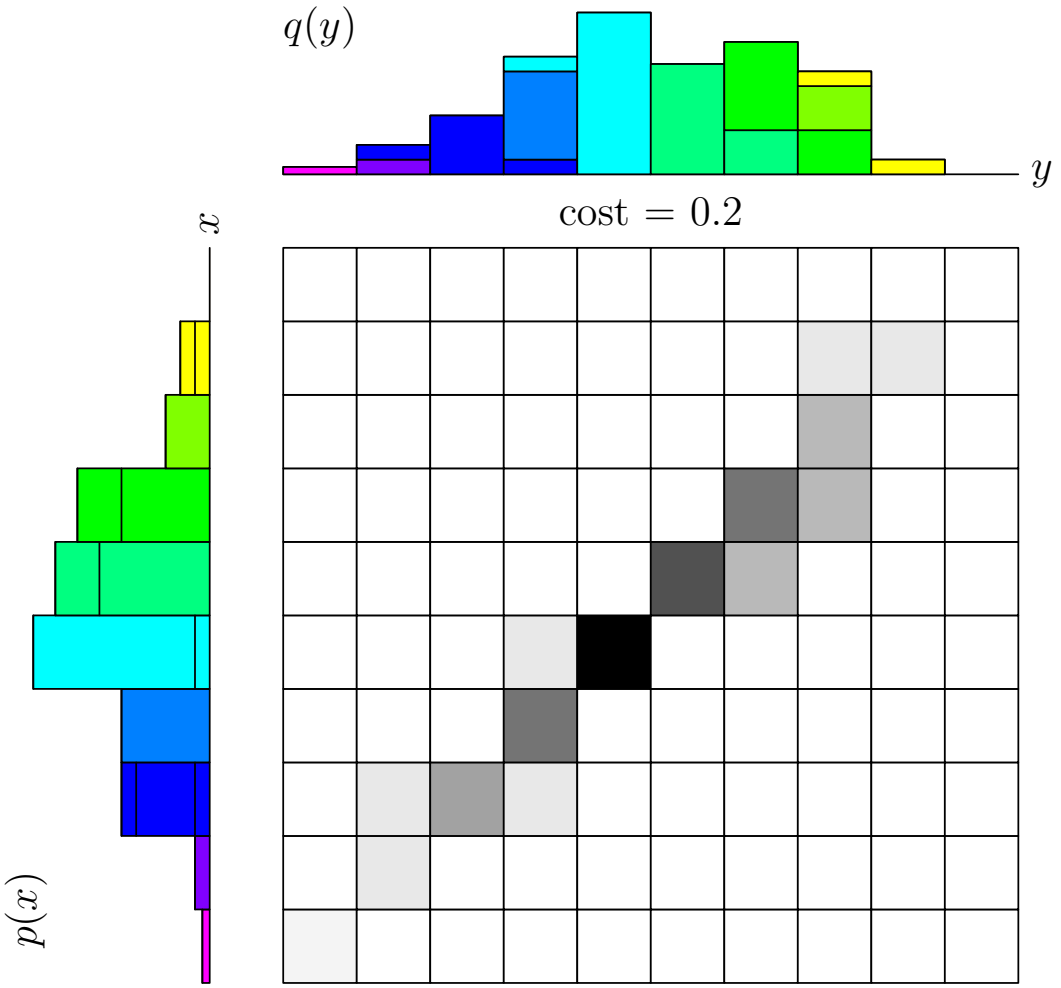
- For two samples of the same size: sort both and match them in order

$$W = \frac{1}{n} \sum_{i=1}^{n} \left| x_{(i)} - y_{(i)} \right|$$

- In higher dimensions there is no such shortcut: we have to solve the optimisation

Outline

1. Comparing Distributions
2. Optimal Transport
3. **Duality**
4. Applications



Computing the Wasserstein Distance

- To compute the Wasserstein distance we have to solve a minimisation task!
- This looks nasty, but it is a (continuous) linear programme, like those in the constrained-optimisation lecture
- Suppose p and q were discrete distributions (i.e. $\boldsymbol{x}$ and $\boldsymbol{y}$ only take discrete points)
- Then we could treat each value of $\gamma(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{\gamma}$ and each value of $d(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{D}$
- Our objective is to choose $\boldsymbol{\gamma}$ to minimise $\boldsymbol{D}^{\mathsf{T}}\boldsymbol{\gamma}$

Computing the Wasserstein Distance

- To compute the Wasserstein distance we have to solve a minimisation task!
- This looks nasty, but it is a (continuous) linear programme, like those in the constrained-optimisation lecture
- Suppose p and q were discrete distributions (i.e. $\boldsymbol{x}$ and $\boldsymbol{y}$ only take discrete points)
- Then we could treat each value of $\gamma(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{\gamma}$ and each value of $d(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{D}$
- Our objective is to choose $\boldsymbol{\gamma}$ to minimise $\boldsymbol{D}^{\mathsf{T}}\boldsymbol{\gamma}$

Computing the Wasserstein Distance

- To compute the Wasserstein distance we have to solve a minimisation task!
- This looks nasty, but it is a (continuous) linear programme, like those in the constrained-optimisation lecture
- Suppose p and q were discrete distributions (i.e. $\boldsymbol{x}$ and $\boldsymbol{y}$ only take discrete points)
- Then we could treat each value of $\gamma(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{\gamma}$ and each value of $d(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{D}$
- Our objective is to choose $\boldsymbol{\gamma}$ to minimise $\boldsymbol{D}^{\mathsf{T}}\boldsymbol{\gamma}$

Computing the Wasserstein Distance

- To compute the Wasserstein distance we have to solve a minimisation task!
- This looks nasty, but it is a (continuous) linear programme, like those in the constrained-optimisation lecture
- Suppose p and q were discrete distributions (i.e. $\boldsymbol{x}$ and $\boldsymbol{y}$ only take discrete points)
- Then we could treat each value of $\gamma(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{\gamma}$ and each value of $d(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{D}$
- Our objective is to choose $\boldsymbol{\gamma}$ to minimise $\boldsymbol{D}^{\mathsf{T}}\boldsymbol{\gamma}$

Computing the Wasserstein Distance

- To compute the Wasserstein distance we have to solve a minimisation task!
- This looks nasty, but it is a (continuous) linear programme, like those in the constrained-optimisation lecture
- Suppose p and q were discrete distributions (i.e. $\boldsymbol{x}$ and $\boldsymbol{y}$ only take discrete points)
- Then we could treat each value of $\gamma(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{\gamma}$ and each value of $d(\boldsymbol{x},\boldsymbol{y})$ as an element of a vector $\boldsymbol{D}$
- Our objective is to choose $\boldsymbol{\gamma}$ to minimise $\boldsymbol{D}^{\mathsf{T}}\boldsymbol{\gamma}$

Constraints

$$\sum_j \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) = p(\boldsymbol{x}_i) \qquad \sum_i \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) = q(\boldsymbol{y}_j)$$

$$\mathbf{A}\boldsymbol{\gamma} = \boldsymbol{P}$$

$$\begin{pmatrix}
1 & 1 & \cdots & 1 & 0 & 0 & \cdots & 0 & \cdots & 0 & 0 & \cdots & 0 \\
0 & 0 & \cdots & 0 & 1 & 1 & \cdots & 1 & \cdots & 0 & 0 & \cdots & 0 \\
\vdots & \vdots & \cdots & \vdots & \vdots & \vdots & \cdots & \vdots & \vdots & \vdots & \cdots & \vdots & \vdots \\
0 & 0 & \cdots & 0 & 0 & 0 & \cdots & 0 & \cdots & 1 & 1 & \cdots & 1 \\
1 & 0 & \cdots & 0 & 1 & 0 & \cdots & 0 & \cdots & 1 & 0 & \cdots & 0 \\
0 & 1 & \cdots & 0 & 0 & 1 & \cdots & 0 & \cdots & 0 & 1 & \cdots & 0 \\
\vdots & \vdots & \cdots & \vdots & \vdots & \vdots & \cdots & \vdots & \vdots & \vdots & \cdots & \vdots & \vdots \\
0 & 0 & \cdots & 1 & 0 & 0 & \cdots & 1 & \cdots & 0 & 0 & \cdots & 1
\end{pmatrix}
\begin{pmatrix}
\gamma(x_1, y_1) \\
\gamma(x_2, y_1) \\
\vdots \\
\gamma(x_n, y_1) \\
\gamma(x_1, y_2) \\
\gamma(x_2, y_2) \\
\vdots \\
\gamma(x_n, y_2) \\
\vdots \\
\vdots \\
\vdots \\
\vdots \\
\gamma(x_1, y_n) \\
\gamma(x_2, y_n) \\
\vdots \\
\gamma(x_n, y_n)
\end{pmatrix}
=
\begin{pmatrix}
q(y_1) \\
q(y_2) \\
\vdots \\
q(y_n) \\
p(x_1) \\
p(x_2) \\
\vdots \\
p(x_n)
\end{pmatrix}$$

The Dual Problem

- With Lagrange multipliers $\alpha(\boldsymbol{x}_i)$ and $\beta(\boldsymbol{y}_j)$ for the two sets of constraints

$$\mathcal{L} = \sum_{i,j} d(\boldsymbol{x}_i, \boldsymbol{y}_j)\gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - \sum_i \alpha(\boldsymbol{x}_i)\left(\sum_j \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - p(\boldsymbol{x}_i)\right) - \sum_j \beta(\boldsymbol{y}_j)\left(\sum_i \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - q(\boldsymbol{y}_j)\right)$$

  subject to $\gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) \geq 0$

- Rearranging

$$\mathcal{L} = \sum_i \alpha(\boldsymbol{x}_i)p(\boldsymbol{x}_i) + \sum_j \beta(\boldsymbol{y}_j)q(\boldsymbol{y}_j) - \sum_{i,j} \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j)(\alpha(\boldsymbol{x}_i) + \beta(\boldsymbol{y}_j) - d(\boldsymbol{x}_i, \boldsymbol{y}_j))$$

- This is equivalent to maximising $\sum_i \alpha(\boldsymbol{x}_i)p(\boldsymbol{x}_i) + \sum_j \beta(\boldsymbol{y}_j)q(\boldsymbol{y}_j)$, subject to

$$\forall i, j \quad \alpha(\boldsymbol{x}_i) + \beta(\boldsymbol{y}_j) \leq d(\boldsymbol{x}_i, \boldsymbol{y}_j)$$

- For continuous distributions sums become integrals and $\alpha(\boldsymbol{x})$, $\beta(\boldsymbol{y})$ become functions

The Dual Problem

- With Lagrange multipliers $\alpha(\boldsymbol{x}_i)$ and $\beta(\boldsymbol{y}_j)$ for the two sets of constraints

$$\mathcal{L} = \sum_{i,j} d(\boldsymbol{x}_i, \boldsymbol{y}_j)\gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - \sum_i \alpha(\boldsymbol{x}_i)\left(\sum_j \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - p(\boldsymbol{x}_i)\right) - \sum_j \beta(\boldsymbol{y}_j)\left(\sum_i \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - q(\boldsymbol{y}_j)\right)$$

  subject to $\gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) \geq 0$

- Rearranging

$$\mathcal{L} = \sum_i \alpha(\boldsymbol{x}_i)p(\boldsymbol{x}_i) + \sum_j \beta(\boldsymbol{y}_j)q(\boldsymbol{y}_j) - \sum_{i,j} \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j)(\alpha(\boldsymbol{x}_i) + \beta(\boldsymbol{y}_j) - d(\boldsymbol{x}_i, \boldsymbol{y}_j))$$

- This is equivalent to maximising $\sum_i \alpha(\boldsymbol{x}_i)p(\boldsymbol{x}_i) + \sum_j \beta(\boldsymbol{y}_j)q(\boldsymbol{y}_j)$, subject to

$$\forall i, j \quad \alpha(\boldsymbol{x}_i) + \beta(\boldsymbol{y}_j) \leq d(\boldsymbol{x}_i, \boldsymbol{y}_j)$$

- For continuous distributions sums become integrals and $\alpha(\boldsymbol{x})$, $\beta(\boldsymbol{y})$ become functions

The Dual Problem

- With Lagrange multipliers $\alpha(\boldsymbol{x}_i)$ and $\beta(\boldsymbol{y}_j)$ for the two sets of constraints

$$\mathcal{L} = \sum_{i,j} d(\boldsymbol{x}_i, \boldsymbol{y}_j)\gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - \sum_i \alpha(\boldsymbol{x}_i)\left(\sum_j \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - p(\boldsymbol{x}_i)\right) - \sum_j \beta(\boldsymbol{y}_j)\left(\sum_i \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - q(\boldsymbol{y}_j)\right)$$

  subject to $\gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) \geq 0$

- Rearranging

$$\mathcal{L} = \sum_i \alpha(\boldsymbol{x}_i)p(\boldsymbol{x}_i) + \sum_j \beta(\boldsymbol{y}_j)q(\boldsymbol{y}_j) - \sum_{i,j} \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j)(\alpha(\boldsymbol{x}_i) + \beta(\boldsymbol{y}_j) - d(\boldsymbol{x}_i, \boldsymbol{y}_j))$$

- This is equivalent to maximising $\sum_i \alpha(\boldsymbol{x}_i)p(\boldsymbol{x}_i) + \sum_j \beta(\boldsymbol{y}_j)q(\boldsymbol{y}_j)$, subject to

$$\forall i, j \quad \alpha(\boldsymbol{x}_i) + \beta(\boldsymbol{y}_j) \leq d(\boldsymbol{x}_i, \boldsymbol{y}_j)$$

- For continuous distributions sums become integrals and $\alpha(\boldsymbol{x})$, $\beta(\boldsymbol{y})$ become functions

The Dual Problem

- With Lagrange multipliers $\alpha(\boldsymbol{x}_i)$ and $\beta(\boldsymbol{y}_j)$ for the two sets of constraints

$$\mathcal{L} = \sum_{i,j} d(\boldsymbol{x}_i, \boldsymbol{y}_j)\gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - \sum_i \alpha(\boldsymbol{x}_i)\left(\sum_j \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - p(\boldsymbol{x}_i)\right) - \sum_j \beta(\boldsymbol{y}_j)\left(\sum_i \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) - q(\boldsymbol{y}_j)\right)$$

  subject to $\gamma(\boldsymbol{x}_i, \boldsymbol{y}_j) \geq 0$

- Rearranging

$$\mathcal{L} = \sum_i \alpha(\boldsymbol{x}_i)p(\boldsymbol{x}_i) + \sum_j \beta(\boldsymbol{y}_j)q(\boldsymbol{y}_j) - \sum_{i,j} \gamma(\boldsymbol{x}_i, \boldsymbol{y}_j)(\alpha(\boldsymbol{x}_i) + \beta(\boldsymbol{y}_j) - d(\boldsymbol{x}_i, \boldsymbol{y}_j))$$

- This is equivalent to maximising $\sum_i \alpha(\boldsymbol{x}_i)p(\boldsymbol{x}_i) + \sum_j \beta(\boldsymbol{y}_j)q(\boldsymbol{y}_j)$, subject to

$$\forall i, j \quad \alpha(\boldsymbol{x}_i) + \beta(\boldsymbol{y}_j) \leq d(\boldsymbol{x}_i, \boldsymbol{y}_j)$$

- For continuous distributions sums become integrals and $\alpha(\boldsymbol{x})$, $\beta(\boldsymbol{y})$ become functions

Dual Form Constraint

- We need $\alpha(\boldsymbol{x}) + \beta(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y})$ for all $\boldsymbol{x}$, $\boldsymbol{y}$, and want β as large as possible
- Taking $\boldsymbol{x} = \boldsymbol{y}$ (where $d(\boldsymbol{y},\boldsymbol{y}) = 0$) gives $\beta(\boldsymbol{y}) \leq -\alpha(\boldsymbol{y})$; so $-\alpha$ is the largest β could possibly be
- Try that largest value, $\beta = -\alpha$: the constraint becomes

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) \quad \text{for all } \boldsymbol{x},\boldsymbol{y}$$

- For any α that satisfies this, $\beta = -\alpha$ is allowed and cannot be beaten
- Restricting to such α loses nothing (the **Kantorovich–Rubinstein** theorem; see the notes)

Dual Form Constraint

- We need $\alpha(\boldsymbol{x}) + \beta(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y})$ for all $\boldsymbol{x}$, $\boldsymbol{y}$, and want β as large as possible
- Taking $\boldsymbol{x} = \boldsymbol{y}$ (where $d(\boldsymbol{y},\boldsymbol{y}) = 0$) gives $\beta(\boldsymbol{y}) \leq -\alpha(\boldsymbol{y})$; so $-\alpha$ is the largest β could possibly be
- Try that largest value, $\beta = -\alpha$: the constraint becomes

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) \quad \text{for all } \boldsymbol{x},\boldsymbol{y}$$

- For any α that satisfies this, $\beta = -\alpha$ is allowed and cannot be beaten
- Restricting to such α loses nothing (the **Kantorovich–Rubinstein** theorem; see the notes)

Dual Form Constraint

- We need $\alpha(\boldsymbol{x}) + \beta(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y})$ for all $\boldsymbol{x}$, $\boldsymbol{y}$, and want β as large as possible
- Taking $\boldsymbol{x} = \boldsymbol{y}$ (where $d(\boldsymbol{y},\boldsymbol{y}) = 0$) gives $\beta(\boldsymbol{y}) \leq -\alpha(\boldsymbol{y})$; so $-\alpha$ is the largest β could possibly be
- Try that largest value, $\beta = -\alpha$: the constraint becomes

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) \quad \text{for all } \boldsymbol{x},\boldsymbol{y}$$

- For any α that satisfies this, $\beta = -\alpha$ is allowed and cannot be beaten
- Restricting to such α loses nothing (the **Kantorovich–Rubinstein** theorem; see the notes)

Dual Form Constraint

- We need $\alpha(\boldsymbol{x}) + \beta(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y})$ for all $\boldsymbol{x}$, $\boldsymbol{y}$, and want β as large as possible
- Taking $\boldsymbol{x} = \boldsymbol{y}$ (where $d(\boldsymbol{y},\boldsymbol{y}) = 0$) gives $\beta(\boldsymbol{y}) \leq -\alpha(\boldsymbol{y})$; so $-\alpha$ is the largest β could possibly be
- Try that largest value, $\beta = -\alpha$: the constraint becomes

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) \quad \text{for all } \boldsymbol{x},\boldsymbol{y}$$

- For any α that satisfies this, $\beta = -\alpha$ is allowed and cannot be beaten
- Restricting to such α loses nothing (the **Kantorovich–Rubinstein** theorem; see the notes)

Dual Form Constraint

- We need $\alpha(\boldsymbol{x}) + \beta(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y})$ for all $\boldsymbol{x}$, $\boldsymbol{y}$, and want β as large as possible
- Taking $\boldsymbol{x} = \boldsymbol{y}$ (where $d(\boldsymbol{y},\boldsymbol{y}) = 0$) gives $\beta(\boldsymbol{y}) \leq -\alpha(\boldsymbol{y})$; so $-\alpha$ is the largest β could possibly be
- Try that largest value, $\beta = -\alpha$: the constraint becomes

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) \quad \text{for all } \boldsymbol{x},\boldsymbol{y}$$

- For any α that satisfies this, $\beta = -\alpha$ is allowed and cannot be beaten
- Restricting to such α loses nothing (the **Kantorovich–Rubinstein** theorem; see the notes)

Dual Form Constraint

- We need $\alpha(\boldsymbol{x}) + \beta(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y})$ for all $\boldsymbol{x}$, $\boldsymbol{y}$, and want β as large as possible
- Taking $\boldsymbol{x} = \boldsymbol{y}$ (where $d(\boldsymbol{y},\boldsymbol{y}) = 0$) gives $\beta(\boldsymbol{y}) \leq -\alpha(\boldsymbol{y})$; so $-\alpha$ is the largest β could possibly be
- Try that largest value, $\beta = -\alpha$: the constraint becomes

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) \quad \text{for all } \boldsymbol{x},\boldsymbol{y}$$

- For any α that satisfies this, $\beta = -\alpha$ is allowed and cannot be beaten
- Restricting to such α loses nothing (the **Kantorovich–Rubinstein** theorem; see the notes)

Kantorovich–Rubinstein Duality

- Putting this together

$$W(p,q) = \max_{\alpha} \int \alpha(\boldsymbol{x}) \big(p(\boldsymbol{x}) - q(\boldsymbol{x})\big) \,\mathrm{d}\boldsymbol{x}$$
$$= \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{X})] - \mathbb{E}_q[\alpha(\boldsymbol{X})] \Big)$$

- where the maximum is over functions with

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x} - \boldsymbol{y}\|$$

  that is, α is **1-Lipschitz** (Lipschitz-1)

- Instead of the huge object $\gamma(\boldsymbol{x},\boldsymbol{y})$ we now need one function $\alpha(\boldsymbol{x})$, and the expectations can be estimated from samples

Kantorovich–Rubinstein Duality

- Putting this together

$$W(p,q) = \max_{\alpha} \int \alpha(\boldsymbol{x}) \big(p(\boldsymbol{x}) - q(\boldsymbol{x})\big) \,\mathrm{d}\boldsymbol{x}$$
$$= \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{X})] - \mathbb{E}_q[\alpha(\boldsymbol{X})] \Big)$$

- where the maximum is over functions with

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x} - \boldsymbol{y}\|$$

  that is, α is **1-Lipschitz** (Lipschitz-1)

- Instead of the huge object $\gamma(\boldsymbol{x},\boldsymbol{y})$ we now need one function $\alpha(\boldsymbol{x})$, and the expectations can be estimated from samples

Kantorovich–Rubinstein Duality

- Putting this together

$$W(p,q) = \max_{\alpha} \int \alpha(\boldsymbol{x}) \big(p(\boldsymbol{x}) - q(\boldsymbol{x})\big) \,\mathrm{d}\boldsymbol{x}$$
$$= \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{X})] - \mathbb{E}_q[\alpha(\boldsymbol{X})] \Big)$$

- where the maximum is over functions with

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x} - \boldsymbol{y}\|$$

  that is, α is **1-Lipschitz** (Lipschitz-1)

- Instead of the huge object $\gamma(\boldsymbol{x},\boldsymbol{y})$ we now need one function $\alpha(\boldsymbol{x})$, and the expectations can be estimated from samples

Kantorovich–Rubinstein Duality

- Putting this together

$$W(p,q) = \max_{\alpha} \int \alpha(\boldsymbol{x}) \big(p(\boldsymbol{x}) - q(\boldsymbol{x})\big) \,\mathrm{d}\boldsymbol{x}$$
$$= \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{X})] - \mathbb{E}_q[\alpha(\boldsymbol{X})] \Big)$$

- where the maximum is over functions with

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x} - \boldsymbol{y}\|$$

  that is, α is **1-Lipschitz** (Lipschitz-1)

- Instead of the huge object $\gamma(\boldsymbol{x},\boldsymbol{y})$ we now need one function $\alpha(\boldsymbol{x})$, and the expectations can be estimated from samples

Kantorovich–Rubinstein Duality

- Putting this together

$$W(p,q) = \max_{\alpha} \int \alpha(\boldsymbol{x}) \big(p(\boldsymbol{x}) - q(\boldsymbol{x})\big) \,\mathrm{d}\boldsymbol{x}$$
$$= \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{X})] - \mathbb{E}_q[\alpha(\boldsymbol{X})] \Big)$$

- where the maximum is over functions with

$$\alpha(\boldsymbol{x}) - \alpha(\boldsymbol{y}) \leq d(\boldsymbol{x},\boldsymbol{y}) = \|\boldsymbol{x} - \boldsymbol{y}\|$$

  that is, α is **1-Lipschitz** (Lipschitz-1)

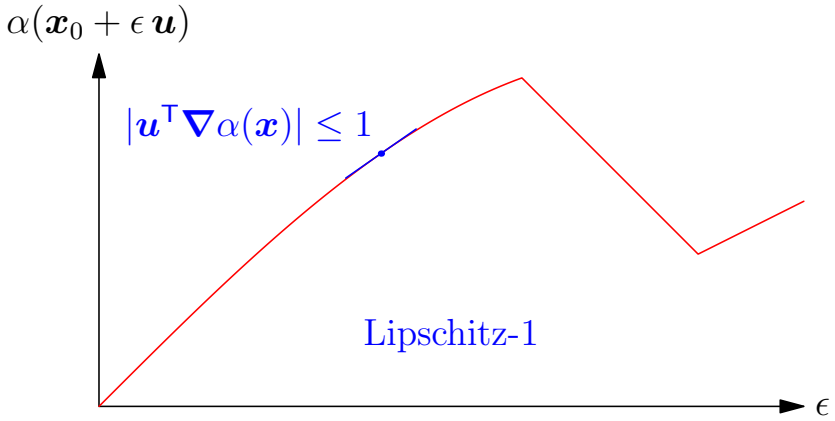
- Instead of the huge object $\gamma(\boldsymbol{x},\boldsymbol{y})$ we now need one function $\alpha(\boldsymbol{x})$, and the expectations can be estimated from samples

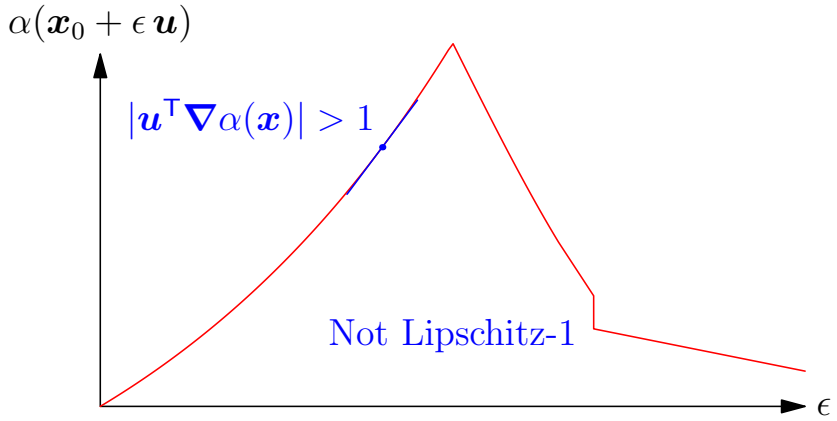
Lipschitz-1 Functions

- We note for a Lipschitz-1 function and any unit vector $\boldsymbol{u}$

$$\boldsymbol{u}^\mathsf{T}\boldsymbol{\nabla}\alpha(\boldsymbol{x}) = \lim_{\epsilon\to 0}\frac{\alpha(\boldsymbol{x}+\epsilon\boldsymbol{u})-\alpha(\boldsymbol{x})}{\epsilon}$$

- That is, at every point the gradient in every direction is at most 1



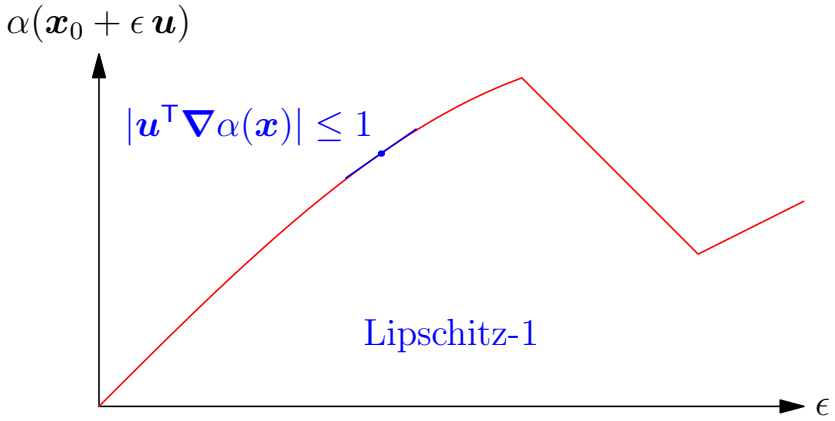


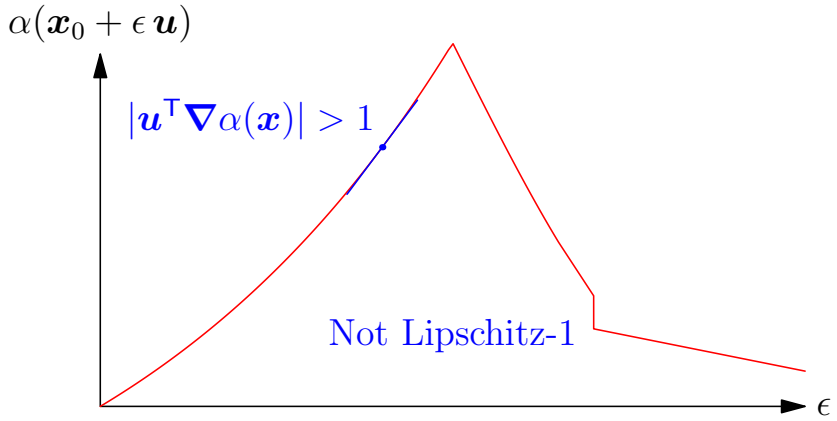
Lipschitz-1 Functions

- We note for a Lipschitz-1 function and any unit vector $\boldsymbol{u}$

$$\boldsymbol{u}^\mathsf{T}\boldsymbol{\nabla}\alpha(\boldsymbol{x}) = \lim_{\epsilon\to 0}\frac{\alpha(\boldsymbol{x}+\epsilon\boldsymbol{u})-\alpha(\boldsymbol{x})}{\epsilon} \leq \frac{\epsilon\|\boldsymbol{u}\|}{\epsilon}$$

- That is, at every point the gradient in every direction is at most 1



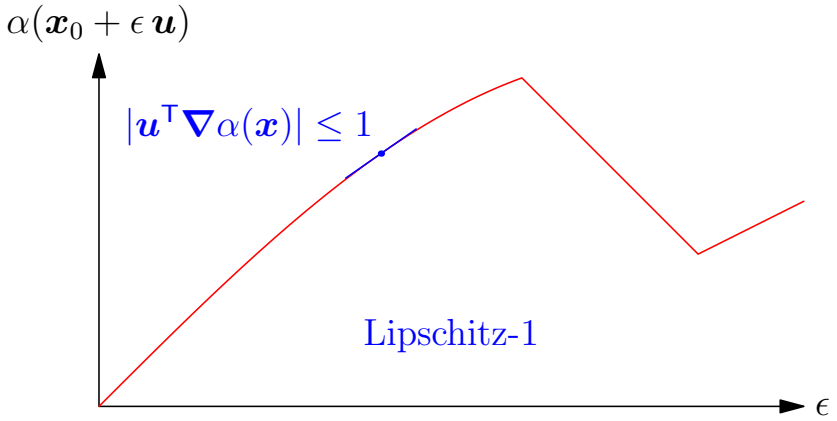


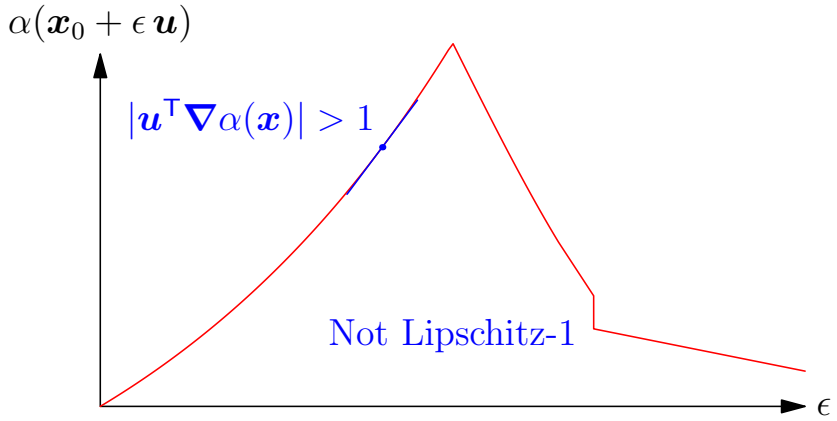
Lipschitz-1 Functions

- We note for a Lipschitz-1 function and any unit vector $\boldsymbol{u}$

$$\boldsymbol{u}^\mathsf{T}\boldsymbol{\nabla}\alpha(\boldsymbol{x}) = \lim_{\epsilon\to 0}\frac{\alpha(\boldsymbol{x}+\epsilon\boldsymbol{u})-\alpha(\boldsymbol{x})}{\epsilon} \leq \frac{\epsilon\|\boldsymbol{u}\|}{\epsilon} = \frac{\epsilon}{\epsilon} = 1$$

- That is, at every point the gradient in every direction is at most 1



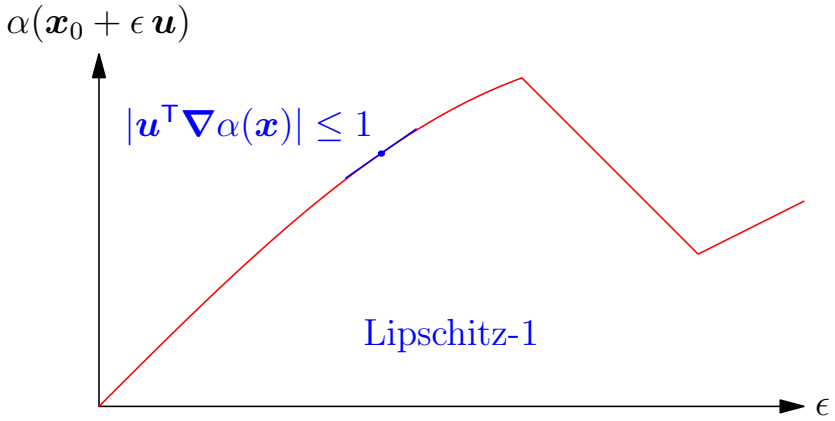


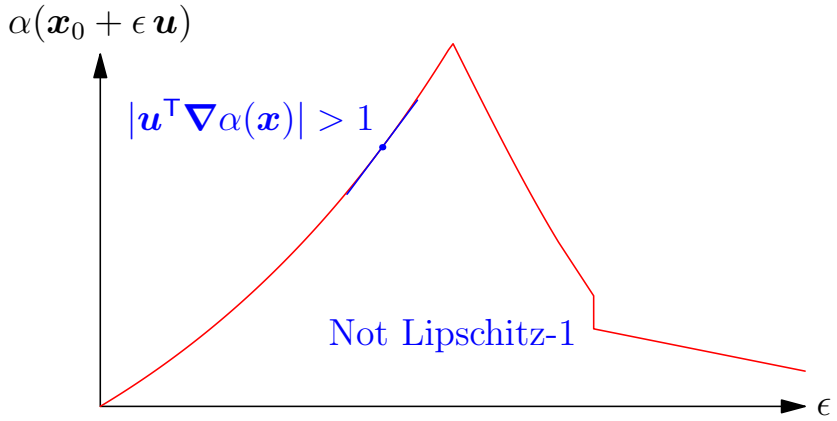
Lipschitz-1 Functions

- We note for a Lipschitz-1 function and any unit vector $\boldsymbol{u}$

$$\boldsymbol{u}^\mathsf{T}\boldsymbol{\nabla}\alpha(\boldsymbol{x}) = \lim_{\epsilon\to 0}\frac{\alpha(\boldsymbol{x}+\epsilon\boldsymbol{u})-\alpha(\boldsymbol{x})}{\epsilon} \leq \frac{\epsilon\|\boldsymbol{u}\|}{\epsilon} = \frac{\epsilon}{\epsilon} = 1$$

- That is, at every point the gradient in every direction is at most 1



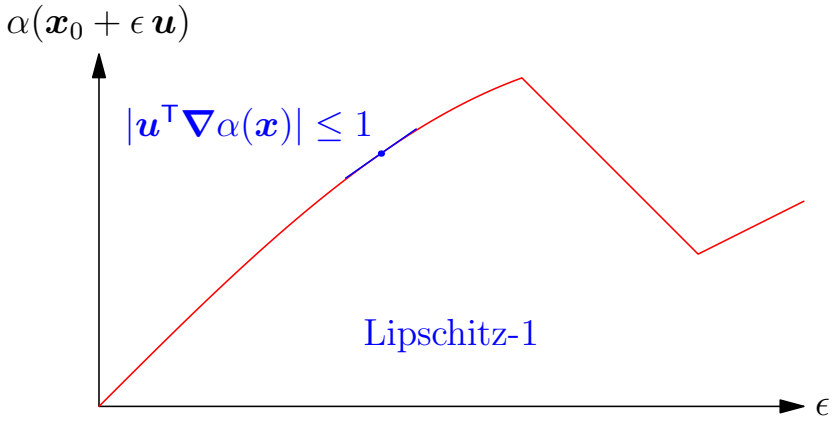


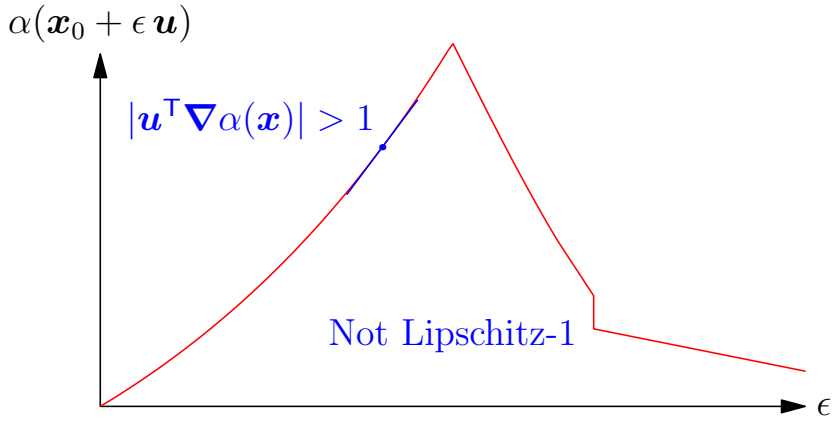
Lipschitz-1 Functions

- We note for a Lipschitz-1 function and any unit vector $\boldsymbol{u}$

$$\boldsymbol{u}^\mathsf{T}\boldsymbol{\nabla}\alpha(\boldsymbol{x}) = \lim_{\epsilon\to 0}\frac{\alpha(\boldsymbol{x}+\epsilon\boldsymbol{u})-\alpha(\boldsymbol{x})}{\epsilon} \leq \frac{\epsilon\|\boldsymbol{u}\|}{\epsilon} = \frac{\epsilon}{\epsilon} = 1$$

- That is, at every point the gradient in every direction is at most 1 (since the gradient defines the direction of greatest increase it is both necessary and sufficient for $\|\boldsymbol{\nabla}\alpha(\boldsymbol{x})\| \leq 1$ everywhere)



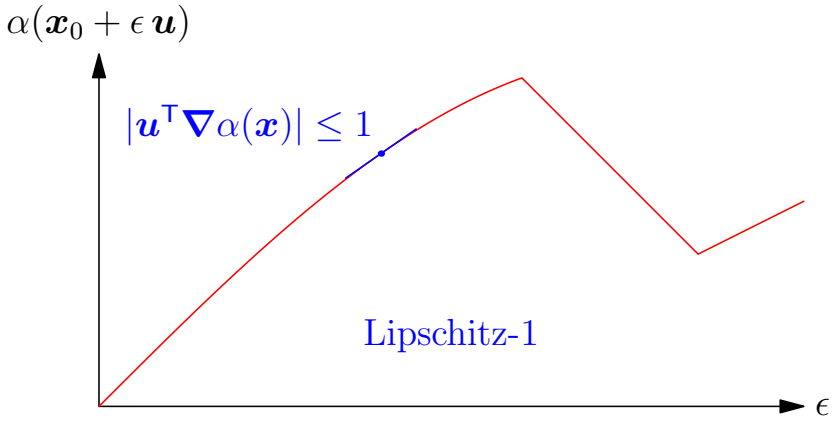


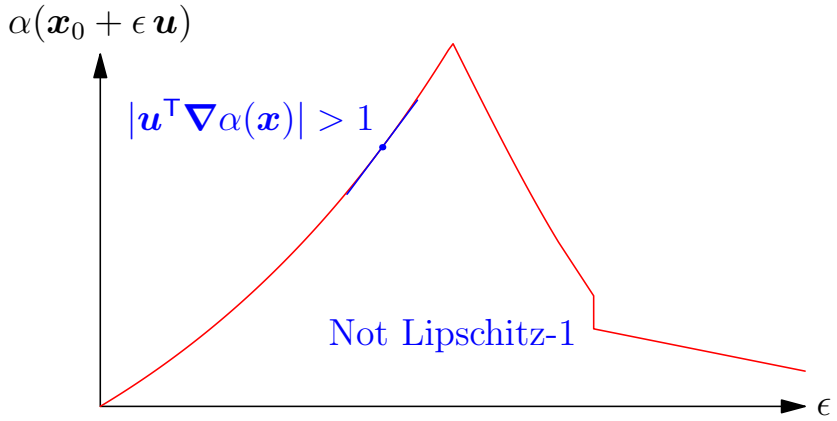
Lipschitz-1 Functions

- We note for a Lipschitz-1 function and any unit vector $\boldsymbol{u}$

$$\boldsymbol{u}^\mathsf{T}\boldsymbol{\nabla}\alpha(\boldsymbol{x}) = \lim_{\epsilon\to 0}\frac{\alpha(\boldsymbol{x}+\epsilon\boldsymbol{u})-\alpha(\boldsymbol{x})}{\epsilon} \leq \frac{\epsilon\|\boldsymbol{u}\|}{\epsilon} = \frac{\epsilon}{\epsilon} = 1$$

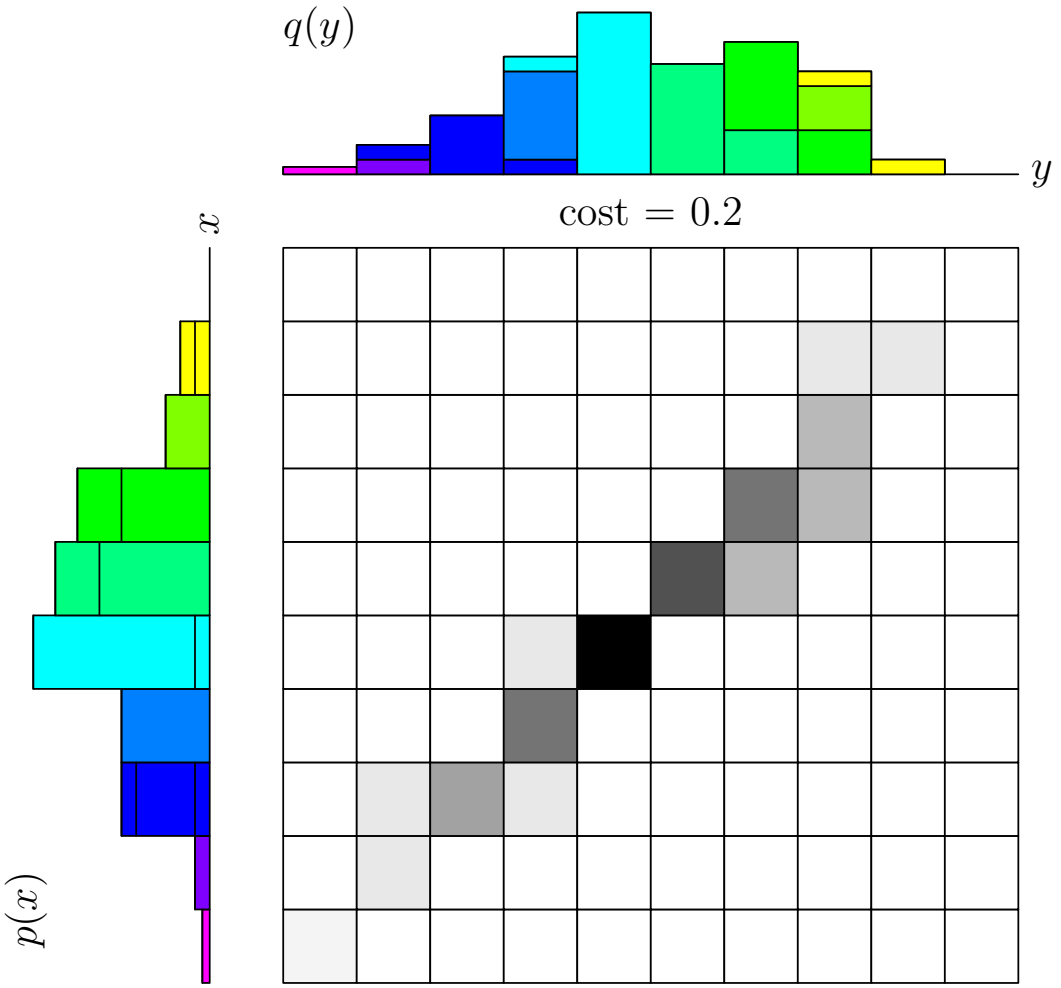
- That is, at every point the gradient in every direction is at most 1 (since the gradient defines the direction of greatest increase it is both necessary and sufficient for $\|\boldsymbol{\nabla}\alpha(\boldsymbol{x})\| \leq 1$ everywhere)



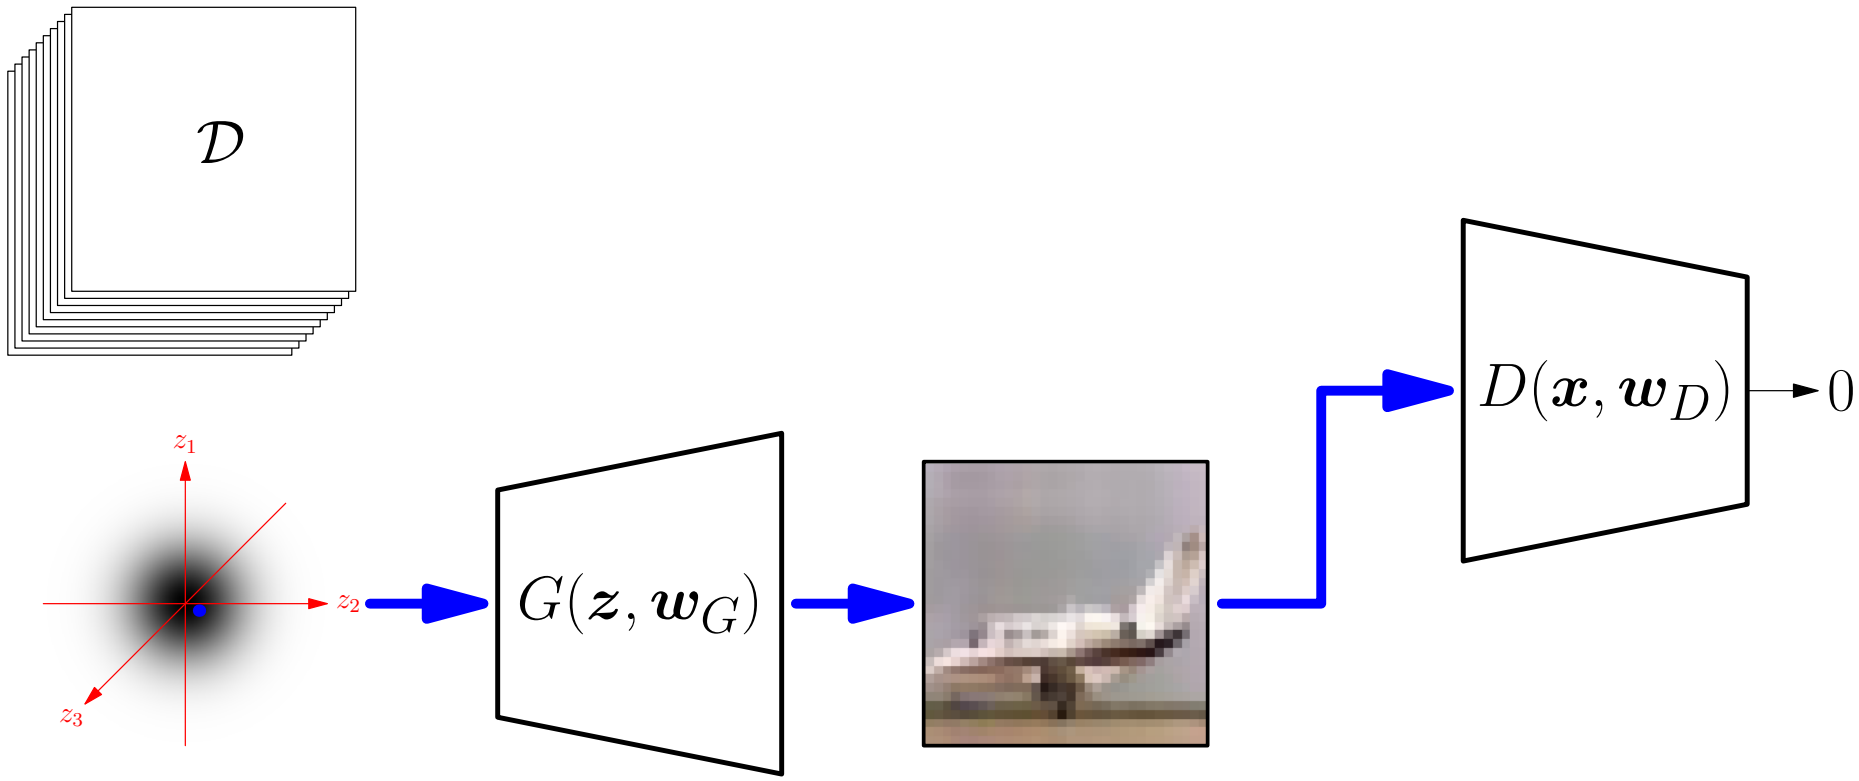


Outline

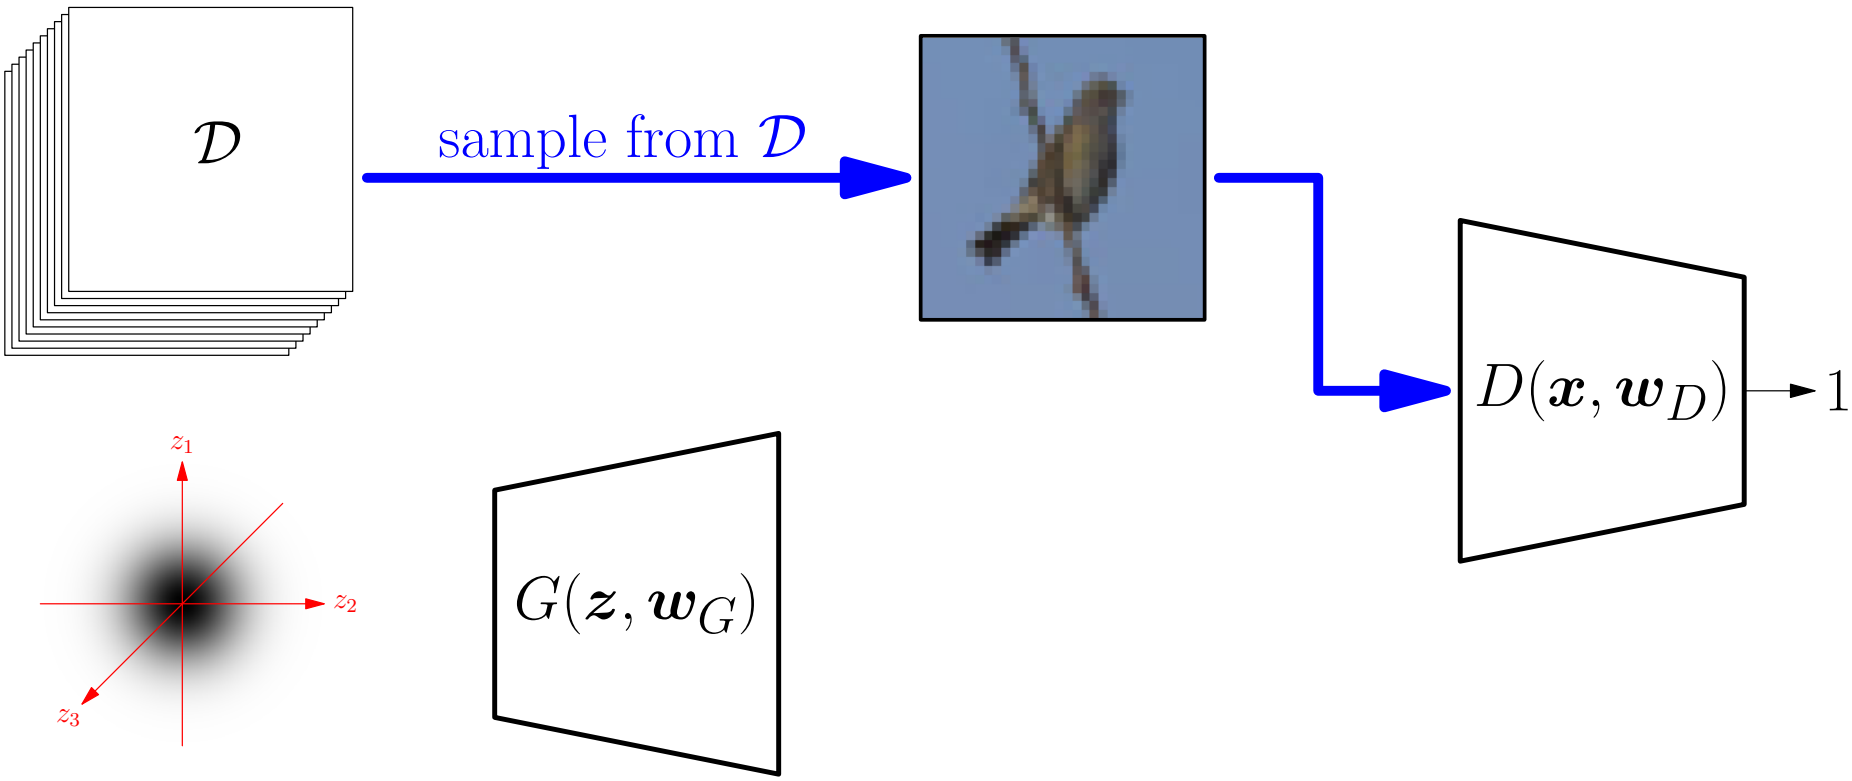
1. Comparing Distributions
2. Optimal Transport
3. Duality
4. **Applications**



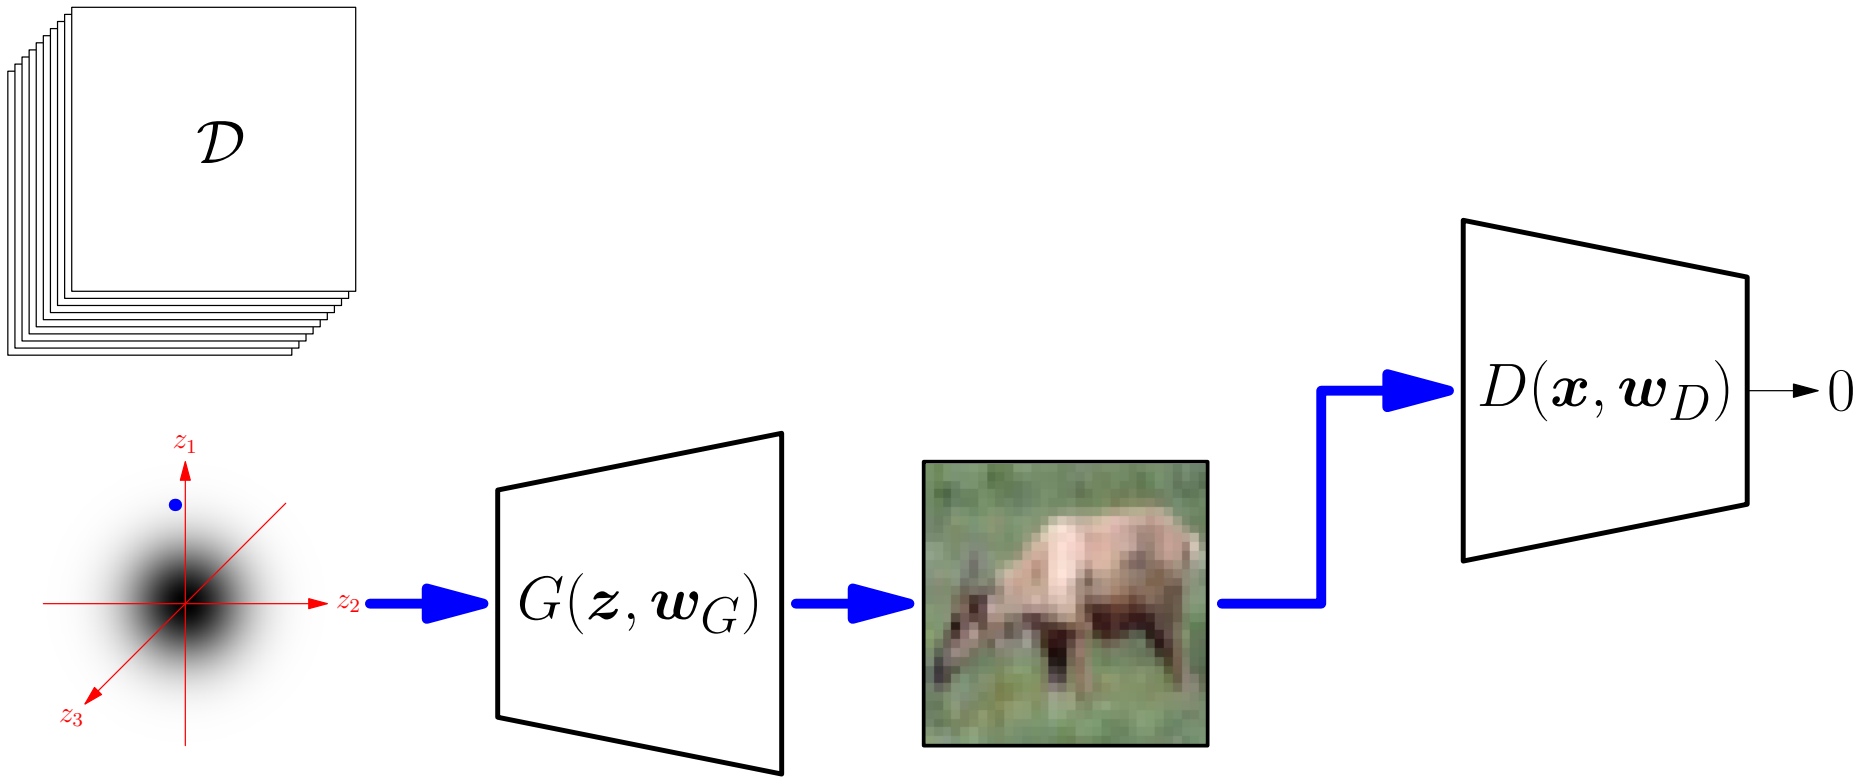
How GANs Work



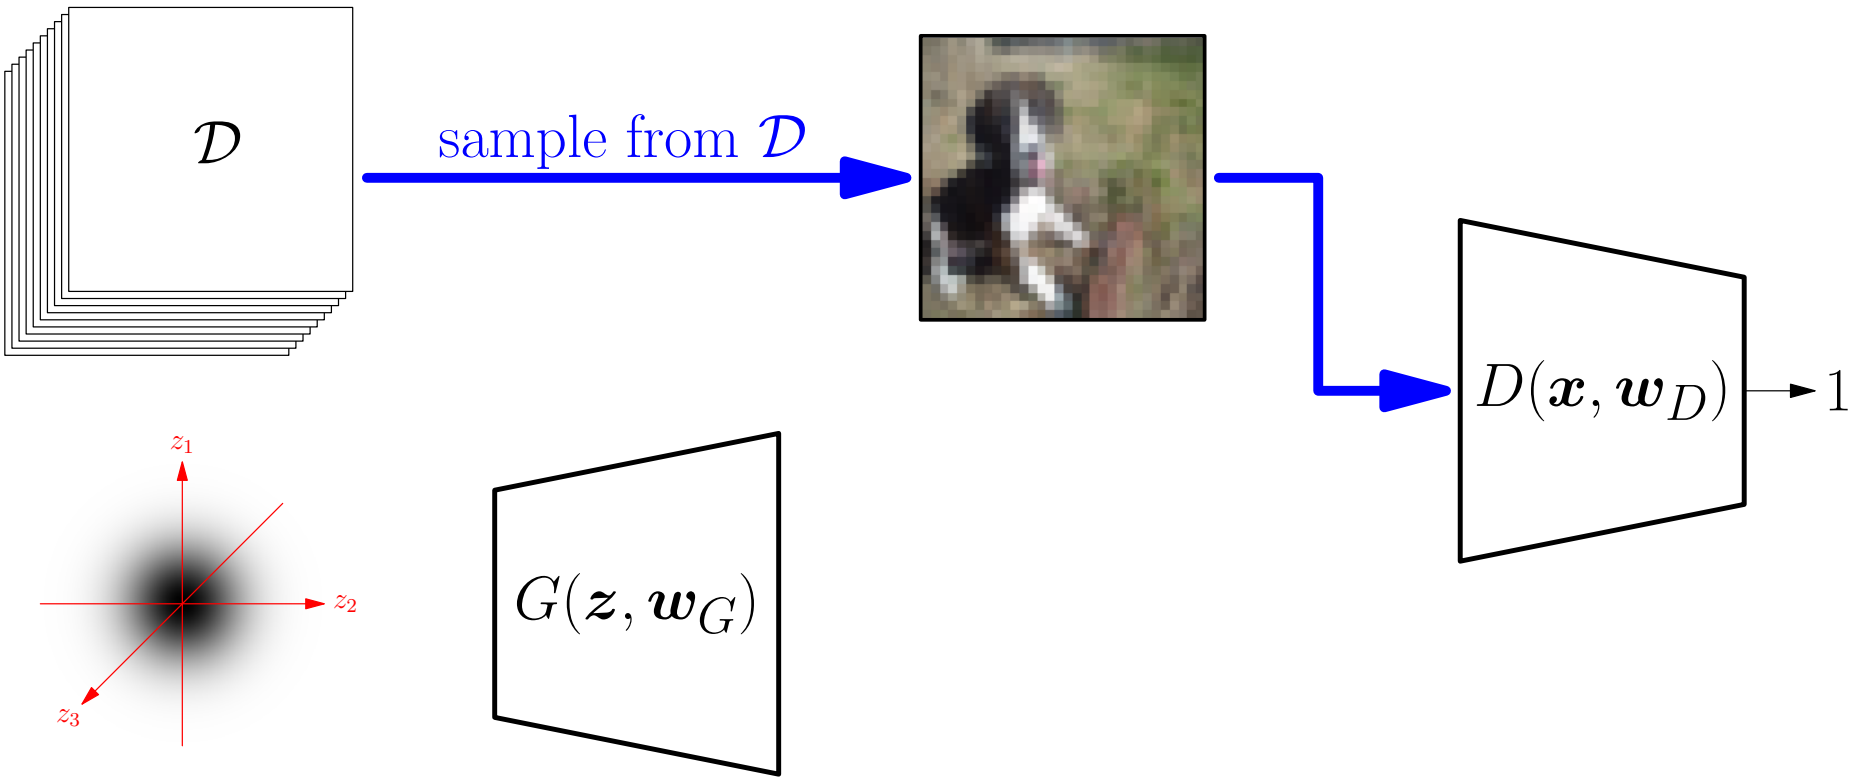
How GANs Work



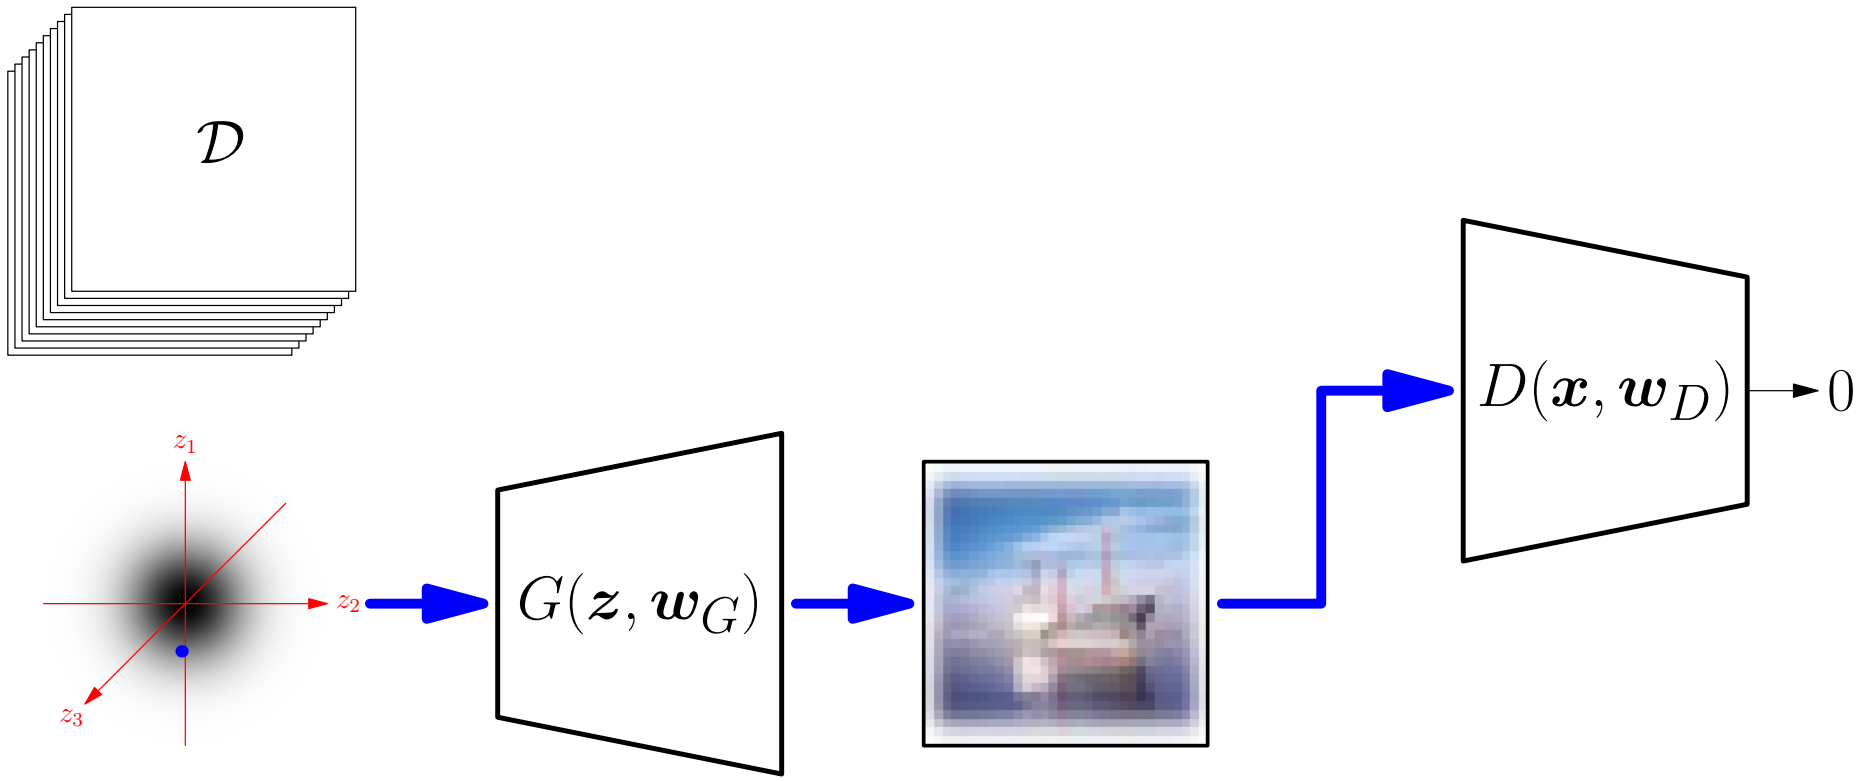
How GANs Work



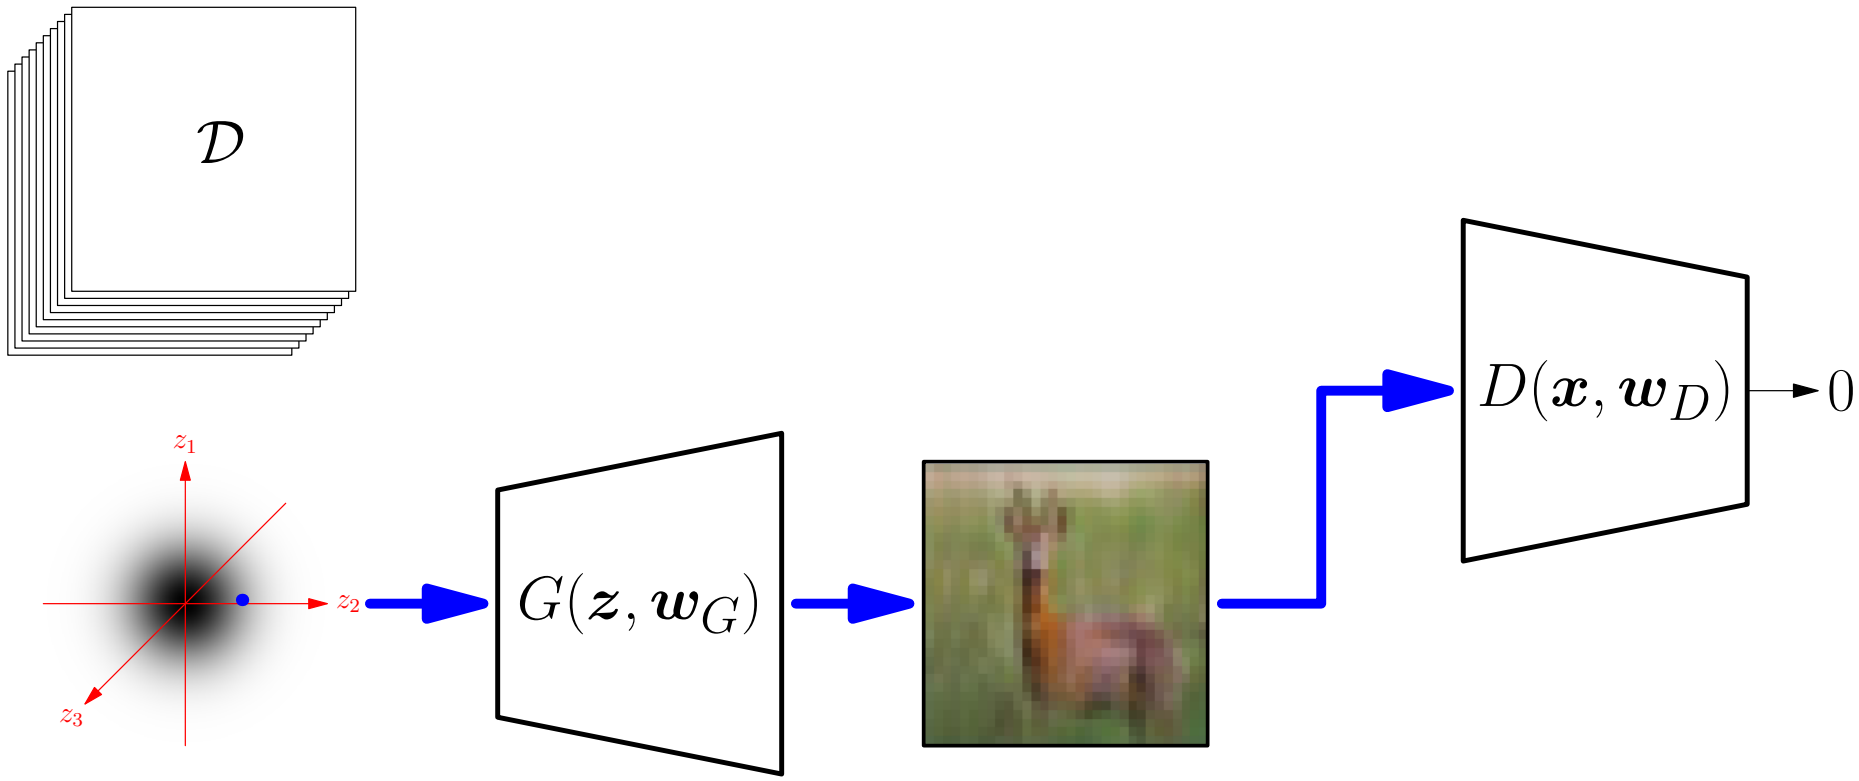
How GANs Work



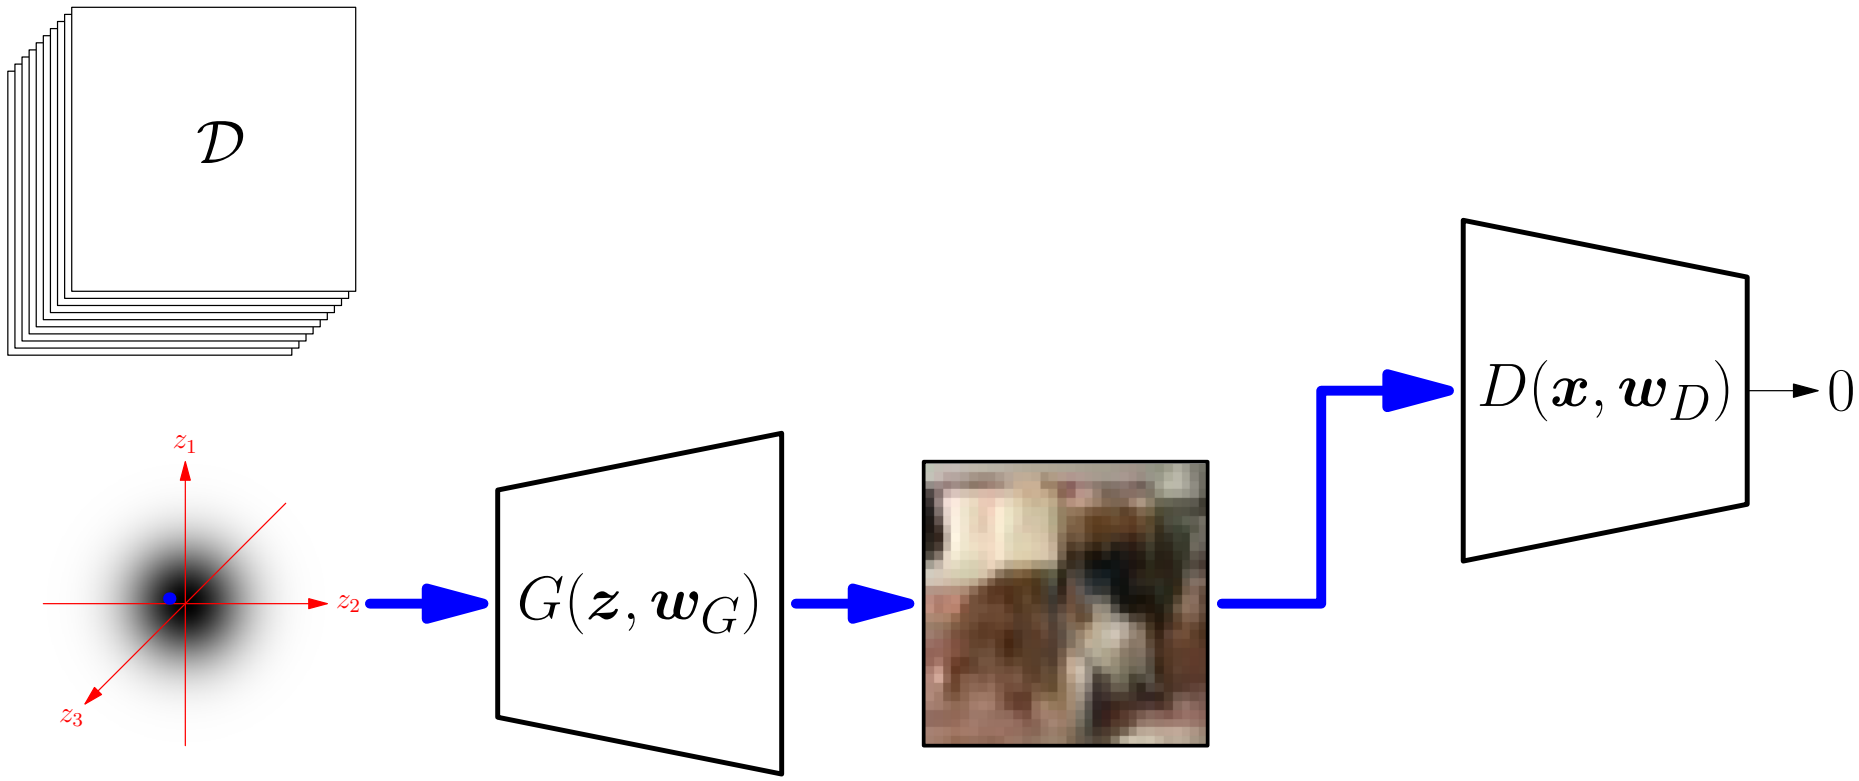
How GANs Work



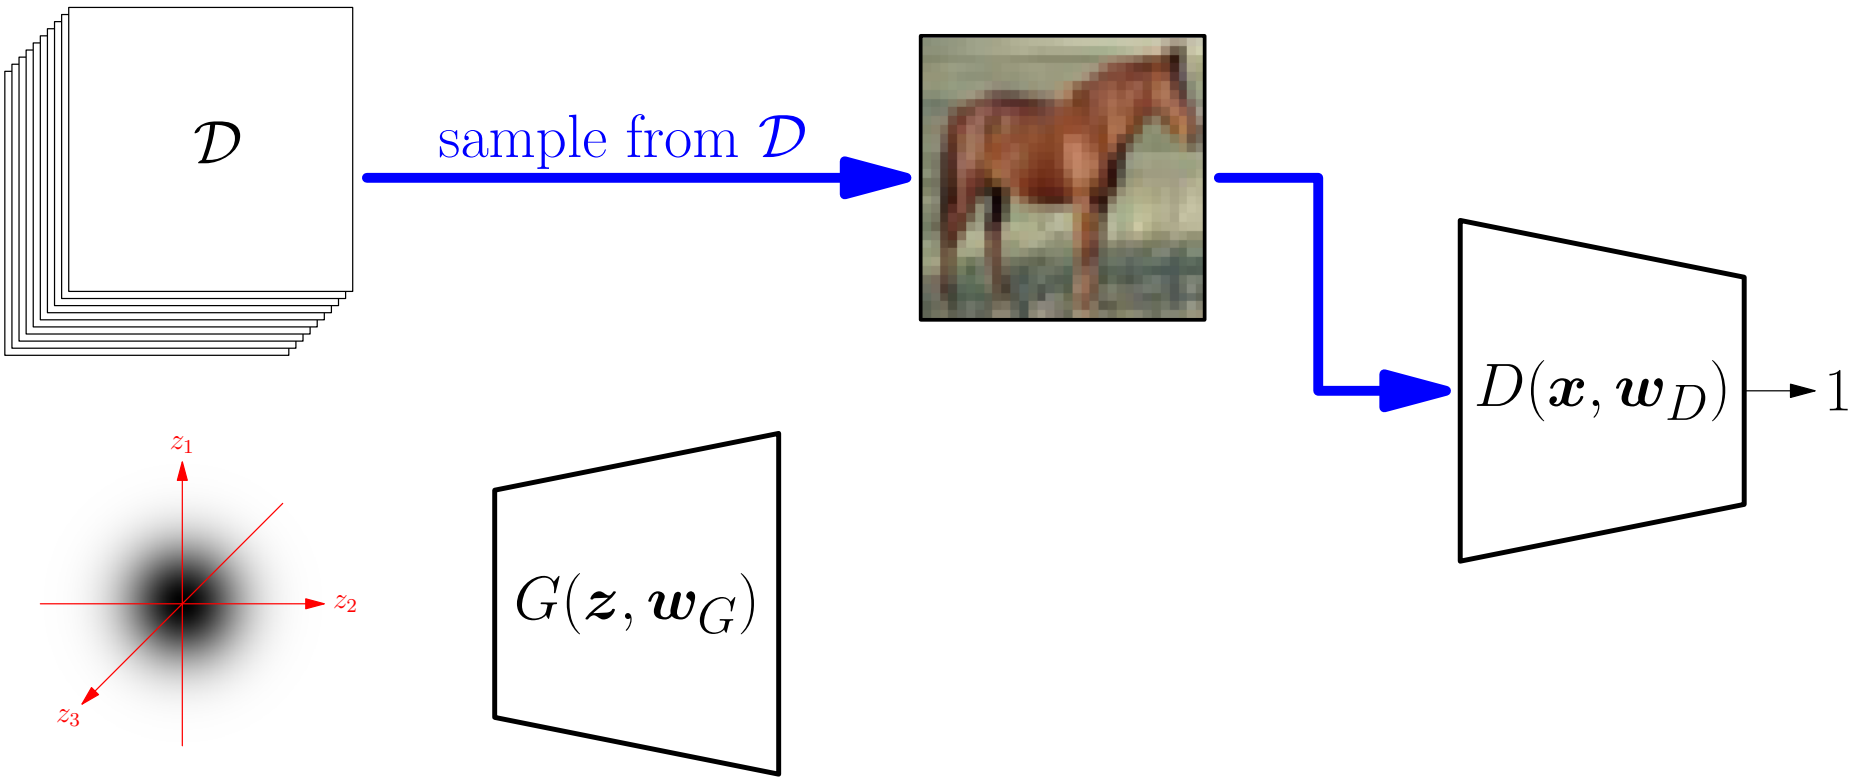
How GANs Work



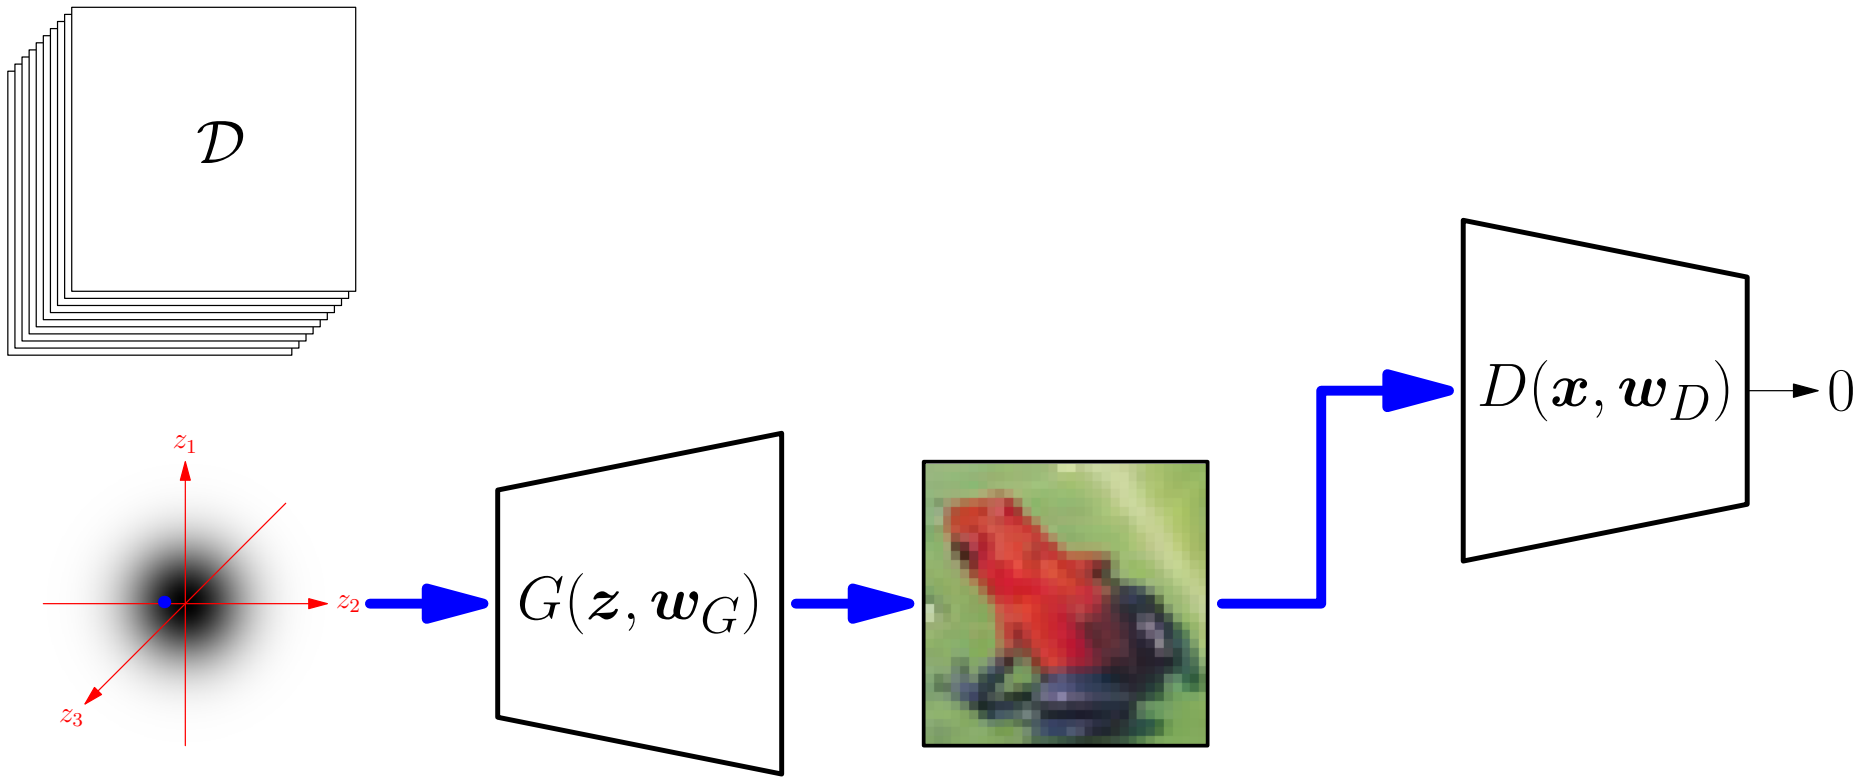
How GANs Work



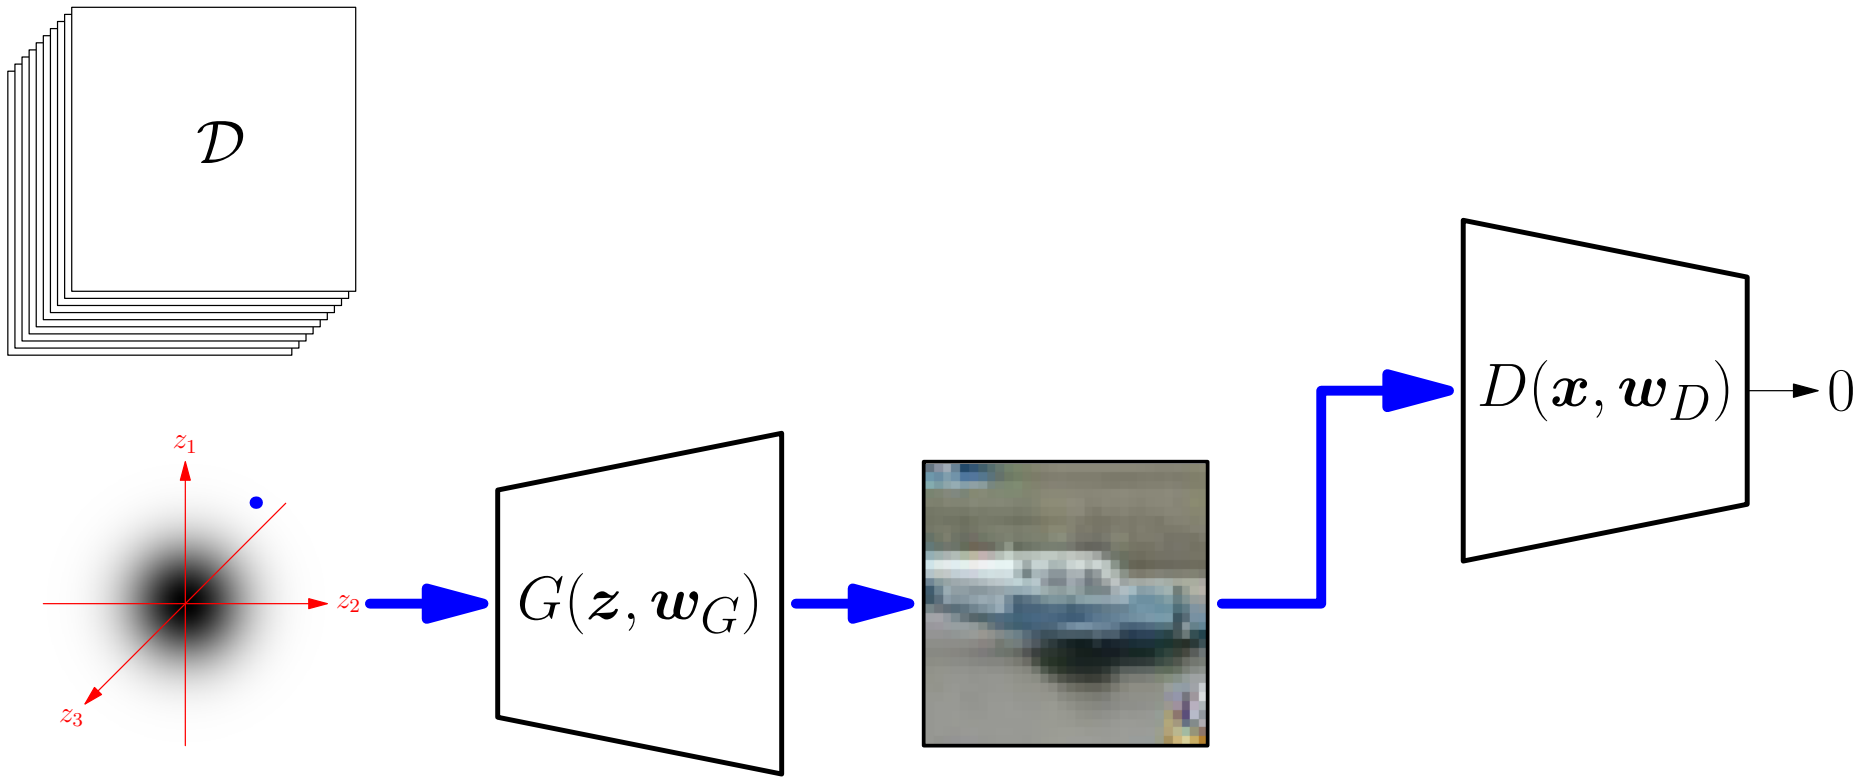
How GANs Work



How GANs Work



How GANs Work



Wasserstein GANs

- A GAN has a **generator** $G(\boldsymbol{z},\boldsymbol{w}_G)$ that turns noise $\boldsymbol{z} \sim \mathcal{N}(\mathbf{0},\mathbf{I})$ into images, and a **discriminator** trained to tell them from real images
- GANs are hard to train: a discriminator that is too good gives the generator no useful gradient
- A Wasserstein GAN instead trains the generator to minimise W between the data distribution p and the generator's distribution q

$$\min_{\boldsymbol{w}_G} \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{x})] - \mathbb{E}_{\boldsymbol{z}}[\alpha(G(\boldsymbol{z},\boldsymbol{w}_G))] \Big)$$

- The **critic** α is a second network in place of the discriminator; it still gives a gradient when the distributions do not overlap

Wasserstein GANs

- A GAN has a **generator** $G(\boldsymbol{z}, \boldsymbol{w}_G)$ that turns noise $\boldsymbol{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ into images, and a **discriminator** trained to tell them from real images
- GANs are hard to train: a discriminator that is too good gives the generator no useful gradient
- A Wasserstein GAN instead trains the generator to minimise W between the data distribution p and the generator's distribution q

$$\min_{\boldsymbol{w}_G} \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{x})] - \mathbb{E}_{\boldsymbol{z}}[\alpha(G(\boldsymbol{z}, \boldsymbol{w}_G))] \Big)$$

- The **critic** α is a second network in place of the discriminator; it still gives a gradient when the distributions do not overlap

Wasserstein GANs

- A GAN has a **generator** $G(\boldsymbol{z},\boldsymbol{w}_G)$ that turns noise $\boldsymbol{z} \sim \mathcal{N}(\mathbf{0},\mathbf{I})$ into images, and a **discriminator** trained to tell them from real images
- GANs are hard to train: a discriminator that is too good gives the generator no useful gradient
- A Wasserstein GAN instead trains the generator to minimise W between the data distribution p and the generator's distribution q

$$\min_{\boldsymbol{w}_G} \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{x})] - \mathbb{E}_{\boldsymbol{z}}[\alpha(G(\boldsymbol{z},\boldsymbol{w}_G))] \Big)$$

- The **critic** α is a second network in place of the discriminator; it still gives a gradient when the distributions do not overlap

Wasserstein GANs

- A GAN has a **generator** $G(\boldsymbol{z}, \boldsymbol{w}_G)$ that turns noise $\boldsymbol{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ into images, and a **discriminator** trained to tell them from real images
- GANs are hard to train: a discriminator that is too good gives the generator no useful gradient
- A Wasserstein GAN instead trains the generator to minimise W between the data distribution p and the generator's distribution q

$$\min_{\boldsymbol{w}_G} \max_{\alpha} \Big(\mathbb{E}_p[\alpha(\boldsymbol{x})] - \mathbb{E}_{\boldsymbol{z}}[\alpha(G(\boldsymbol{z}, \boldsymbol{w}_G))] \Big)$$

- The **critic** α is a second network in place of the discriminator; it still gives a gradient when the distributions do not overlap

The Critic

- Estimate the expectations from a minibatch $\mathcal{B}$ of real images and n generated ones

$$C = \frac{1}{|\mathcal{B}|}\sum_{\boldsymbol{x}\in\mathcal{B}} \alpha(\boldsymbol{x},\boldsymbol{w}_\alpha) - \frac{1}{n}\sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i,\boldsymbol{w}_G),\boldsymbol{w}_\alpha)$$

- The critic's weights $\boldsymbol{w}_\alpha$ are trained to maximise C; the generator's weights $\boldsymbol{w}_G$ to minimise it

- The critic should be 1-Lipschitz, which in practice is only roughly enforced: by clipping the weights (the original WGAN), penalising $\|\boldsymbol{\nabla}\alpha\|$ away from 1, or scaling each weight matrix to spectral norm 1

The Critic

- Estimate the expectations from a minibatch $\mathcal{B}$ of real images and n generated ones

$$C = \frac{1}{|\mathcal{B}|}\sum_{\boldsymbol{x}\in\mathcal{B}} \alpha(\boldsymbol{x},\boldsymbol{w}_\alpha) - \frac{1}{n}\sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i,\boldsymbol{w}_G),\boldsymbol{w}_\alpha)$$

- The critic's weights $\boldsymbol{w}_\alpha$ are trained to maximise C; the generator's weights $\boldsymbol{w}_G$ to minimise it

- The critic should be 1-Lipschitz, which in practice is only roughly enforced: by clipping the weights (the original WGAN), penalising $\|\boldsymbol{\nabla}\alpha\|$ away from 1, or scaling each weight matrix to spectral norm 1

The Critic

- Estimate the expectations from a minibatch $\mathcal{B}$ of real images and n generated ones

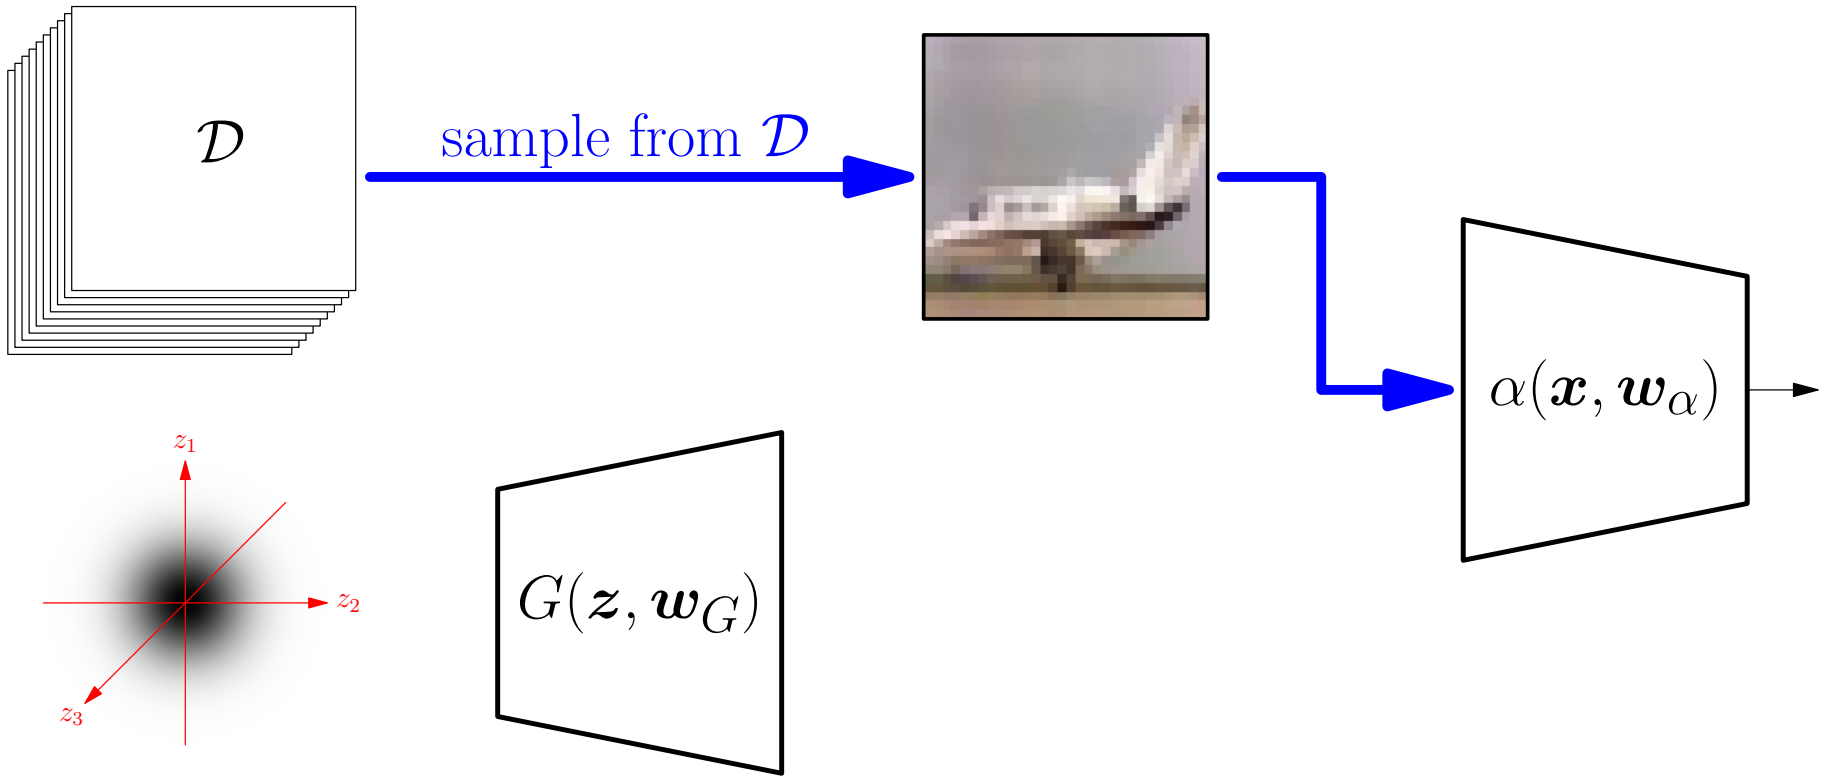
$$C = \frac{1}{|\mathcal{B}|}\sum_{\boldsymbol{x}\in\mathcal{B}} \alpha(\boldsymbol{x},\boldsymbol{w}_\alpha) - \frac{1}{n}\sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i,\boldsymbol{w}_G),\boldsymbol{w}_\alpha)$$

- The critic's weights $\boldsymbol{w}_\alpha$ are trained to maximise C; the generator's weights $\boldsymbol{w}_G$ to minimise it

- The critic should be 1-Lipschitz, which in practice is only roughly enforced: by clipping the weights (the original WGAN), penalising $\|\boldsymbol{\nabla}\alpha\|$ away from 1, or scaling each weight matrix to spectral norm 1

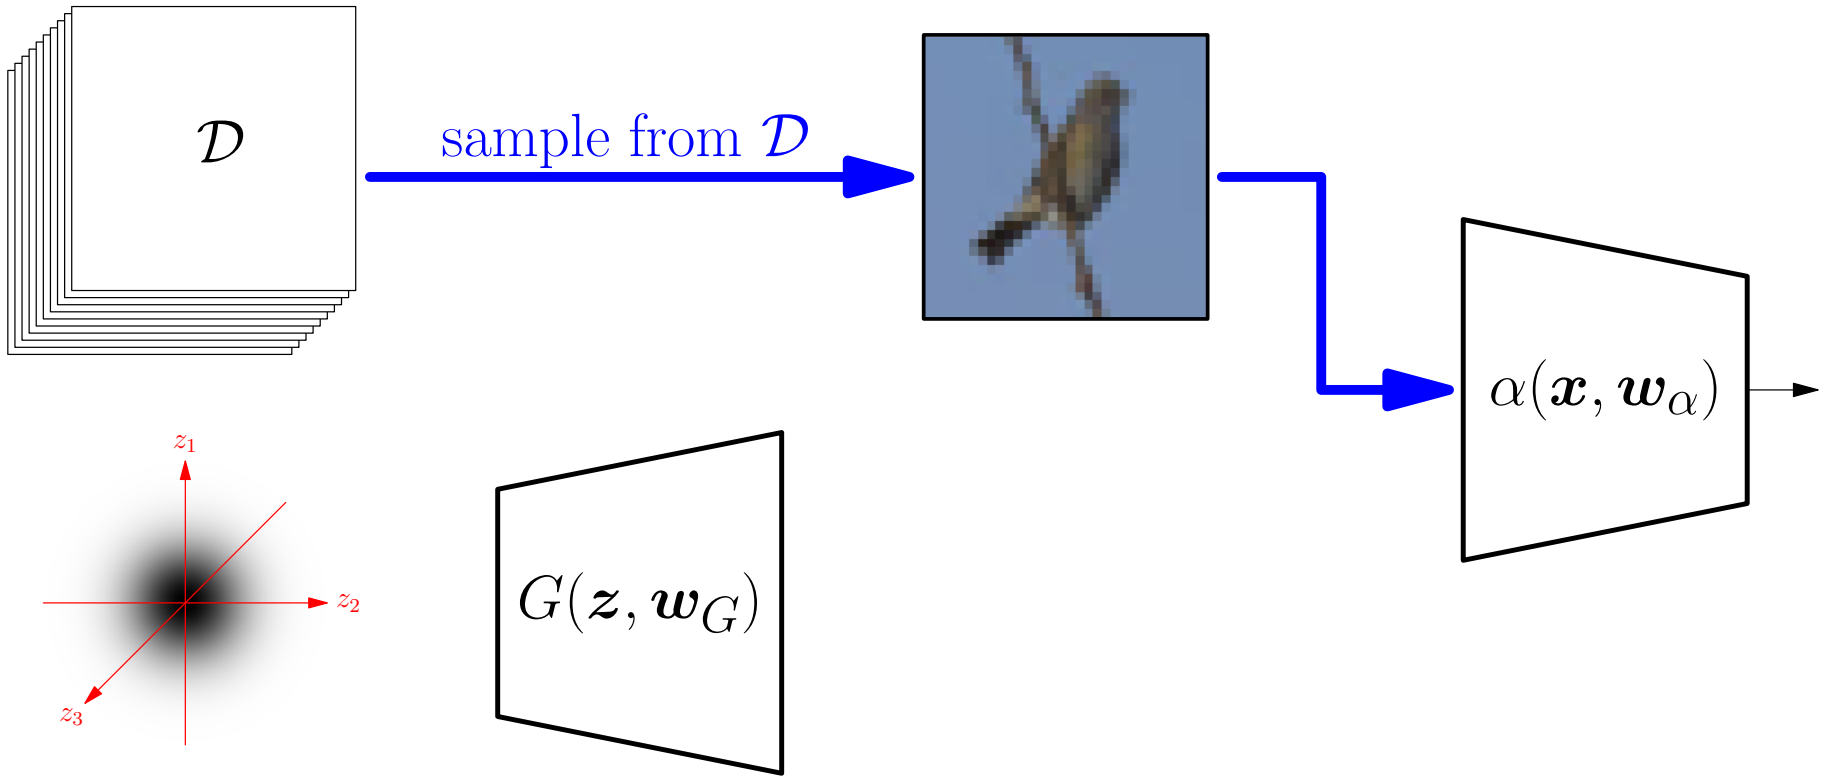
Wasserstein GANs

$$\max_{\boldsymbol{w}_\alpha} \min_{\boldsymbol{w}_G} \frac{1}{|\mathcal{B}|} \sum_{\boldsymbol{x} \in \mathcal{B}} \alpha(\boldsymbol{x}, \boldsymbol{w}_\alpha) - \frac{1}{n} \sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i, \boldsymbol{w}_G), \boldsymbol{w}_\alpha)$$



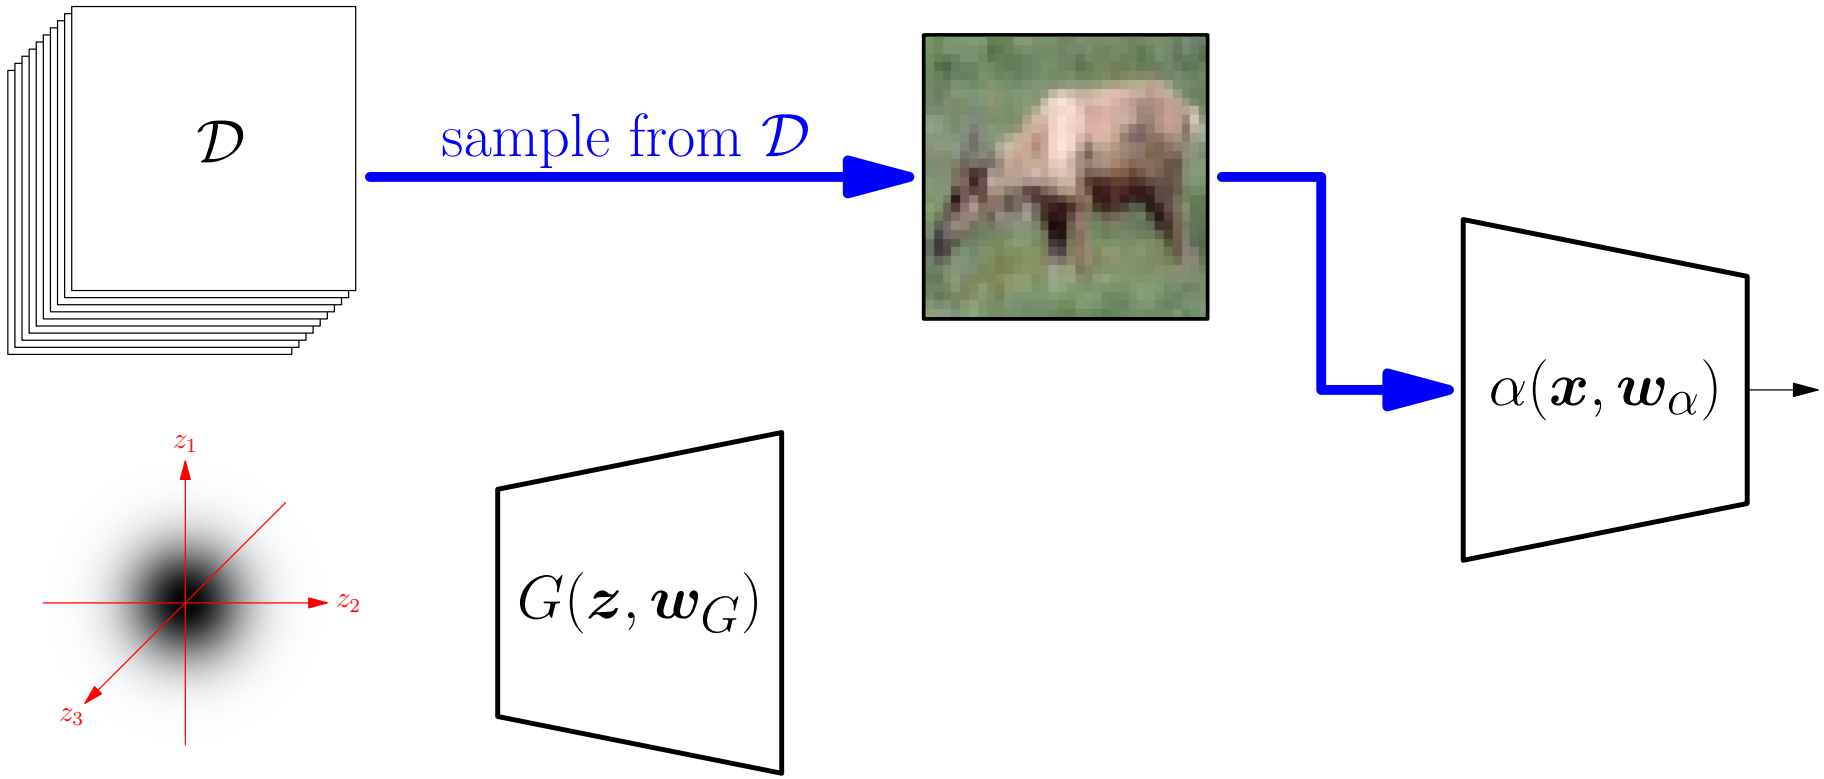
Wasserstein GANs

$$\max_{\boldsymbol{w}_\alpha} \min_{\boldsymbol{w}_G} \frac{1}{|\mathcal{B}|} \sum_{\boldsymbol{x} \in \mathcal{B}} \alpha(\boldsymbol{x}, \boldsymbol{w}_\alpha) - \frac{1}{n} \sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i, \boldsymbol{w}_G), \boldsymbol{w}_\alpha)$$



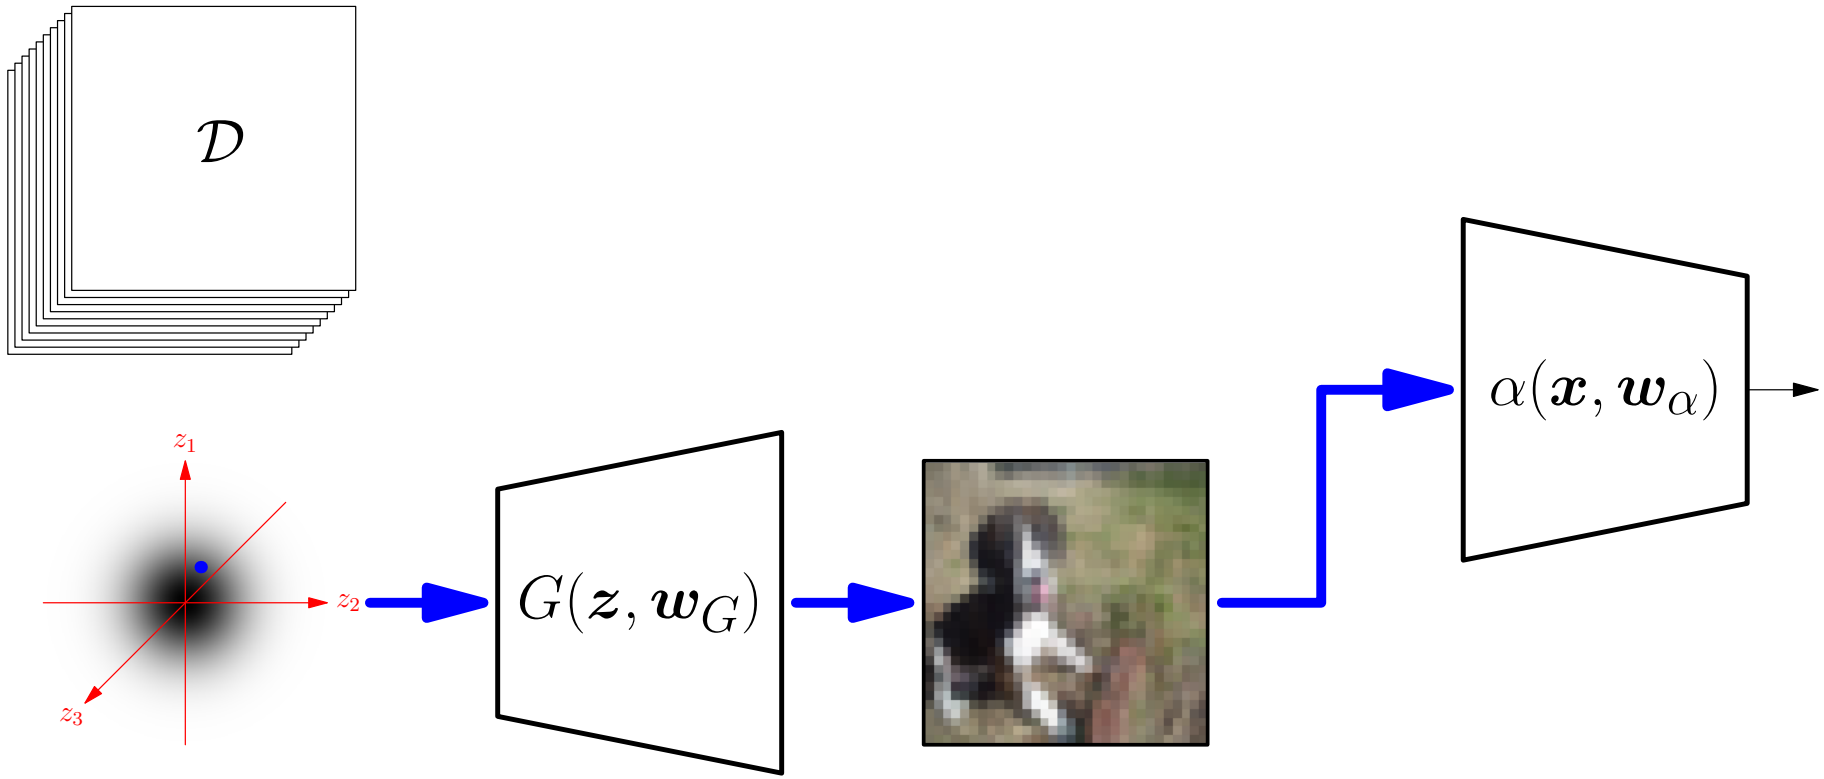
Wasserstein GANs

$$\max_{\boldsymbol{w}_\alpha} \min_{\boldsymbol{w}_G} \frac{1}{|\mathcal{B}|} \sum_{\boldsymbol{x} \in \mathcal{B}} \alpha(\boldsymbol{x}, \boldsymbol{w}_\alpha) - \frac{1}{n} \sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i, \boldsymbol{w}_G), \boldsymbol{w}_\alpha)$$



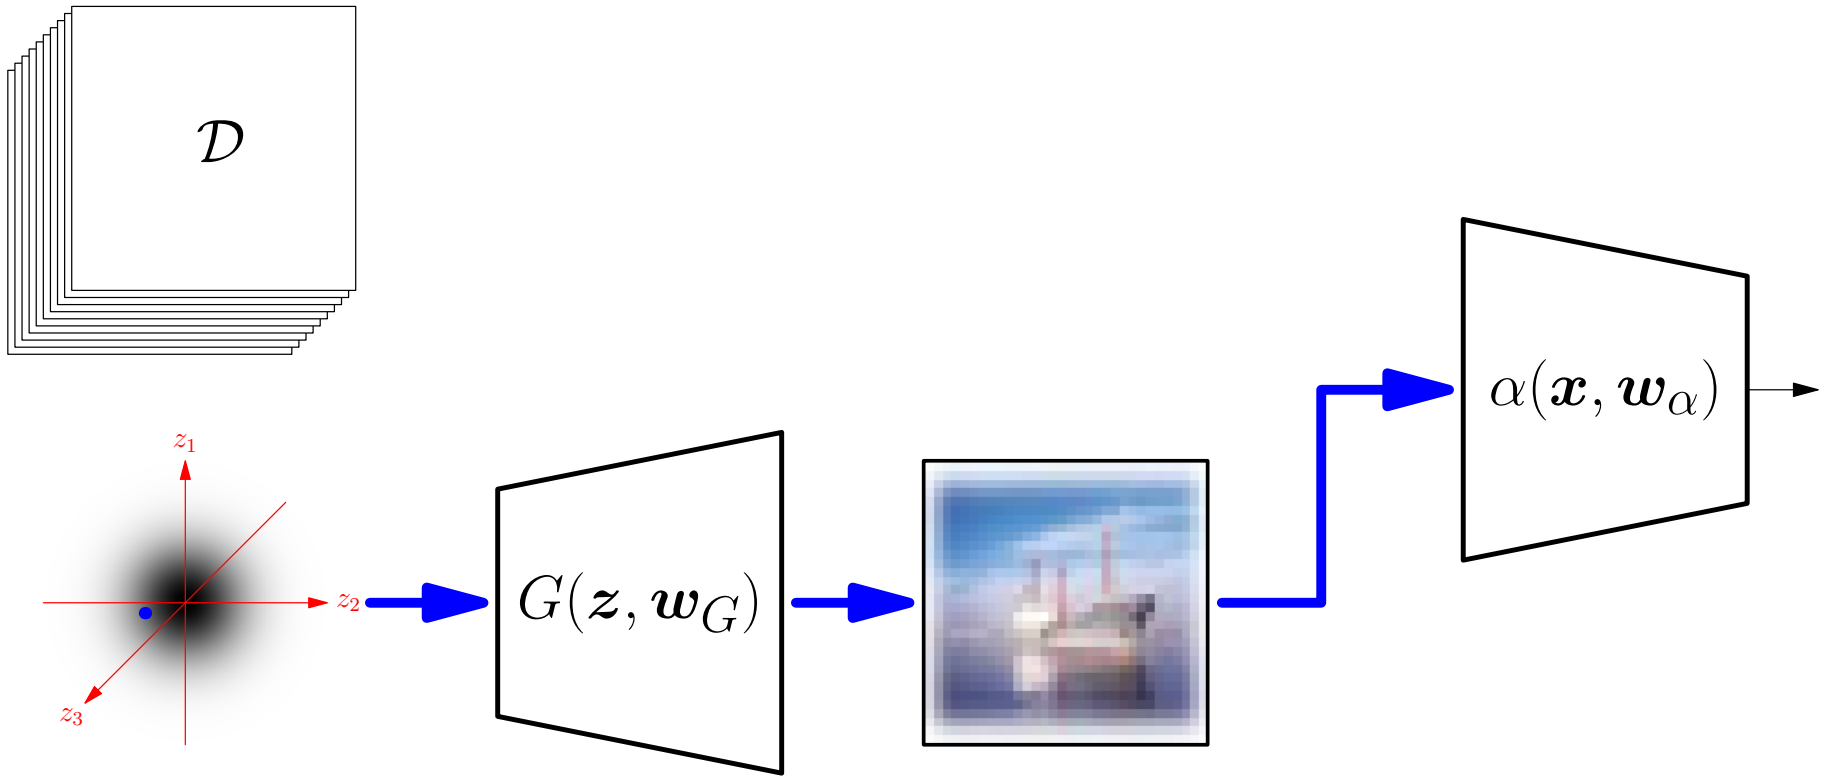
Wasserstein GANs

$$\max_{\boldsymbol{w}_\alpha} \min_{\boldsymbol{w}_G} \frac{1}{|\mathcal{B}|} \sum_{\boldsymbol{x} \in \mathcal{B}} \alpha(\boldsymbol{x}, \boldsymbol{w}_\alpha) - \frac{1}{n} \sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i, \boldsymbol{w}_G), \boldsymbol{w}_\alpha)$$



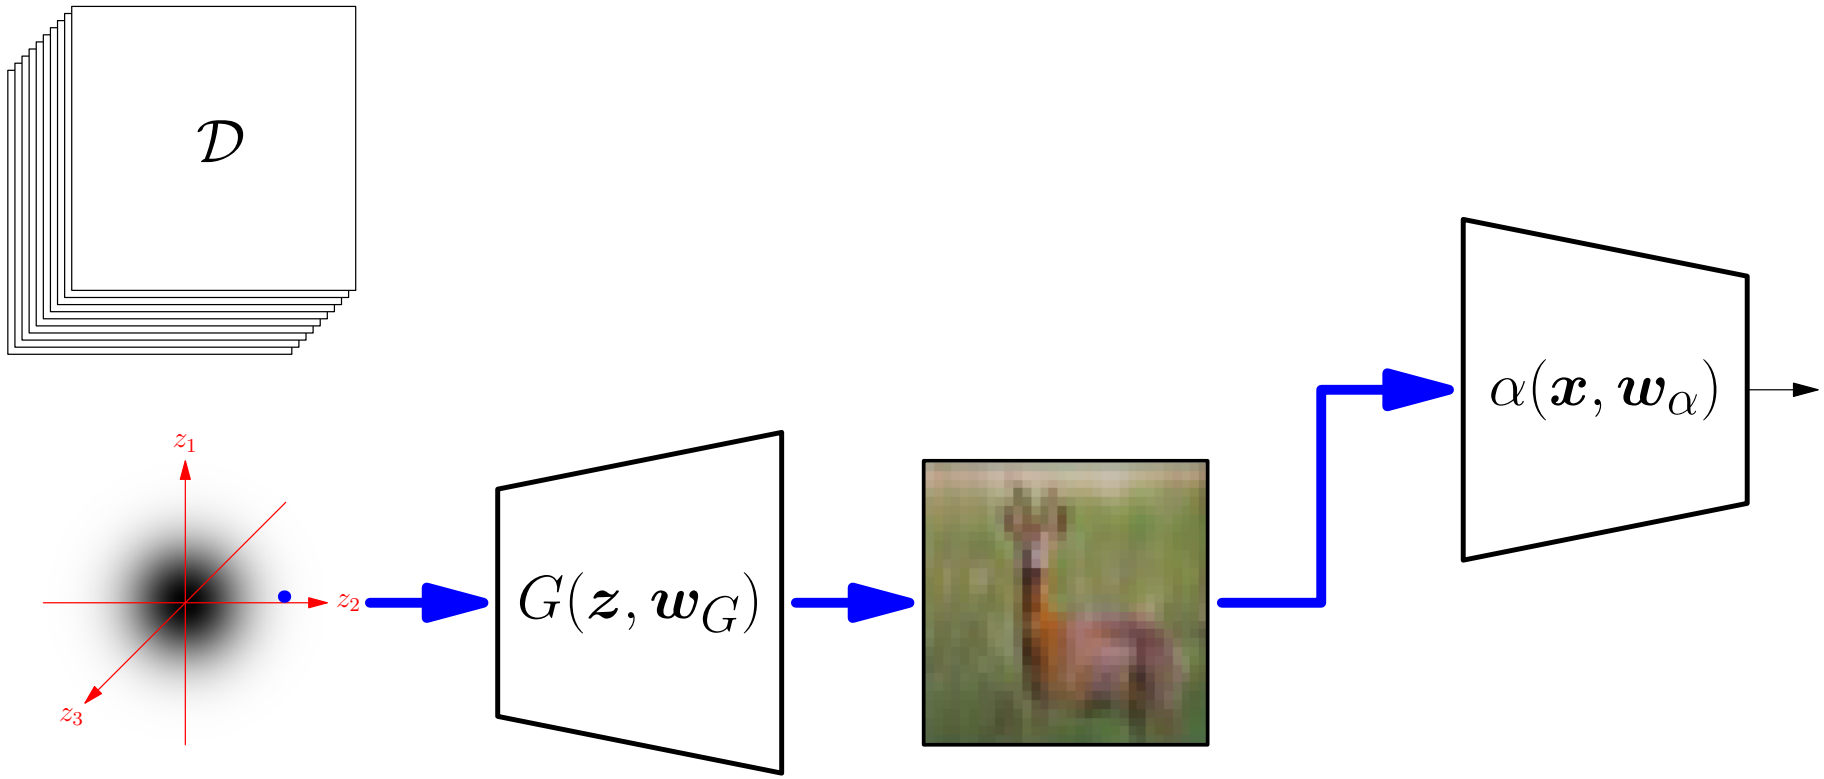
Wasserstein GANs

$$\max_{\boldsymbol{w}_\alpha} \min_{\boldsymbol{w}_G} \frac{1}{|\mathcal{B}|} \sum_{\boldsymbol{x} \in \mathcal{B}} \alpha(\boldsymbol{x}, \boldsymbol{w}_\alpha) - \frac{1}{n} \sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i, \boldsymbol{w}_G), \boldsymbol{w}_\alpha)$$



Wasserstein GANs

$$\max_{\boldsymbol{w}_\alpha} \min_{\boldsymbol{w}_G} \frac{1}{|\mathcal{B}|} \sum_{\boldsymbol{x} \in \mathcal{B}} \alpha(\boldsymbol{x}, \boldsymbol{w}_\alpha) - \frac{1}{n} \sum_{i=1}^{n} \alpha(G(\boldsymbol{z}_i, \boldsymbol{w}_G), \boldsymbol{w}_\alpha)$$



Entropic Regularisation

- Solving the linear programme for n points costs roughly $O(n^3)$: too slow for large problems

- Cuturi (2013) added an entropy term, with $\epsilon > 0$

$$\min_{\gamma} \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij} \log \gamma_{ij}$$

  with the same constraints

- The optimisation is given by: $\max_{\boldsymbol{\alpha},\boldsymbol{\beta}} \min_{\gamma} \mathcal{L}(\gamma, \boldsymbol{\alpha}, \boldsymbol{\beta})$

$$\mathcal{L} = \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij} \log \gamma_{ij} - \sum_i \alpha_i \left(\sum_j \gamma_{ij} - p_i \right) - \sum_j \beta_j \left(\sum_i \gamma_{ij} - q_j \right)$$

Entropic Regularisation

- Solving the linear programme for n points costs roughly $O(n^3)$: too slow for large problems

- Cuturi (2013) added an entropy term, with $\epsilon > 0$

$$\min_{\gamma} \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij} \log \gamma_{ij}$$

  with the same constraints

- The optimisation is given by: $\max_{\boldsymbol{\alpha},\boldsymbol{\beta}} \min_{\gamma} \mathcal{L}(\gamma, \boldsymbol{\alpha}, \boldsymbol{\beta})$

$$\mathcal{L} = \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij} \log \gamma_{ij} - \sum_i \alpha_i \left(\sum_j \gamma_{ij} - p_i \right) - \sum_j \beta_j \left(\sum_i \gamma_{ij} - q_j \right)$$

Entropic Regularisation

- Solving the linear programme for n points costs roughly $O(n^3)$: too slow for large problems

- Cuturi (2013) added an entropy term, with $\epsilon > 0$

$$\min_{\gamma} \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij} \log \gamma_{ij}$$

  with the same constraints

- The optimisation is given by: $\max_{\boldsymbol{\alpha},\boldsymbol{\beta}} \min_{\gamma} \mathcal{L}(\gamma, \boldsymbol{\alpha}, \boldsymbol{\beta})$

$$\mathcal{L} = \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij} \log \gamma_{ij} - \sum_i \alpha_i \left(\sum_j \gamma_{ij} - p_i \right) - \sum_j \beta_j \left(\sum_i \gamma_{ij} - q_j \right)$$

Solving the Primal Problem

- Using

$$\mathcal{L} = \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij} \log \gamma_{ij} - \sum_i \alpha_i \left(\sum_j \gamma_{ij} - p_i \right) - \sum_j \beta_j \left(\sum_i \gamma_{ij} - q_j \right)$$

- Setting the derivative of the Lagrangian to zero gives

$$\gamma_{ij} = \mathrm{e}^{(\alpha_i - d_{ij} + \beta_j)/\epsilon - 1} = u_i K_{ij} v_j$$

where $u_i = \mathrm{e}^{\alpha_i/\epsilon - 1}$, $v_j = \mathrm{e}^{\beta_j/\epsilon}$ and $K_{ij} = \mathrm{e}^{-d_{ij}/\epsilon}$

- We have to choose u_i and v_j (equivalently the Lagrange multipliers) so that

$$\sum_j u_i K_{ij} v_j = p_i \qquad \sum_i u_i K_{ij} v_j = q_j$$

Solving the Primal Problem

- Using

$$\mathcal{L} = \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij}\log\gamma_{ij} - \sum_i \alpha_i \left(\sum_j \gamma_{ij} - p_i \right) - \sum_j \beta_j \left(\sum_i \gamma_{ij} - q_j \right)$$

- Setting the derivative of the Lagrangian to zero gives

$$\gamma_{ij} = \mathrm{e}^{(\alpha_i - d_{ij} + \beta_j)/\epsilon - 1} = u_i K_{ij} v_j$$

where $u_i = \mathrm{e}^{\alpha_i/\epsilon - 1}$, $v_j = \mathrm{e}^{\beta_j/\epsilon}$ and $K_{ij} = \mathrm{e}^{-d_{ij}/\epsilon}$

- We have to choose u_i and v_j (equivalently the Lagrange multipliers) so that

$$\sum_j u_i K_{ij} v_j = p_i \qquad \sum_i u_i K_{ij} v_j = q_j$$

Solving the Primal Problem

- Using

$$\mathcal{L} = \sum_{i,j} d_{ij}\gamma_{ij} + \epsilon \sum_{i,j} \gamma_{ij}\log\gamma_{ij} - \sum_i \alpha_i \left(\sum_j \gamma_{ij} - p_i \right) - \sum_j \beta_j \left(\sum_i \gamma_{ij} - q_j \right)$$

- Setting the derivative of the Lagrangian to zero gives

$$\gamma_{ij} = \mathrm{e}^{(\alpha_i - d_{ij} + \beta_j)/\epsilon - 1} = u_i K_{ij} v_j$$

where $u_i = \mathrm{e}^{\alpha_i/\epsilon - 1}$, $v_j = \mathrm{e}^{\beta_j/\epsilon}$ and $K_{ij} = \mathrm{e}^{-d_{ij}/\epsilon}$

- We have to choose u_i and v_j (equivalently the Lagrange multipliers) so that

$$\sum_j u_i K_{ij} v_j = p_i \qquad \sum_i u_i K_{ij} v_j = q_j$$

The Sinkhorn Algorithm

- Choose $\boldsymbol{u}$ and $\boldsymbol{v}$ so the row and column sums are $\boldsymbol{p}$ and $\boldsymbol{q}$: alternately rescale ($\oslash$ is element-by-element division)

$$\boldsymbol{v} \leftarrow \boldsymbol{q} \oslash (\mathbf{K}^{\mathsf{T}}\boldsymbol{u}), \qquad \boldsymbol{u} \leftarrow \boldsymbol{p} \oslash (\mathbf{K}\boldsymbol{v})$$

- This is Sinkhorn's matrix-scaling algorithm: each step is two matrix–vector products
- Fast on GPUs and differentiable, so it can sit inside a loss
- As $\epsilon \to 0$ we recover the exact transport cost; larger ϵ is faster but blurs the plan
- This is how optimal transport is mostly computed in ML today

The Sinkhorn Algorithm

- Choose $\boldsymbol{u}$ and $\boldsymbol{v}$ so the row and column sums are $\boldsymbol{p}$ and $\boldsymbol{q}$: alternately rescale ($\oslash$ is element-by-element division)

$$\boldsymbol{v} \leftarrow \boldsymbol{q} \oslash (\mathbf{K}^{\mathsf{T}}\boldsymbol{u}), \qquad \boldsymbol{u} \leftarrow \boldsymbol{p} \oslash (\mathbf{K}\boldsymbol{v})$$

- This is Sinkhorn's matrix-scaling algorithm: each step is two matrix–vector products
- Fast on GPUs and differentiable, so it can sit inside a loss
- As $\epsilon \to 0$ we recover the exact transport cost; larger ϵ is faster but blurs the plan
- This is how optimal transport is mostly computed in ML today

The Sinkhorn Algorithm

- Choose $\boldsymbol{u}$ and $\boldsymbol{v}$ so the row and column sums are $\boldsymbol{p}$ and $\boldsymbol{q}$: alternately rescale ($\oslash$ is element-by-element division)

$$\boldsymbol{v} \leftarrow \boldsymbol{q} \oslash (\mathbf{K}^{\mathsf{T}}\boldsymbol{u}), \qquad \boldsymbol{u} \leftarrow \boldsymbol{p} \oslash (\mathbf{K}\boldsymbol{v})$$

- This is Sinkhorn's matrix-scaling algorithm: each step is two matrix–vector products
- Fast on GPUs and differentiable, so it can sit inside a loss
- As $\epsilon \to 0$ we recover the exact transport cost; larger ϵ is faster but blurs the plan
- This is how optimal transport is mostly computed in ML today

The Sinkhorn Algorithm

- Choose $\boldsymbol{u}$ and $\boldsymbol{v}$ so the row and column sums are $\boldsymbol{p}$ and $\boldsymbol{q}$: alternately rescale ($\oslash$ is element-by-element division)

$$\boldsymbol{v} \leftarrow \boldsymbol{q} \oslash (\mathbf{K}^{\mathsf{T}}\boldsymbol{u}), \qquad \boldsymbol{u} \leftarrow \boldsymbol{p} \oslash (\mathbf{K}\boldsymbol{v})$$

- This is Sinkhorn's matrix-scaling algorithm: each step is two matrix–vector products
- Fast on GPUs and differentiable, so it can sit inside a loss
- As $\epsilon \to 0$ we recover the exact transport cost; larger ϵ is faster but blurs the plan
- This is how optimal transport is mostly computed in ML today

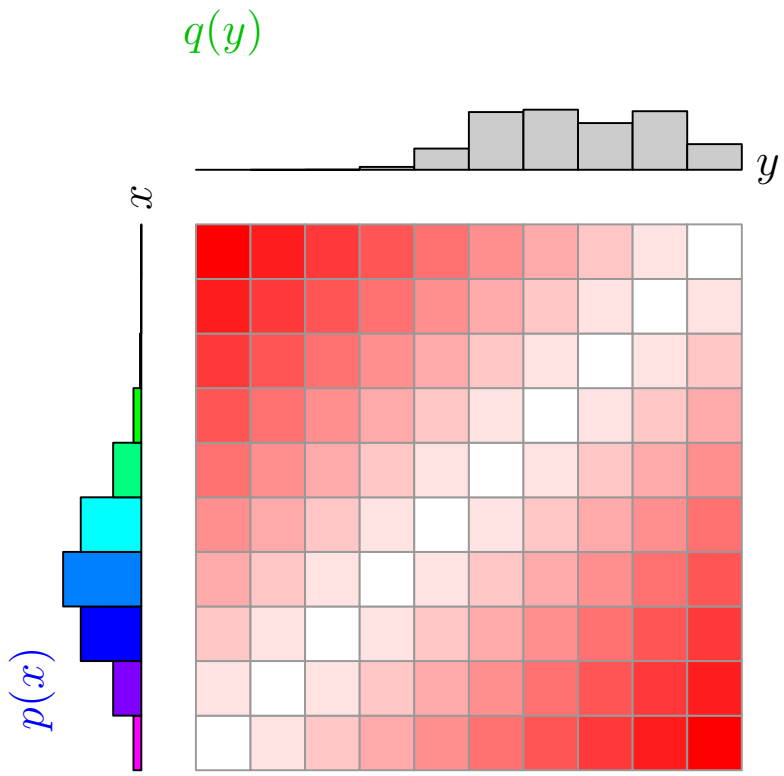
The Sinkhorn Algorithm

- Choose $\boldsymbol{u}$ and $\boldsymbol{v}$ so the row and column sums are $\boldsymbol{p}$ and $\boldsymbol{q}$: alternately rescale ($\oslash$ is element-by-element division)

$$\boldsymbol{v} \leftarrow \boldsymbol{q} \oslash (\mathbf{K}^{\mathsf{T}}\boldsymbol{u}), \qquad \boldsymbol{u} \leftarrow \boldsymbol{p} \oslash (\mathbf{K}\boldsymbol{v})$$

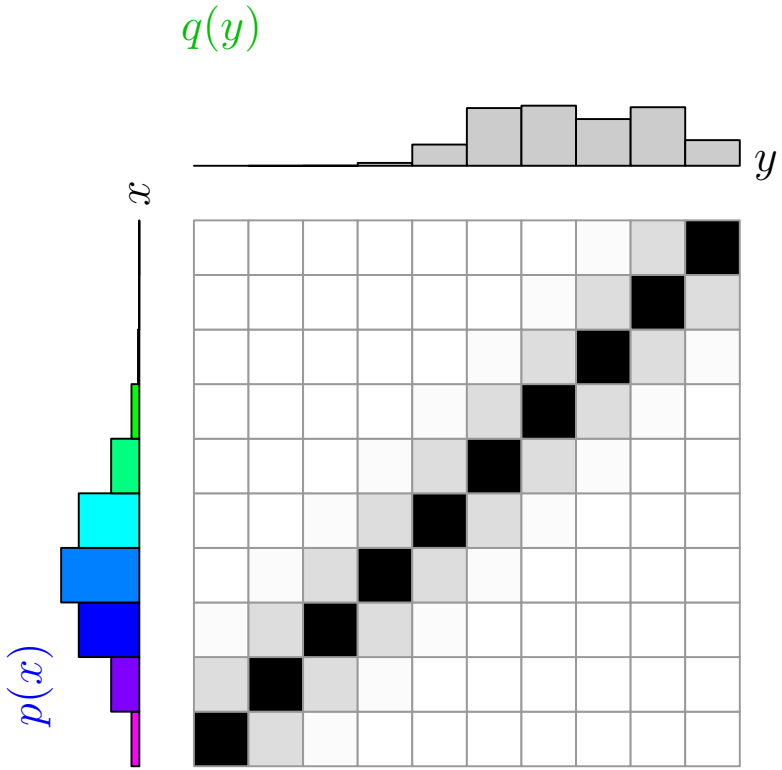
- This is Sinkhorn's matrix-scaling algorithm: each step is two matrix–vector products
- Fast on GPUs and differentiable, so it can sit inside a loss
- As $\epsilon \to 0$ we recover the exact transport cost; larger ϵ is faster but blurs the plan
- This is how optimal transport is mostly computed in ML today

Sinkhorn in Action



$$\text{cost } d_{ij} = |x_i - y_j|$$

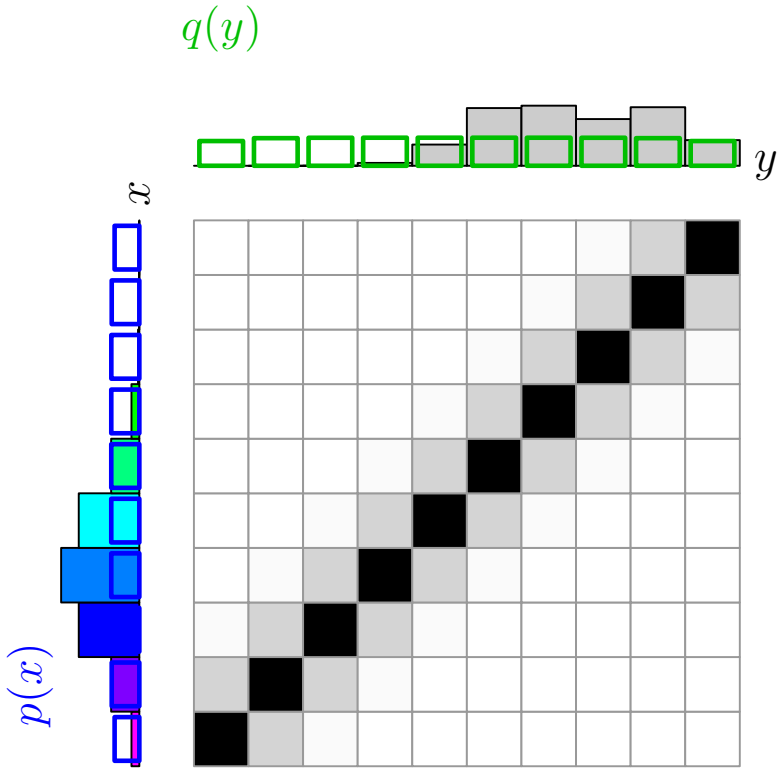
Sinkhorn in Action



kernel $K_{ij} = \mathrm{e}^{-d_{ij}/\epsilon}$, $\epsilon = 0.5$

$$\gamma = \mathrm{diag}(\boldsymbol{u})\,\mathbf{K}\,\mathrm{diag}(\boldsymbol{v})$$

Sinkhorn in Action

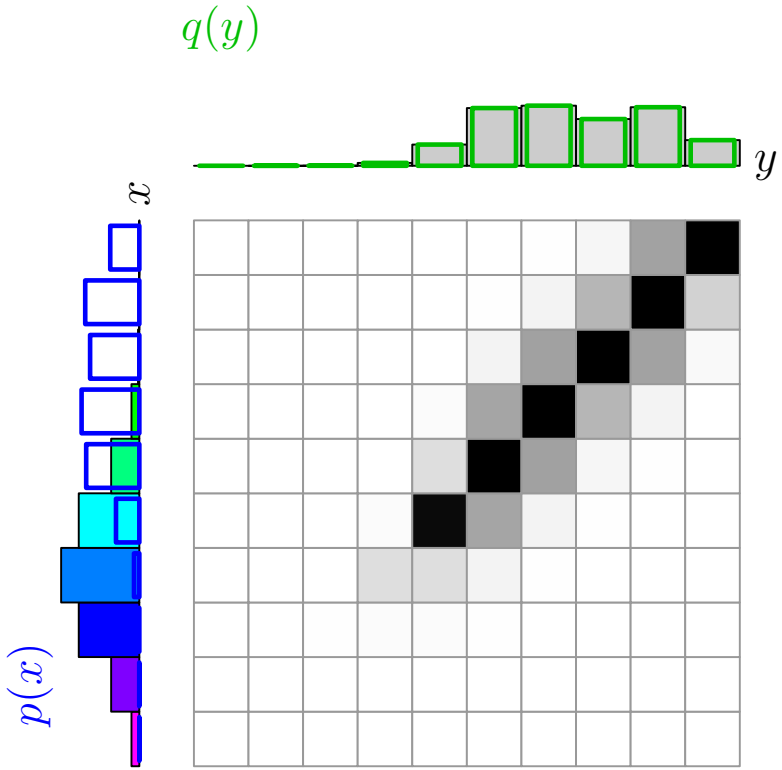


start: $\gamma \propto \mathbf{K}$

$\boldsymbol{u} = \mathbf{1}/\sum_{ij} K_{ij}$, $\boldsymbol{v} = \mathbf{1}$

$\gamma = \mathrm{diag}(\boldsymbol{u})\,\mathbf{K}\,\mathrm{diag}(\boldsymbol{v})$

Sinkhorn in Action

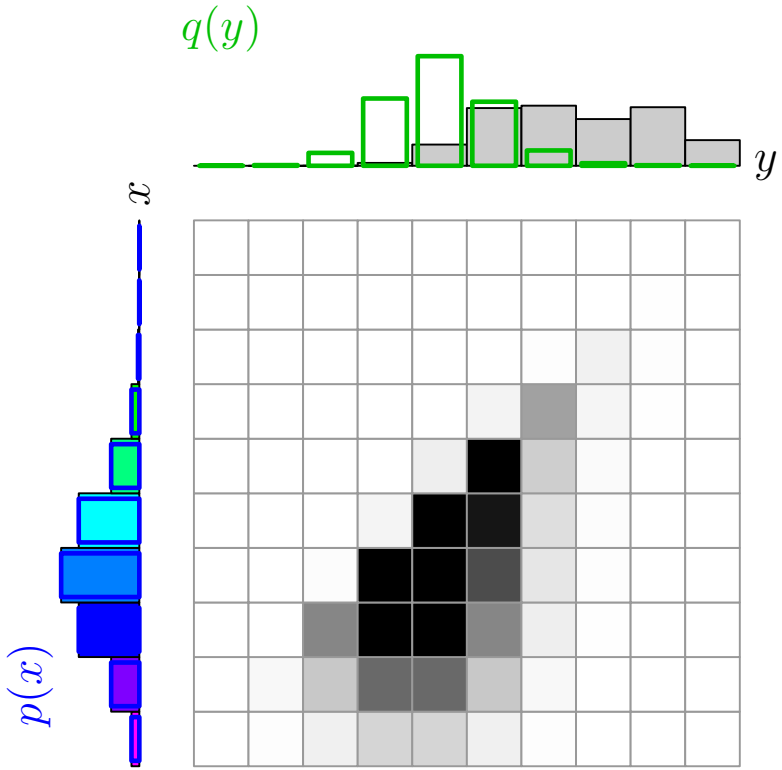


iteration 1: columns now sum to $\boldsymbol{q}$

$$\boldsymbol{v} \leftarrow \boldsymbol{q} \oslash (\mathbf{K}^{\mathsf{T}} \boldsymbol{u})$$

$$\gamma = \mathrm{diag}(\boldsymbol{u})\, \mathbf{K}\, \mathrm{diag}(\boldsymbol{v})$$

Sinkhorn in Action

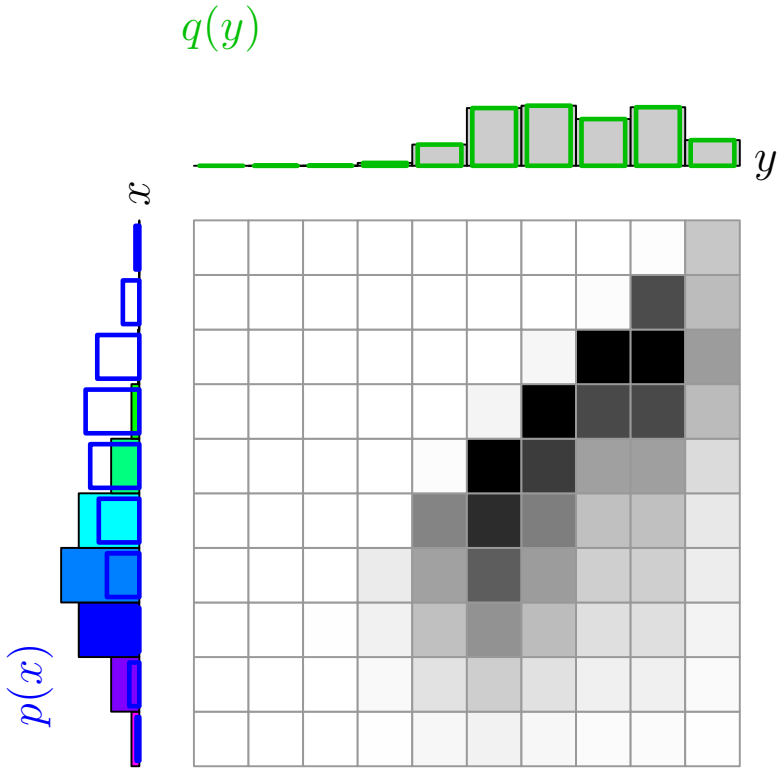


iteration 1: rows now sum to $\boldsymbol{p}$

$$\boldsymbol{u} \leftarrow \boldsymbol{p} \oslash (\mathbf{K}\boldsymbol{v})$$

$$\gamma = \mathrm{diag}(\boldsymbol{u})\,\mathbf{K}\,\mathrm{diag}(\boldsymbol{v})$$

Sinkhorn in Action

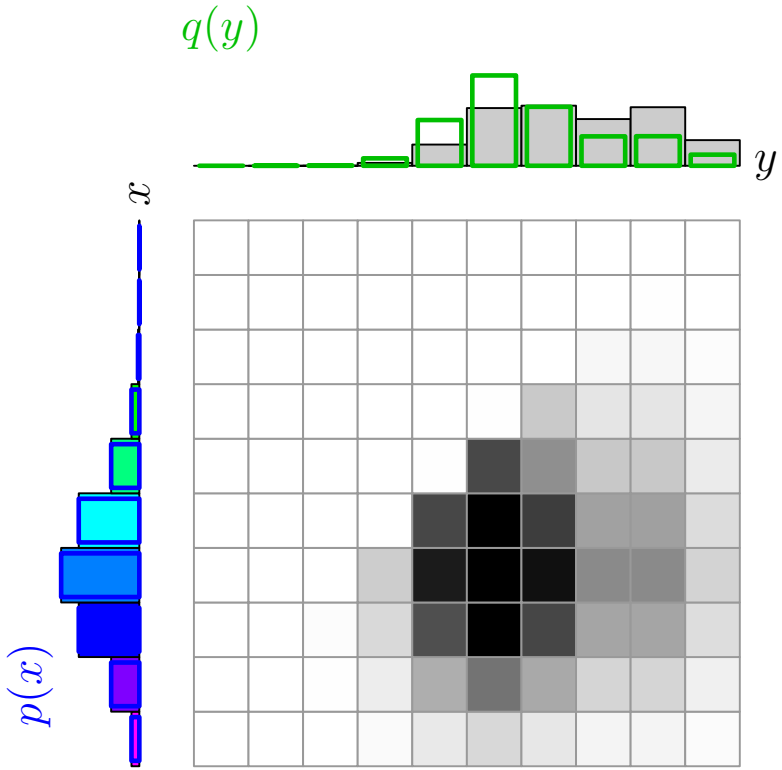


iteration 2: columns now sum to $\boldsymbol{q}$

$$\boldsymbol{v} \leftarrow \boldsymbol{q} \oslash (\mathbf{K}^{\mathsf{T}} \boldsymbol{u})$$

$$\gamma = \mathrm{diag}(\boldsymbol{u})\, \mathbf{K}\, \mathrm{diag}(\boldsymbol{v})$$

Sinkhorn in Action

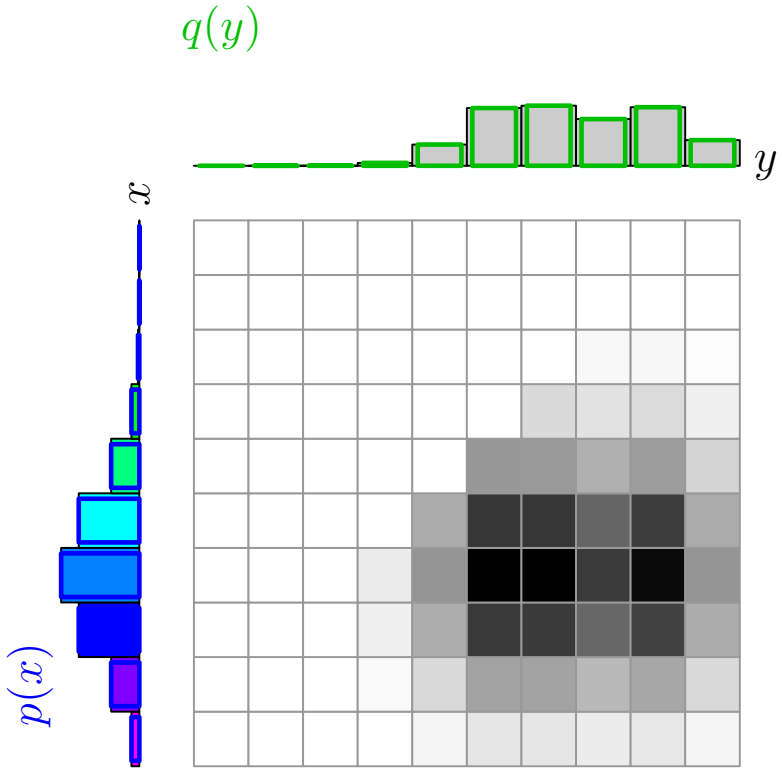


iteration 2: rows now sum to $\boldsymbol{p}$

$$\boldsymbol{u} \leftarrow \boldsymbol{p} \oslash (\mathbf{K}\boldsymbol{v})$$

$$\gamma = \mathrm{diag}(\boldsymbol{u})\,\mathbf{K}\,\mathrm{diag}(\boldsymbol{v})$$

Sinkhorn in Action



after 10 iterations: cost $= 3.460$

exact Wasserstein distance $= 3.460$

$$\gamma = \mathrm{diag}(\boldsymbol{u})\,\mathbf{K}\,\mathrm{diag}(\boldsymbol{v})$$

Evaluating Generative Models: FID

- For Gaussians, W_2 has a closed form

$$W_2^2 = \|\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2\|^2 + \mathrm{Tr}\Big(\boldsymbol{\Sigma}_1 + \boldsymbol{\Sigma}_2 - 2\big(\boldsymbol{\Sigma}_1\boldsymbol{\Sigma}_2\big)^{1/2}\Big)$$

  (in one dimension $(\mu_1 - \mu_2)^2 + (\sigma_1 - \sigma_2)^2$)

- The **Fréchet Inception Distance** (Heusel et al., 2017) fits Gaussians to a pretrained network's features of real and of generated images, and reports this distance

- It is the standard score for image generators, GANs and diffusion models alike (lower is better)

- Caveat: it only compares means and covariances of the features

Evaluating Generative Models: FID

- For Gaussians, W_2 has a closed form

$$W_2^2 = \|\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2\|^2 + \mathrm{Tr}\Big(\boldsymbol{\Sigma}_1 + \boldsymbol{\Sigma}_2 - 2\big(\boldsymbol{\Sigma}_1\boldsymbol{\Sigma}_2\big)^{1/2}\Big)$$

  (in one dimension $(\mu_1 - \mu_2)^2 + (\sigma_1 - \sigma_2)^2$)

- The **Fréchet Inception Distance** (Heusel et al., 2017) fits Gaussians to a pretrained network's features of real and of generated images, and reports this distance

- It is the standard score for image generators, GANs and diffusion models alike (lower is better)

- Caveat: it only compares means and covariances of the features

Evaluating Generative Models: FID

- For Gaussians, W_2 has a closed form

$$W_2^2 = \|\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2\|^2 + \mathrm{Tr}\Big(\boldsymbol{\Sigma}_1 + \boldsymbol{\Sigma}_2 - 2\big(\boldsymbol{\Sigma}_1\boldsymbol{\Sigma}_2\big)^{1/2}\Big)$$

  (in one dimension $(\mu_1 - \mu_2)^2 + (\sigma_1 - \sigma_2)^2$)

- The **Fréchet Inception Distance** (Heusel et al., 2017) fits Gaussians to a pretrained network's features of real and of generated images, and reports this distance

- It is the standard score for image generators, GANs and diffusion models alike (lower is better)

- Caveat: it only compares means and covariances of the features

Evaluating Generative Models: FID

- For Gaussians, W_2 has a closed form

$$W_2^2 = \|\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2\|^2 + \mathrm{Tr}\Big(\boldsymbol{\Sigma}_1 + \boldsymbol{\Sigma}_2 - 2\big(\boldsymbol{\Sigma}_1\boldsymbol{\Sigma}_2\big)^{1/2}\Big)$$

  (in one dimension $(\mu_1 - \mu_2)^2 + (\sigma_1 - \sigma_2)^2$)

- The **Fréchet Inception Distance** (Heusel et al., 2017) fits Gaussians to a pretrained network's features of real and of generated images, and reports this distance

- It is the standard score for image generators, GANs and diffusion models alike (lower is better)

- Caveat: it only compares means and covariances of the features

Evaluating Generative Models: FID

- For Gaussians, W_2 has a closed form

$$W_2^2 = \|\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2\|^2 + \mathrm{Tr}\Big(\boldsymbol{\Sigma}_1 + \boldsymbol{\Sigma}_2 - 2\big(\boldsymbol{\Sigma}_1\boldsymbol{\Sigma}_2\big)^{1/2}\Big)$$

  (in one dimension $(\mu_1 - \mu_2)^2 + (\sigma_1 - \sigma_2)^2$)

- The **Fréchet Inception Distance** (Heusel et al., 2017) fits Gaussians to a pretrained network's features of real and of generated images, and reports this distance

- It is the standard score for image generators, GANs and diffusion models alike (lower is better)

- Caveat: it only compares means and covariances of the features

Lessons

- The Wasserstein distance compares distributions by how far probability mass must move, so it behaves well when they do not overlap
- Computing it is a linear programme, whose dual is a maximum over 1-Lipschitz functions
- That dual turned GANs into Wasserstein GANs: an elegant piece of theory
- In practice optimal transport is computed with Sinkhorn, and the W_2 distance between Gaussians is the FID used to evaluate generative models

Lessons

- The Wasserstein distance compares distributions by how far probability mass must move, so it behaves well when they do not overlap
- Computing it is a linear programme, whose dual is a maximum over 1-Lipschitz functions
- That dual turned GANs into Wasserstein GANs: an elegant piece of theory
- In practice optimal transport is computed with Sinkhorn, and the W_2 distance between Gaussians is the FID used to evaluate generative models

Lessons

- The Wasserstein distance compares distributions by how far probability mass must move, so it behaves well when they do not overlap
- Computing it is a linear programme, whose dual is a maximum over 1-Lipschitz functions
- That dual turned GANs into Wasserstein GANs: an elegant piece of theory
- In practice optimal transport is computed with Sinkhorn, and the W_2 distance between Gaussians is the FID used to evaluate generative models

Lessons

- The Wasserstein distance compares distributions by how far probability mass must move, so it behaves well when they do not overlap
- Computing it is a linear programme, whose dual is a maximum over 1-Lipschitz functions
- That dual turned GANs into Wasserstein GANs: an elegant piece of theory
- In practice optimal transport is computed with Sinkhorn, and the W_2 distance between Gaussians is the FID used to evaluate generative models